

INF721

2024/2



Deep Learning

L17: Transformers

Logistics

Last Lecture

- ▶ Machine Translation
- ▶ Decoding
 - ▶ Greedy Search
 - ▶ Beam Search
- ▶ Attention in RNNs

Lecture Outline

- ▶ Machine Translation
- ▶ Problems with RNNs
- ▶ Transformers
 - ▶ Self-Attention
 - ▶ Multi-head Attention
 - ▶ Encoder & Decoder
 - ▶ Positional Encoding
 - ▶ Masked Multi-head Attention

Machine Translation

Given a dataset of sentence pairs:

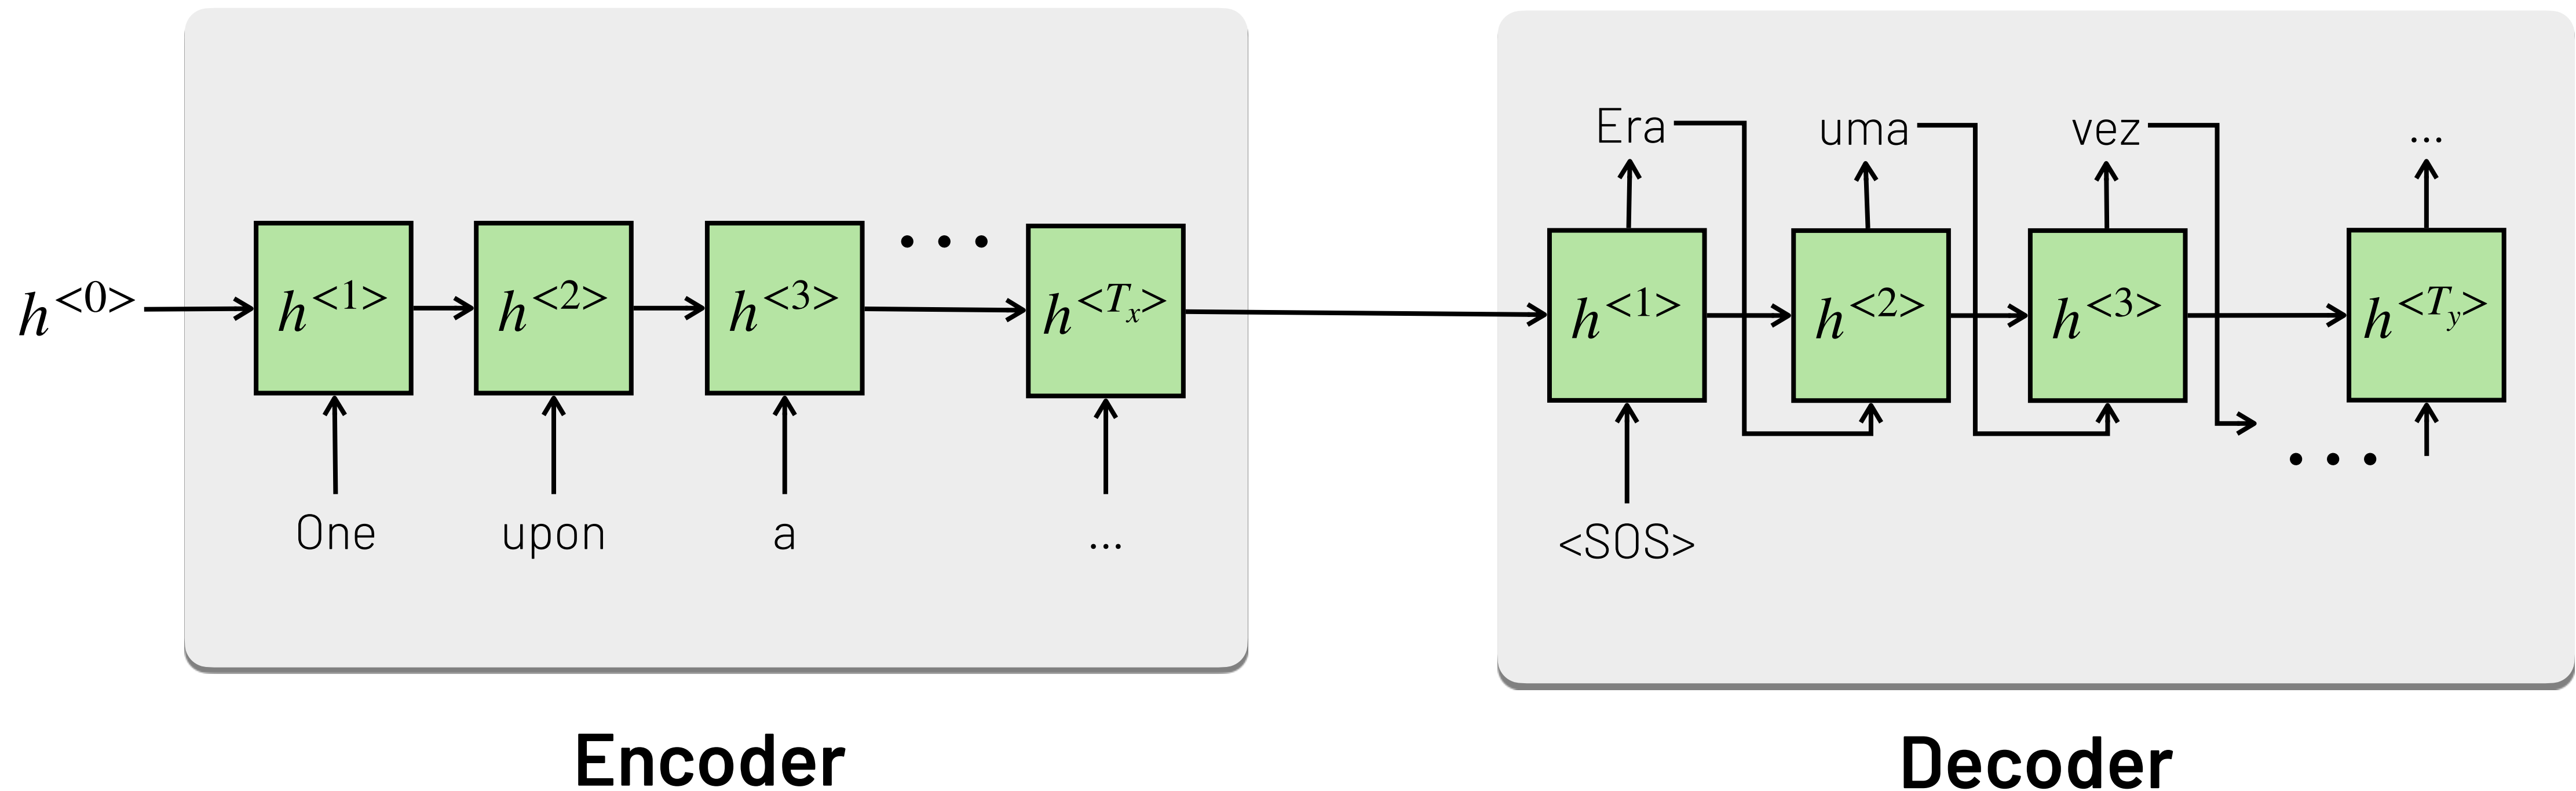
$$(x = \{x^{<1>}, x^{<2>}, \dots, x^{<T_x>}\}, y = \{y^{<1>}, y^{<2>}, \dots, y^{<T_y>}\}),$$

we want to learn a model that maps x into y .

Portuguese	English
Olá, como vai você?	Hello, how are you?
O livro está em cima da mesa.	The book is on the table.
Lucas irá viajar ao Rio em Dezembro.	Lucas is travelling to Rio in December.
Em Dezembro, Lucas irá viajar ao Rio.	Lucas is travelling to Rio in December.
....

Problems with RNNs

- ▶ Struggle to capture long dependencies in sequences
- ▶ Hard to parallelize

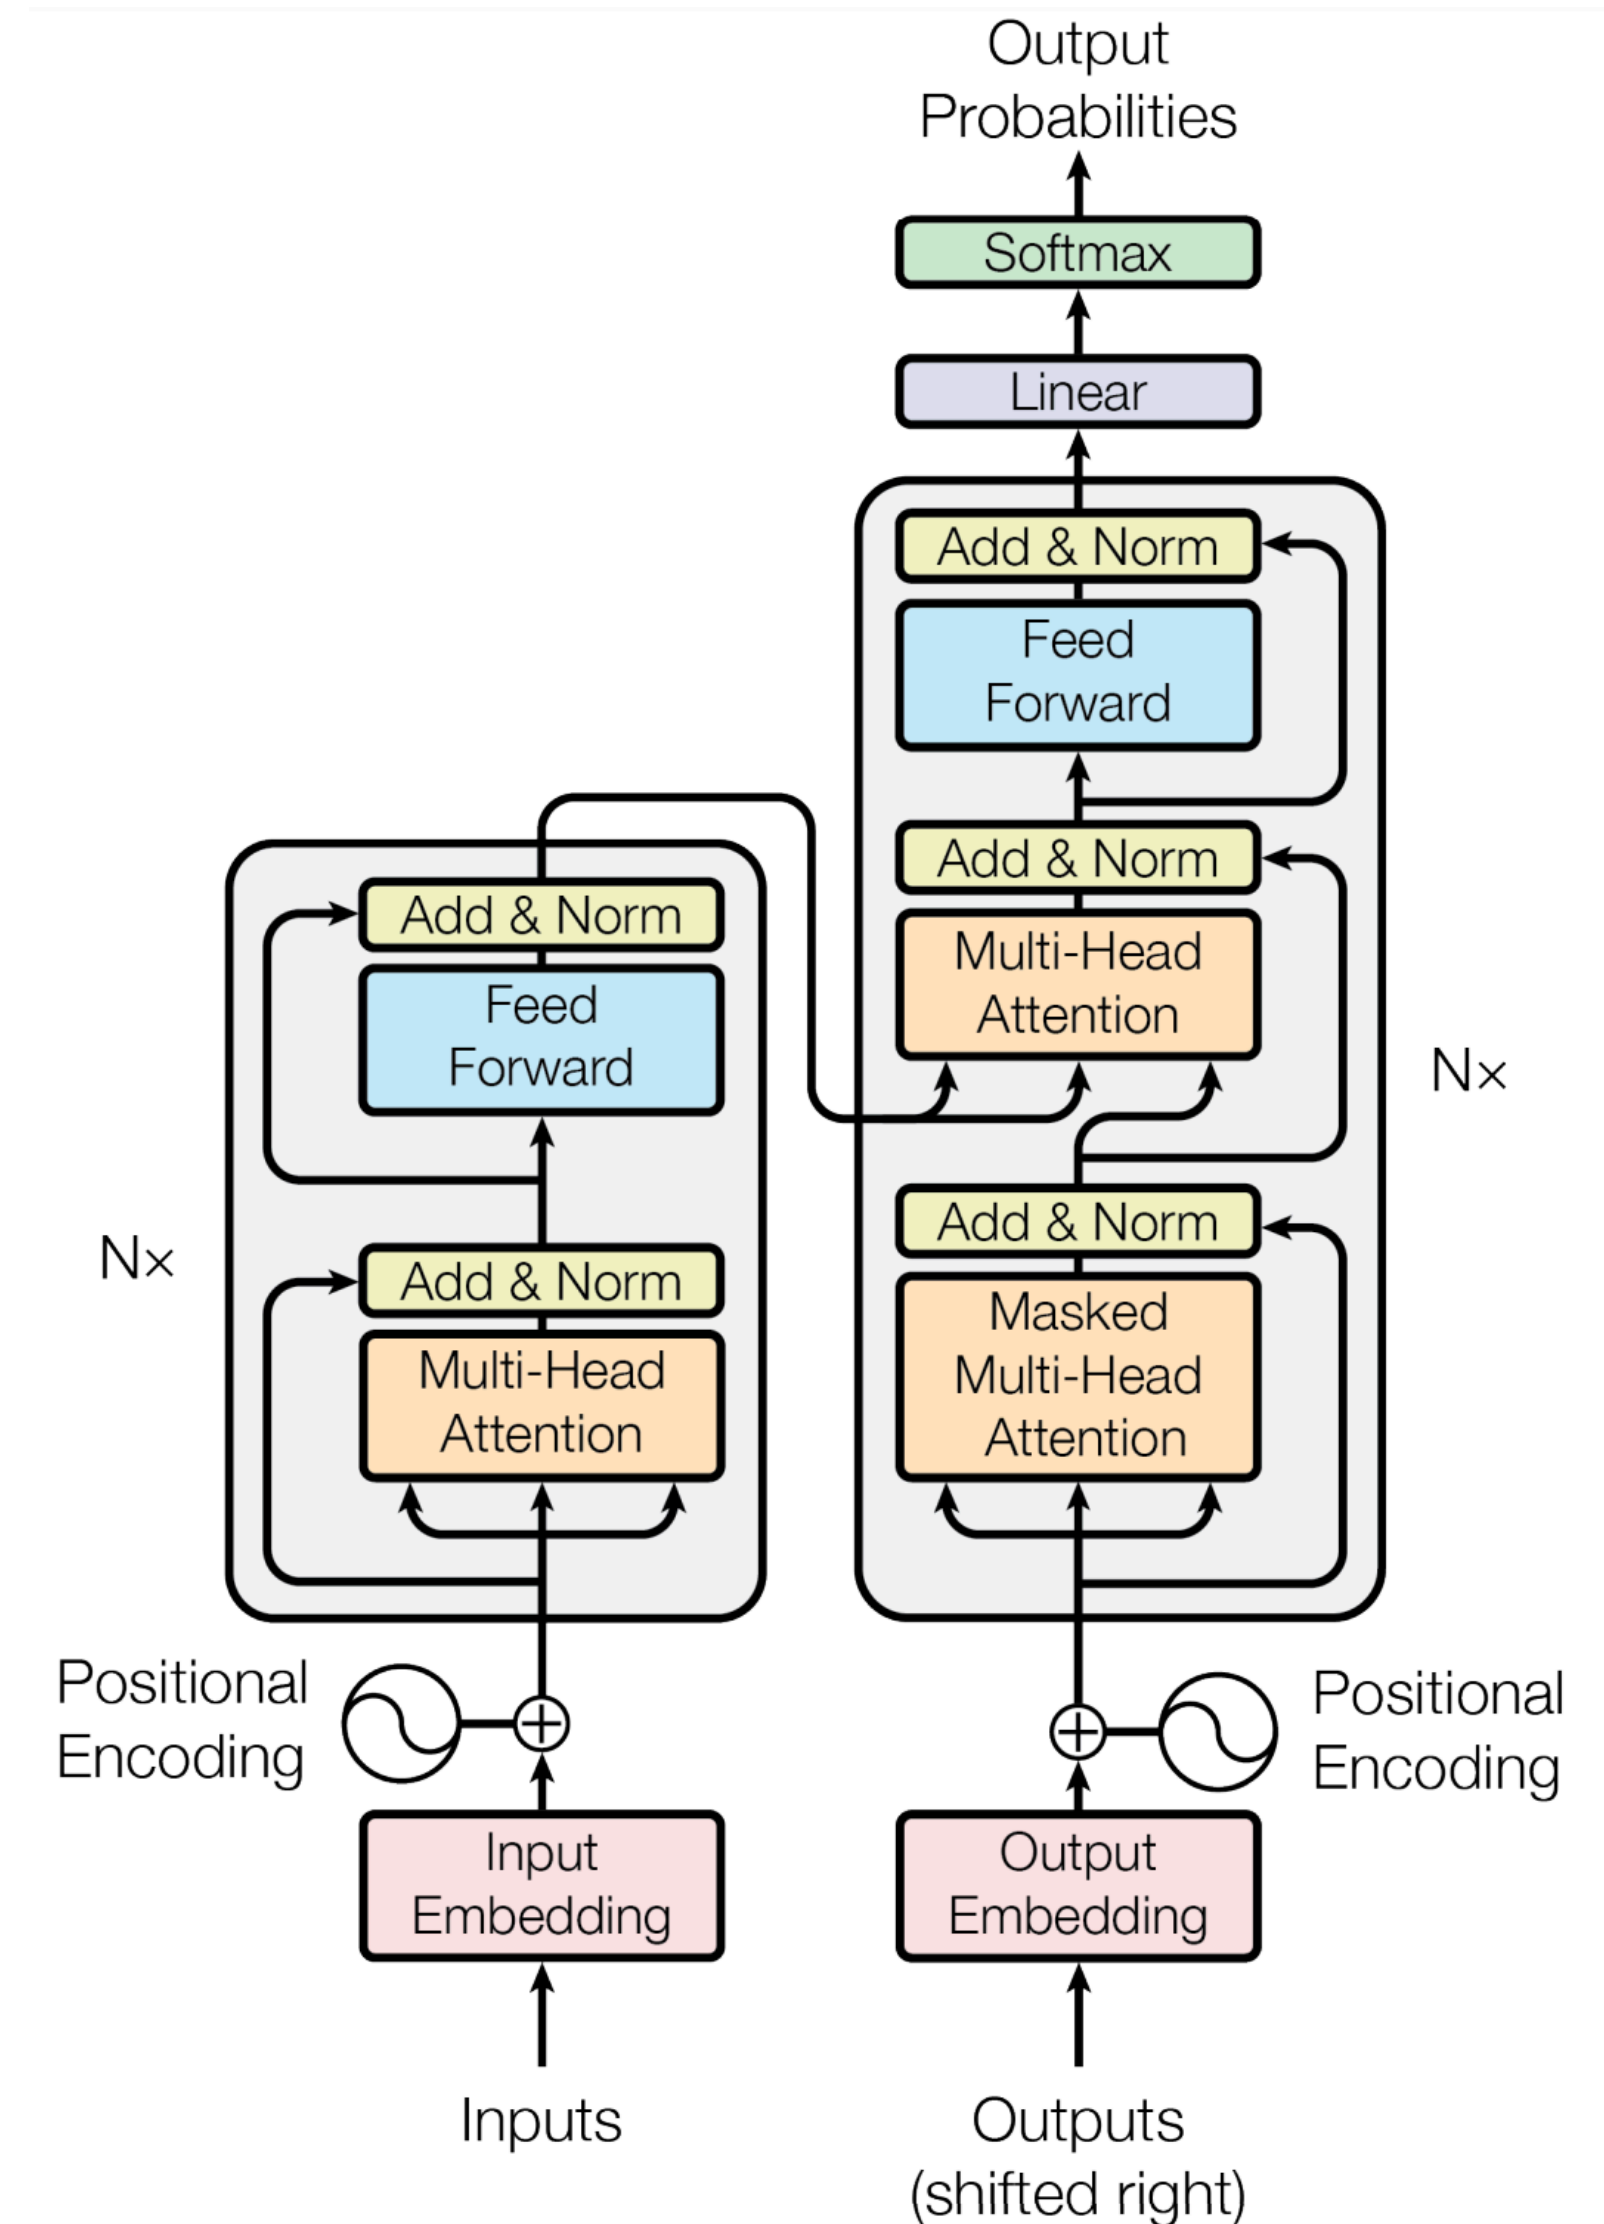


Transformers

Transformers are an encoder-decoder architecture to process sequences using only attention (eliminating recurrence).

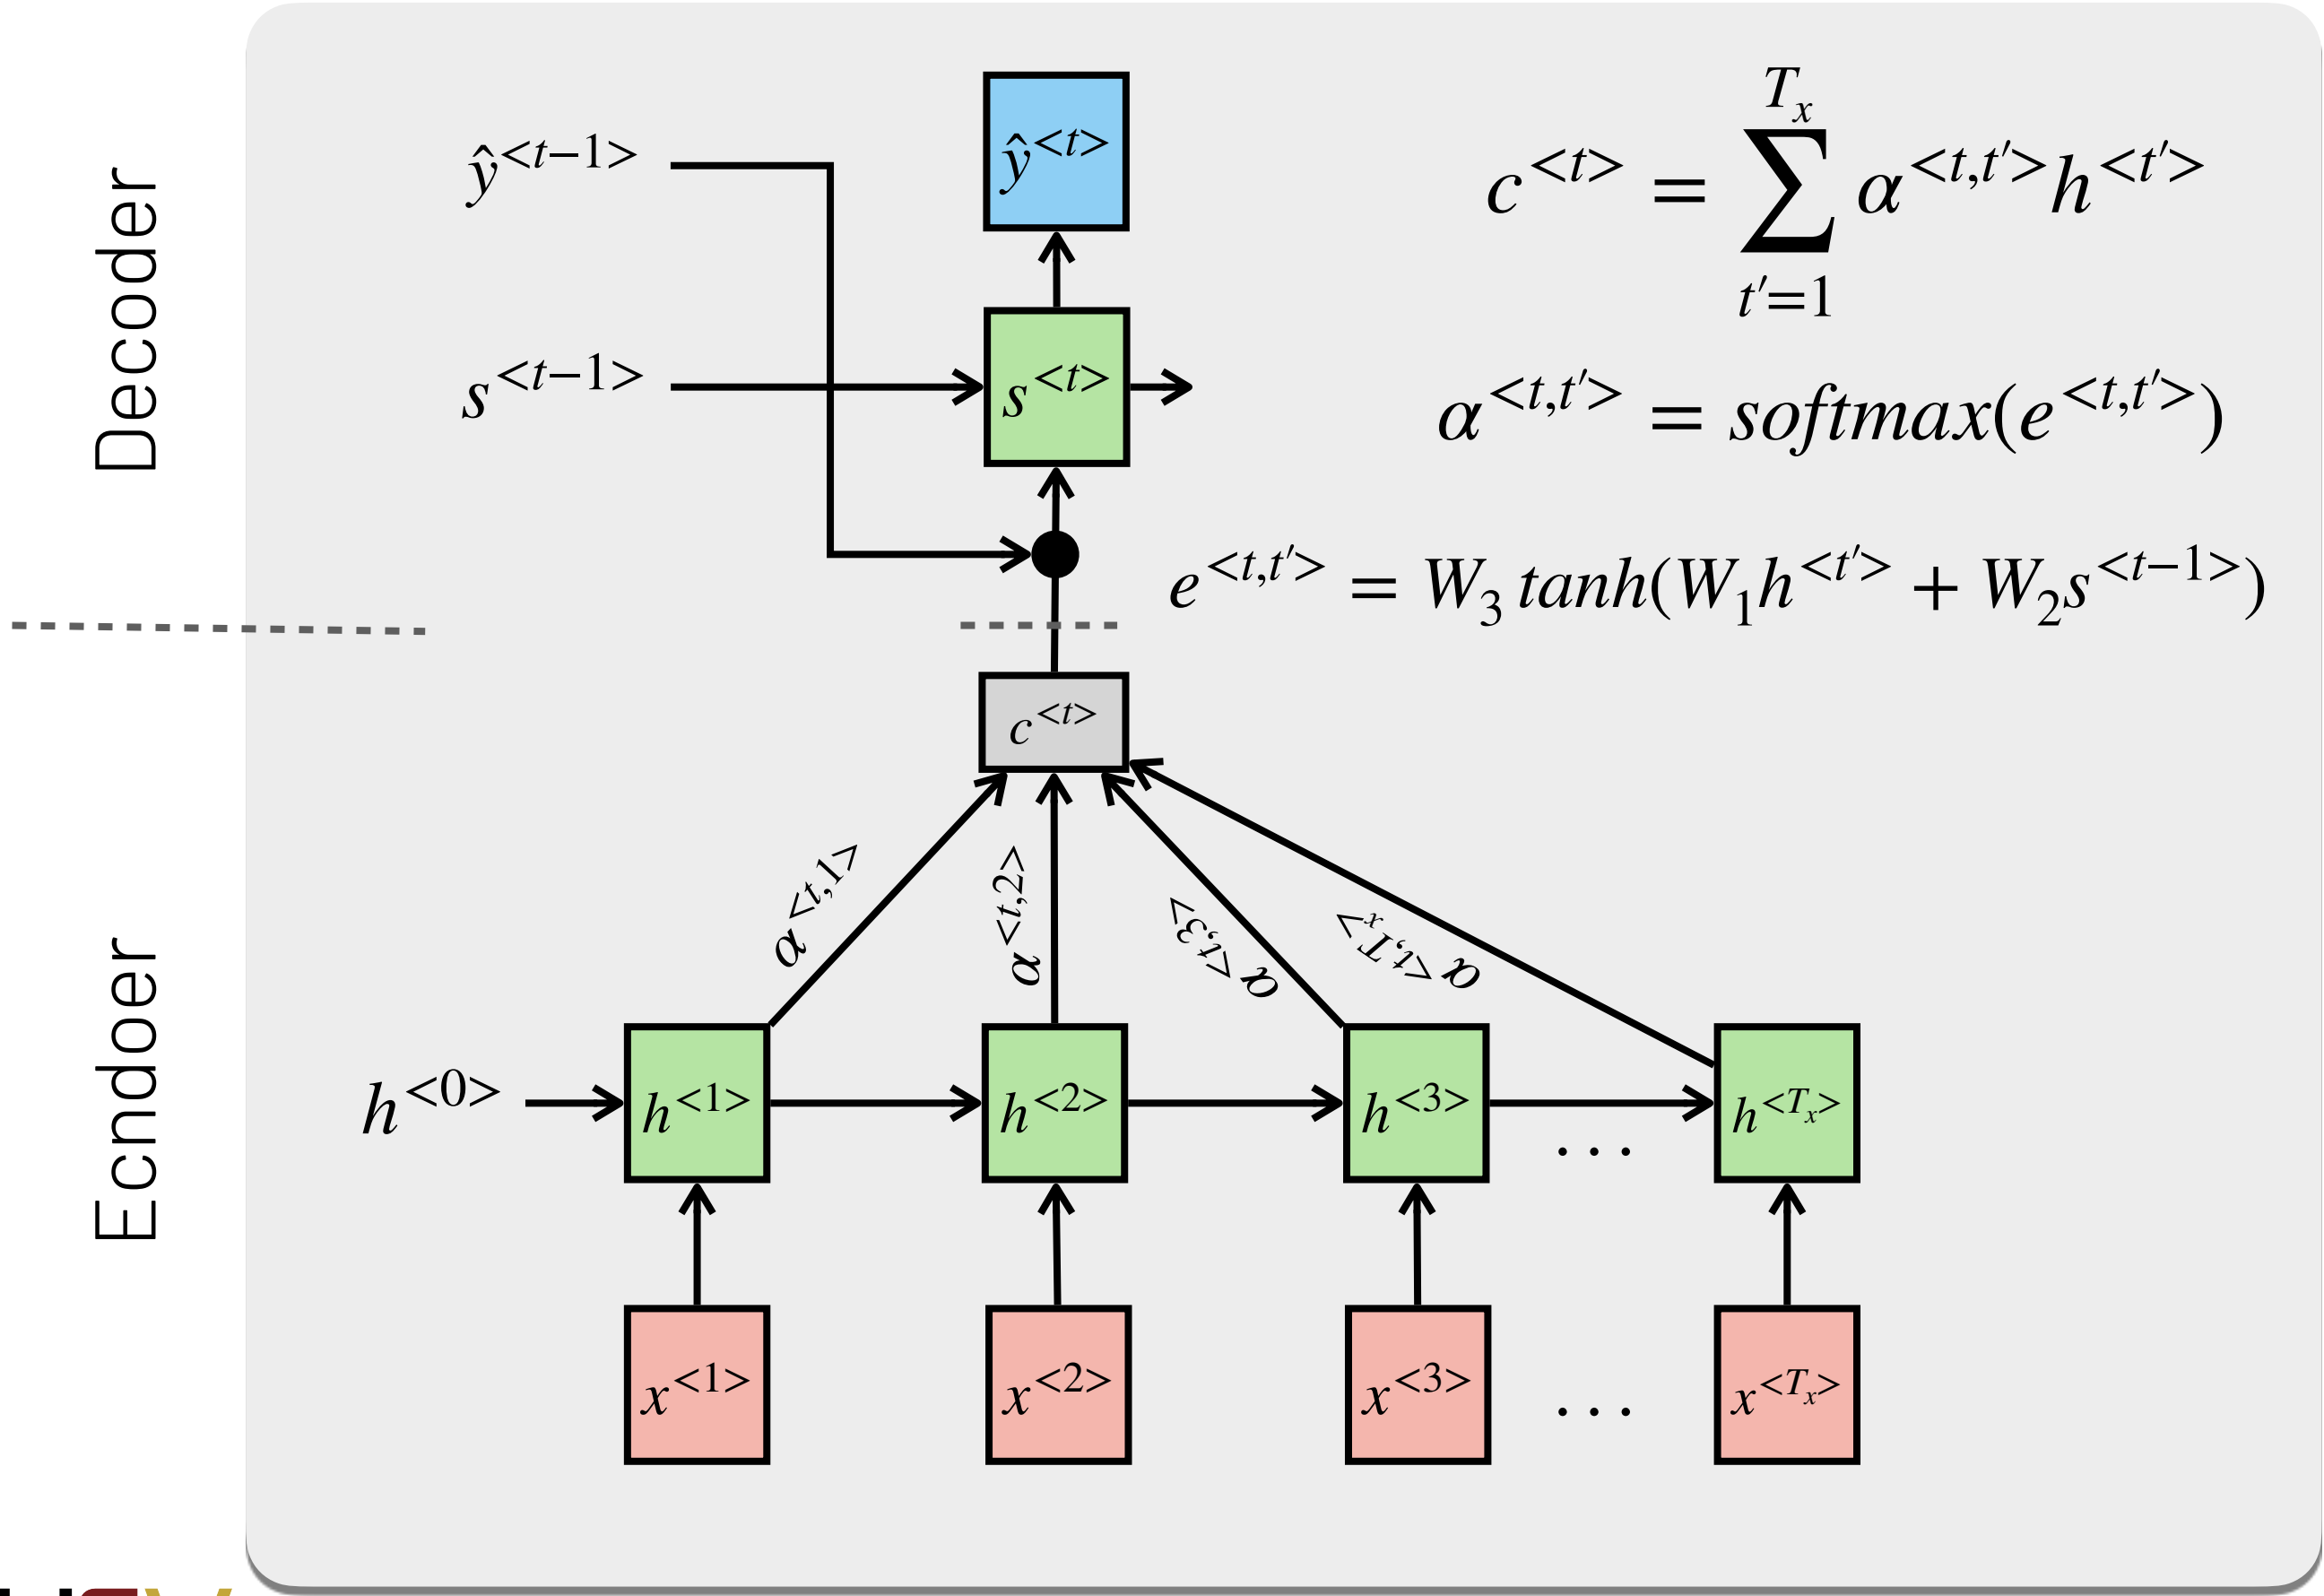
Initially proposed for machine translation, but proved to be very effective in many other problems in:

- ▶ Natural Language Processing
- ▶ Computer Vision
- ▶ Reinforcement Learning
- ▶ ...

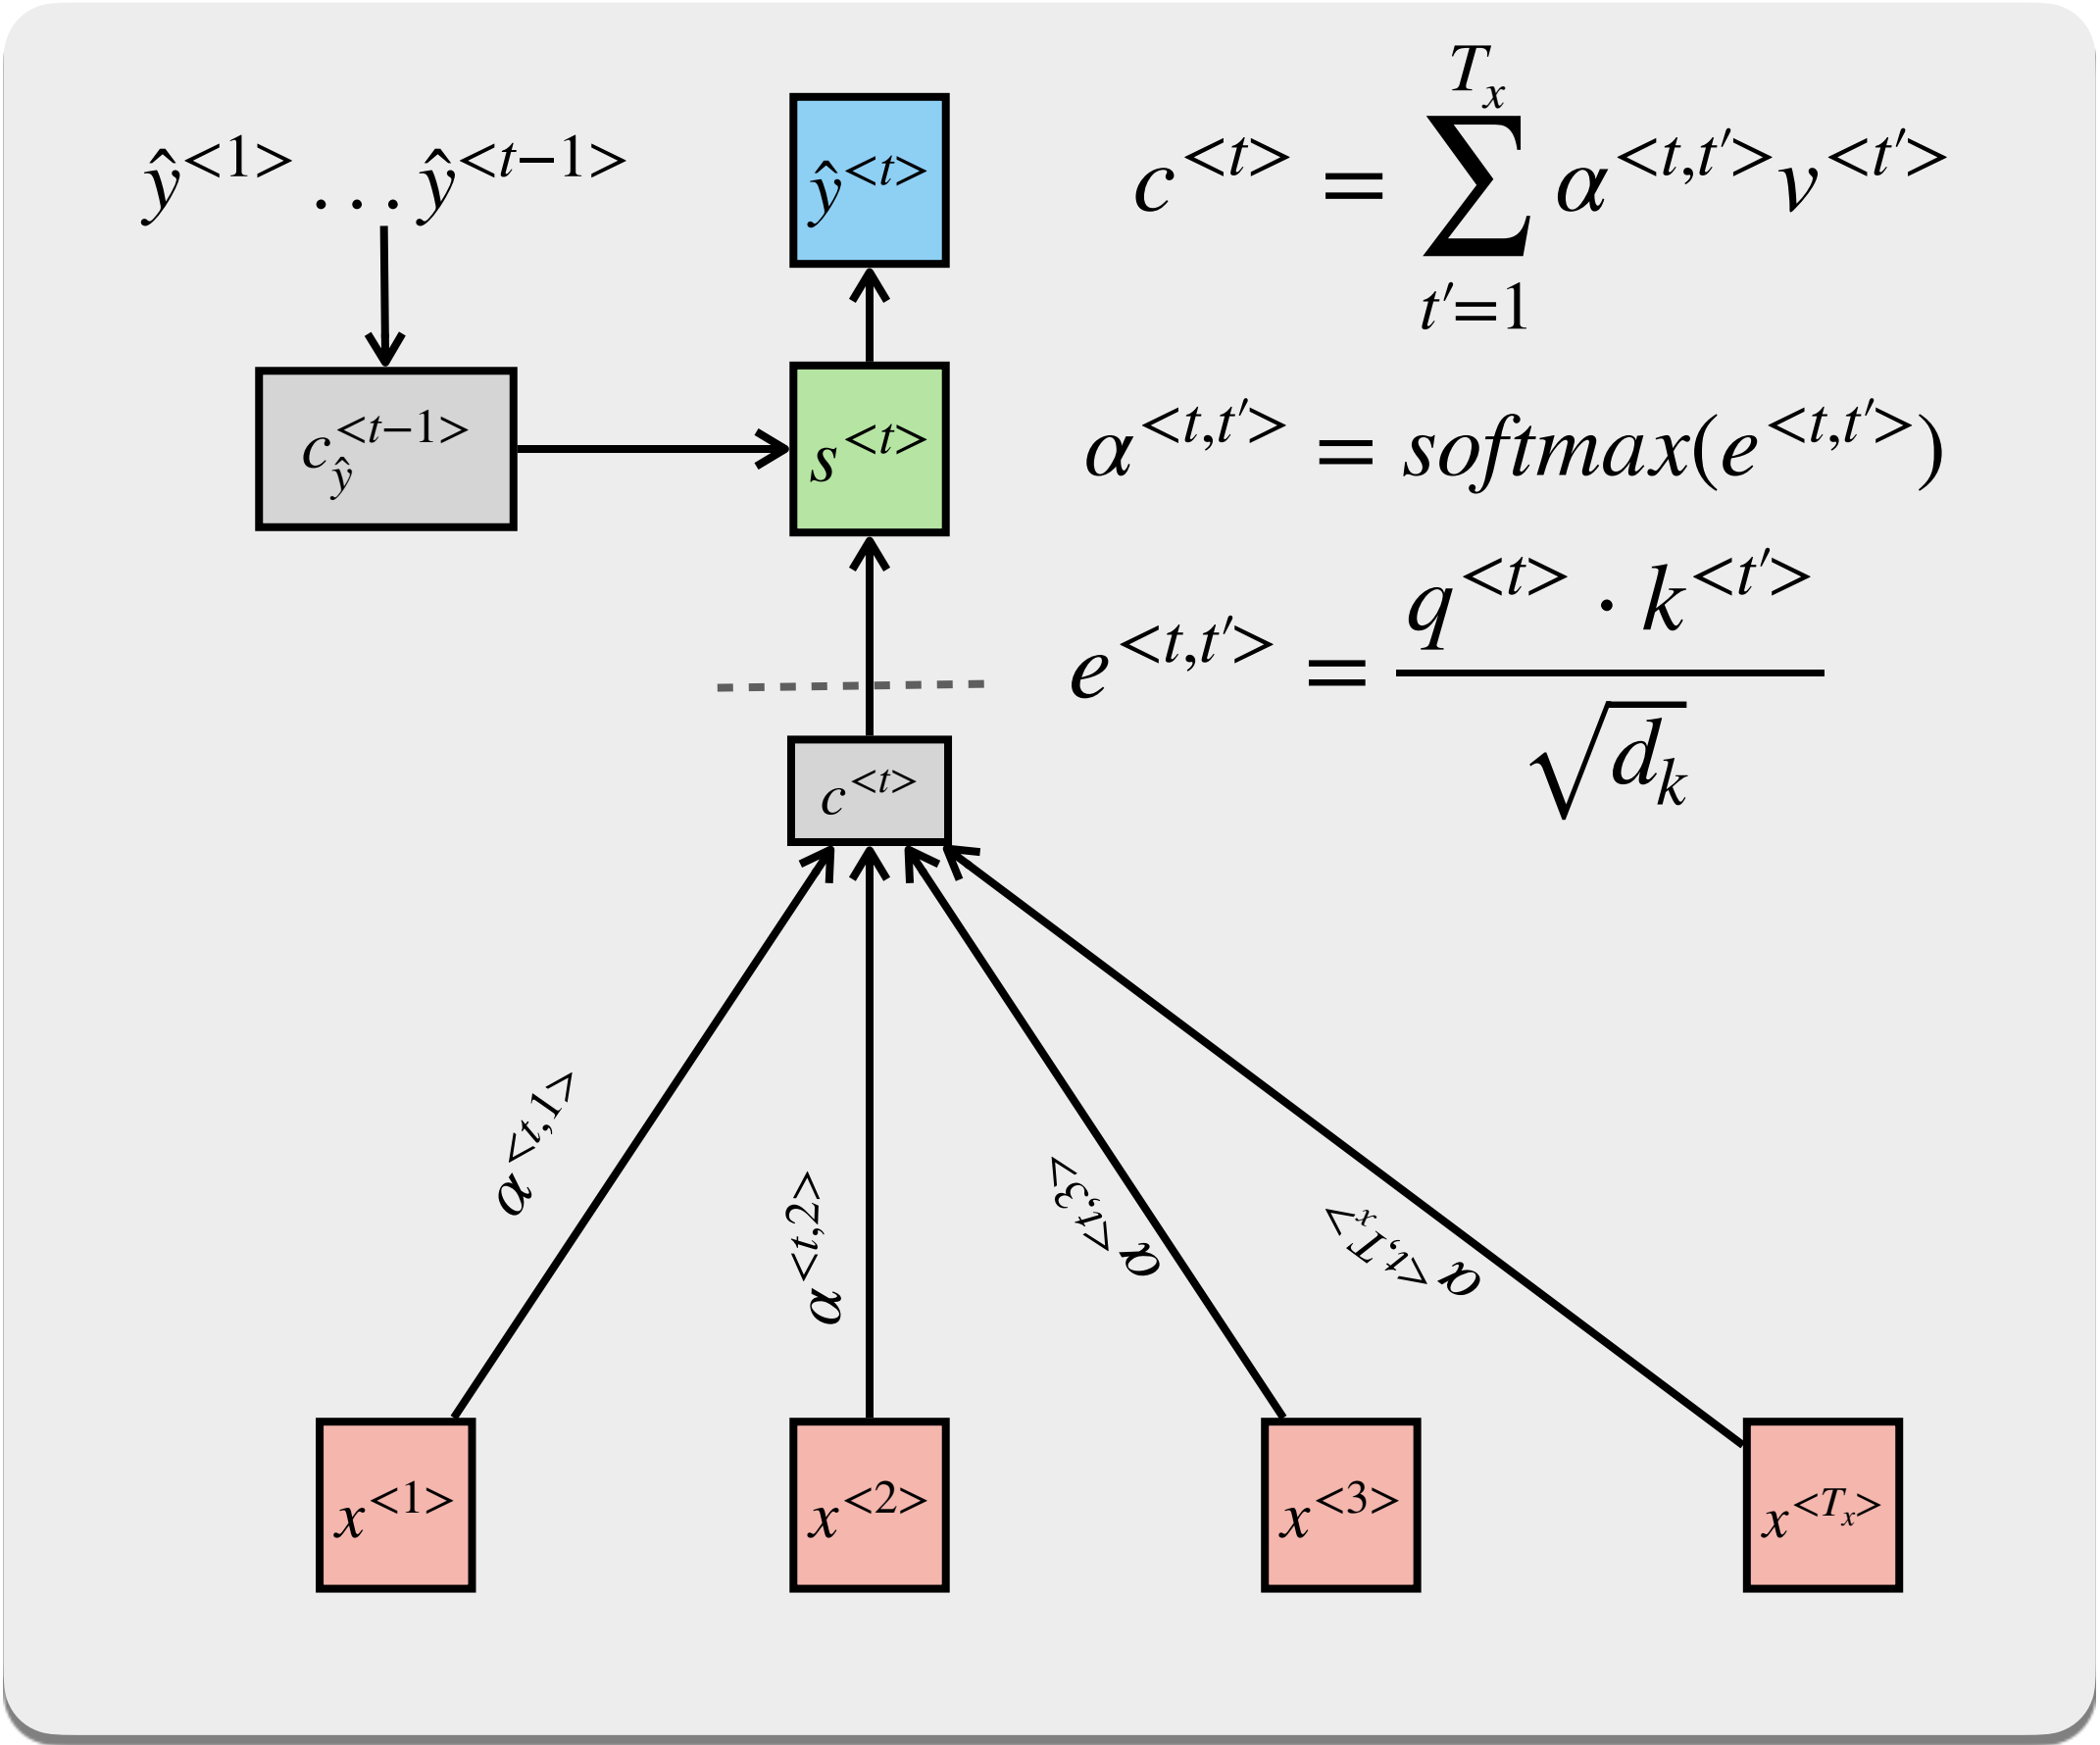


Attention in RNNs vs. Transformers

RNNs Badahnau Attention

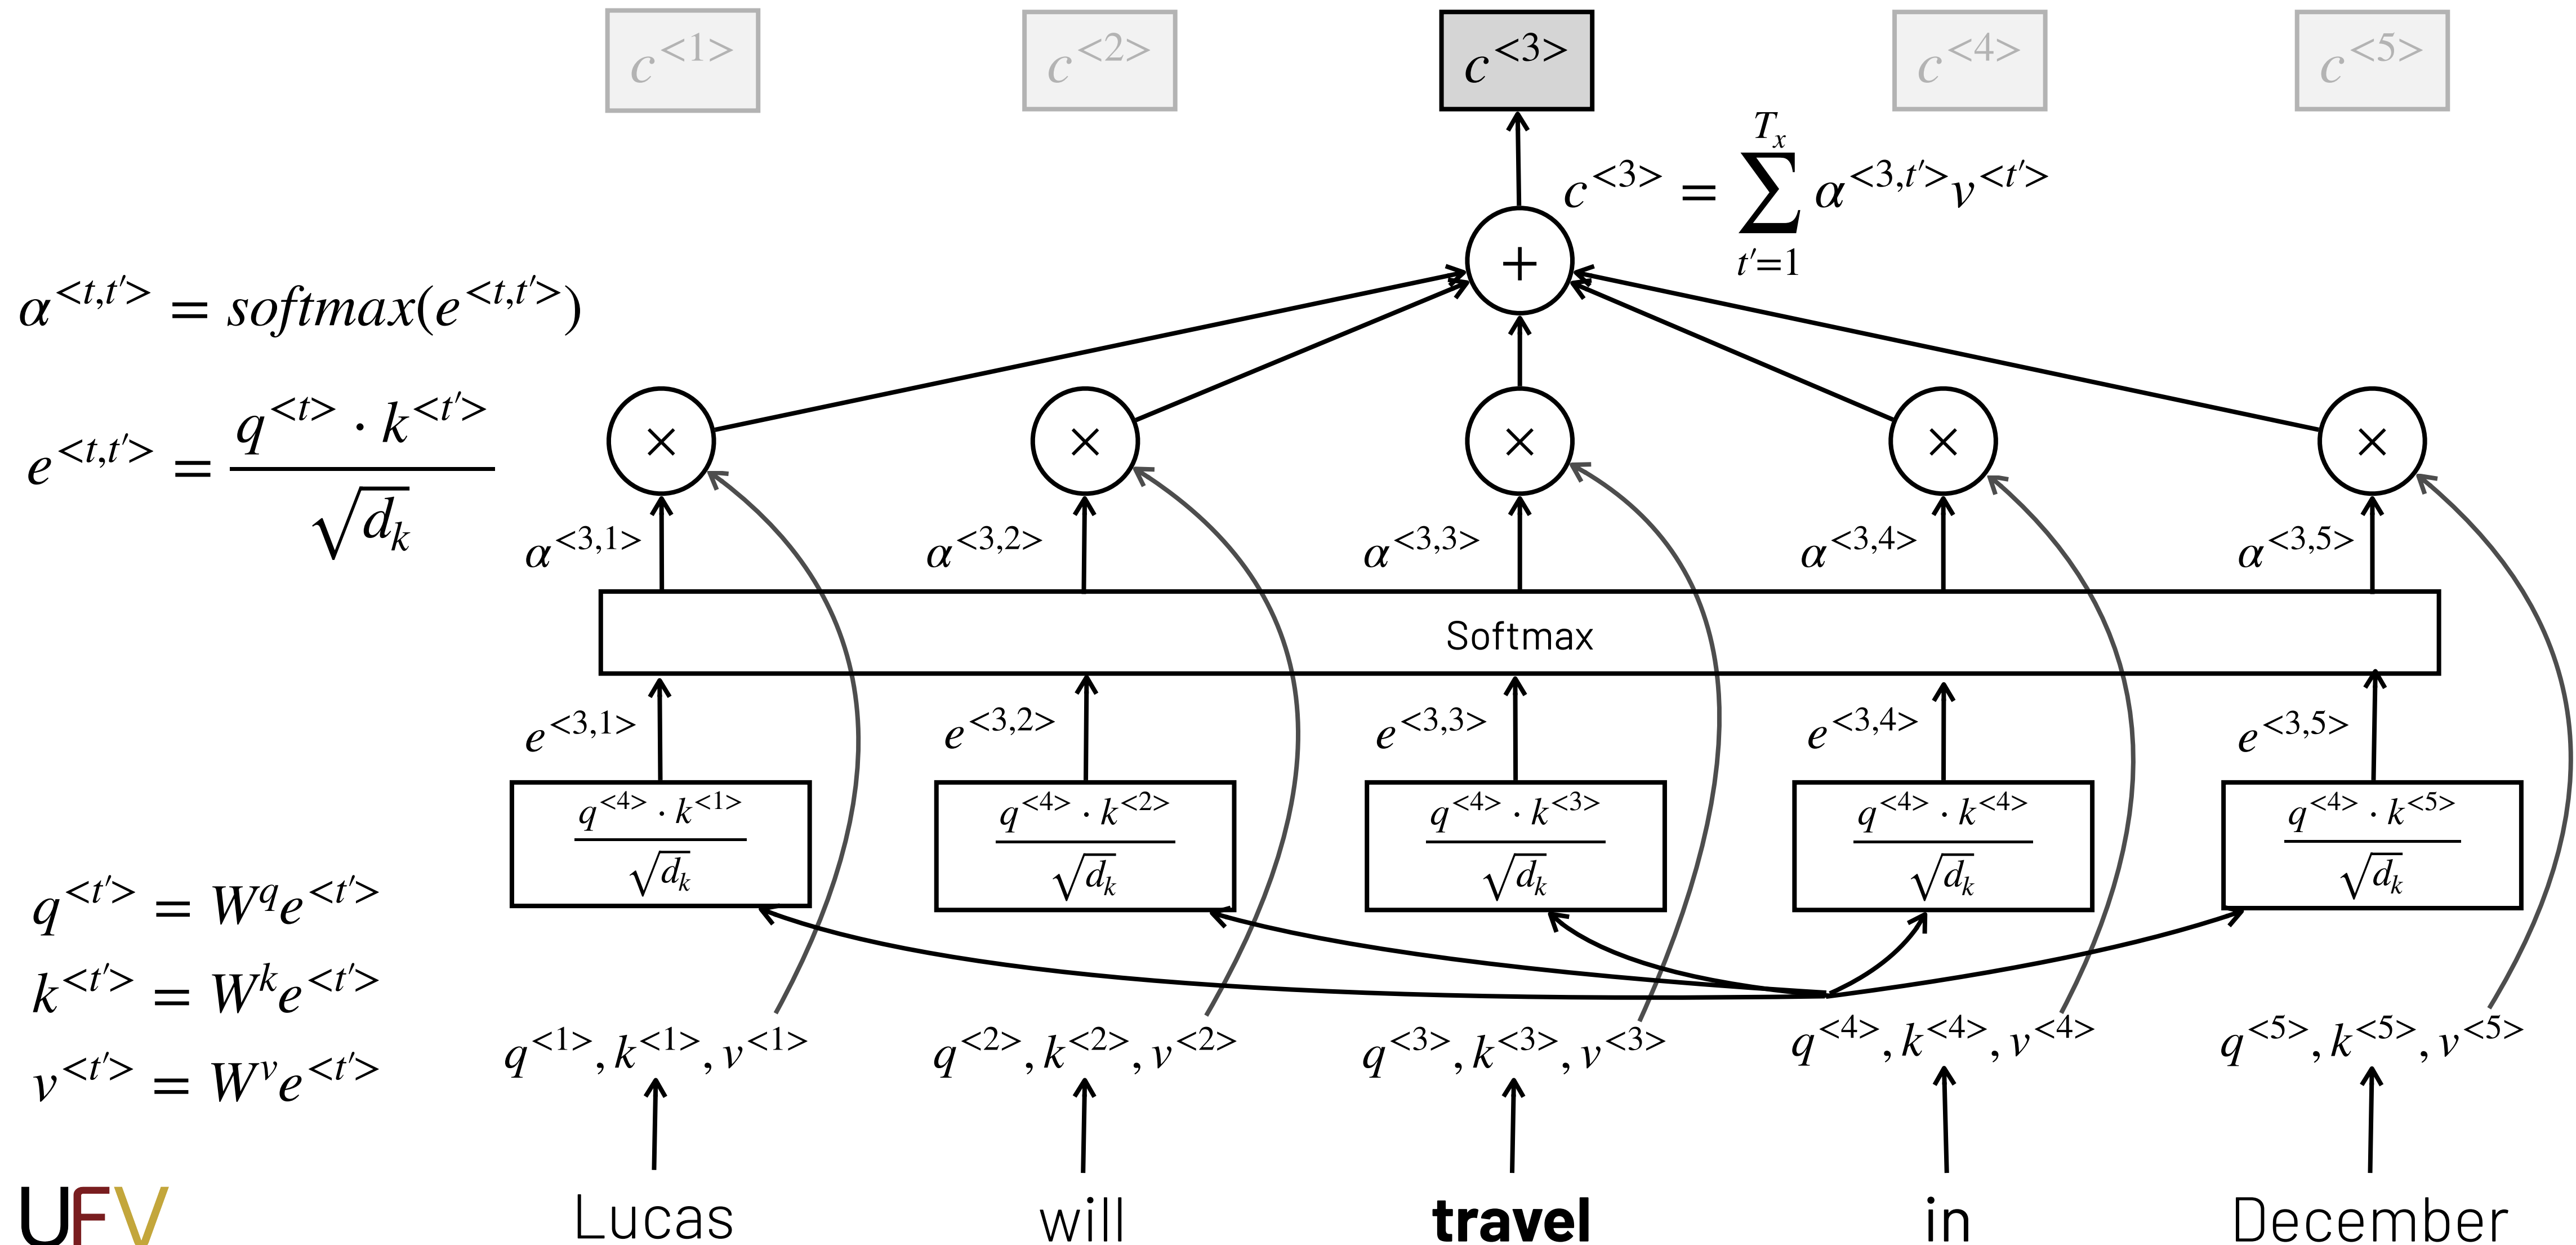


Transformers Scaled Dot-Product Attention



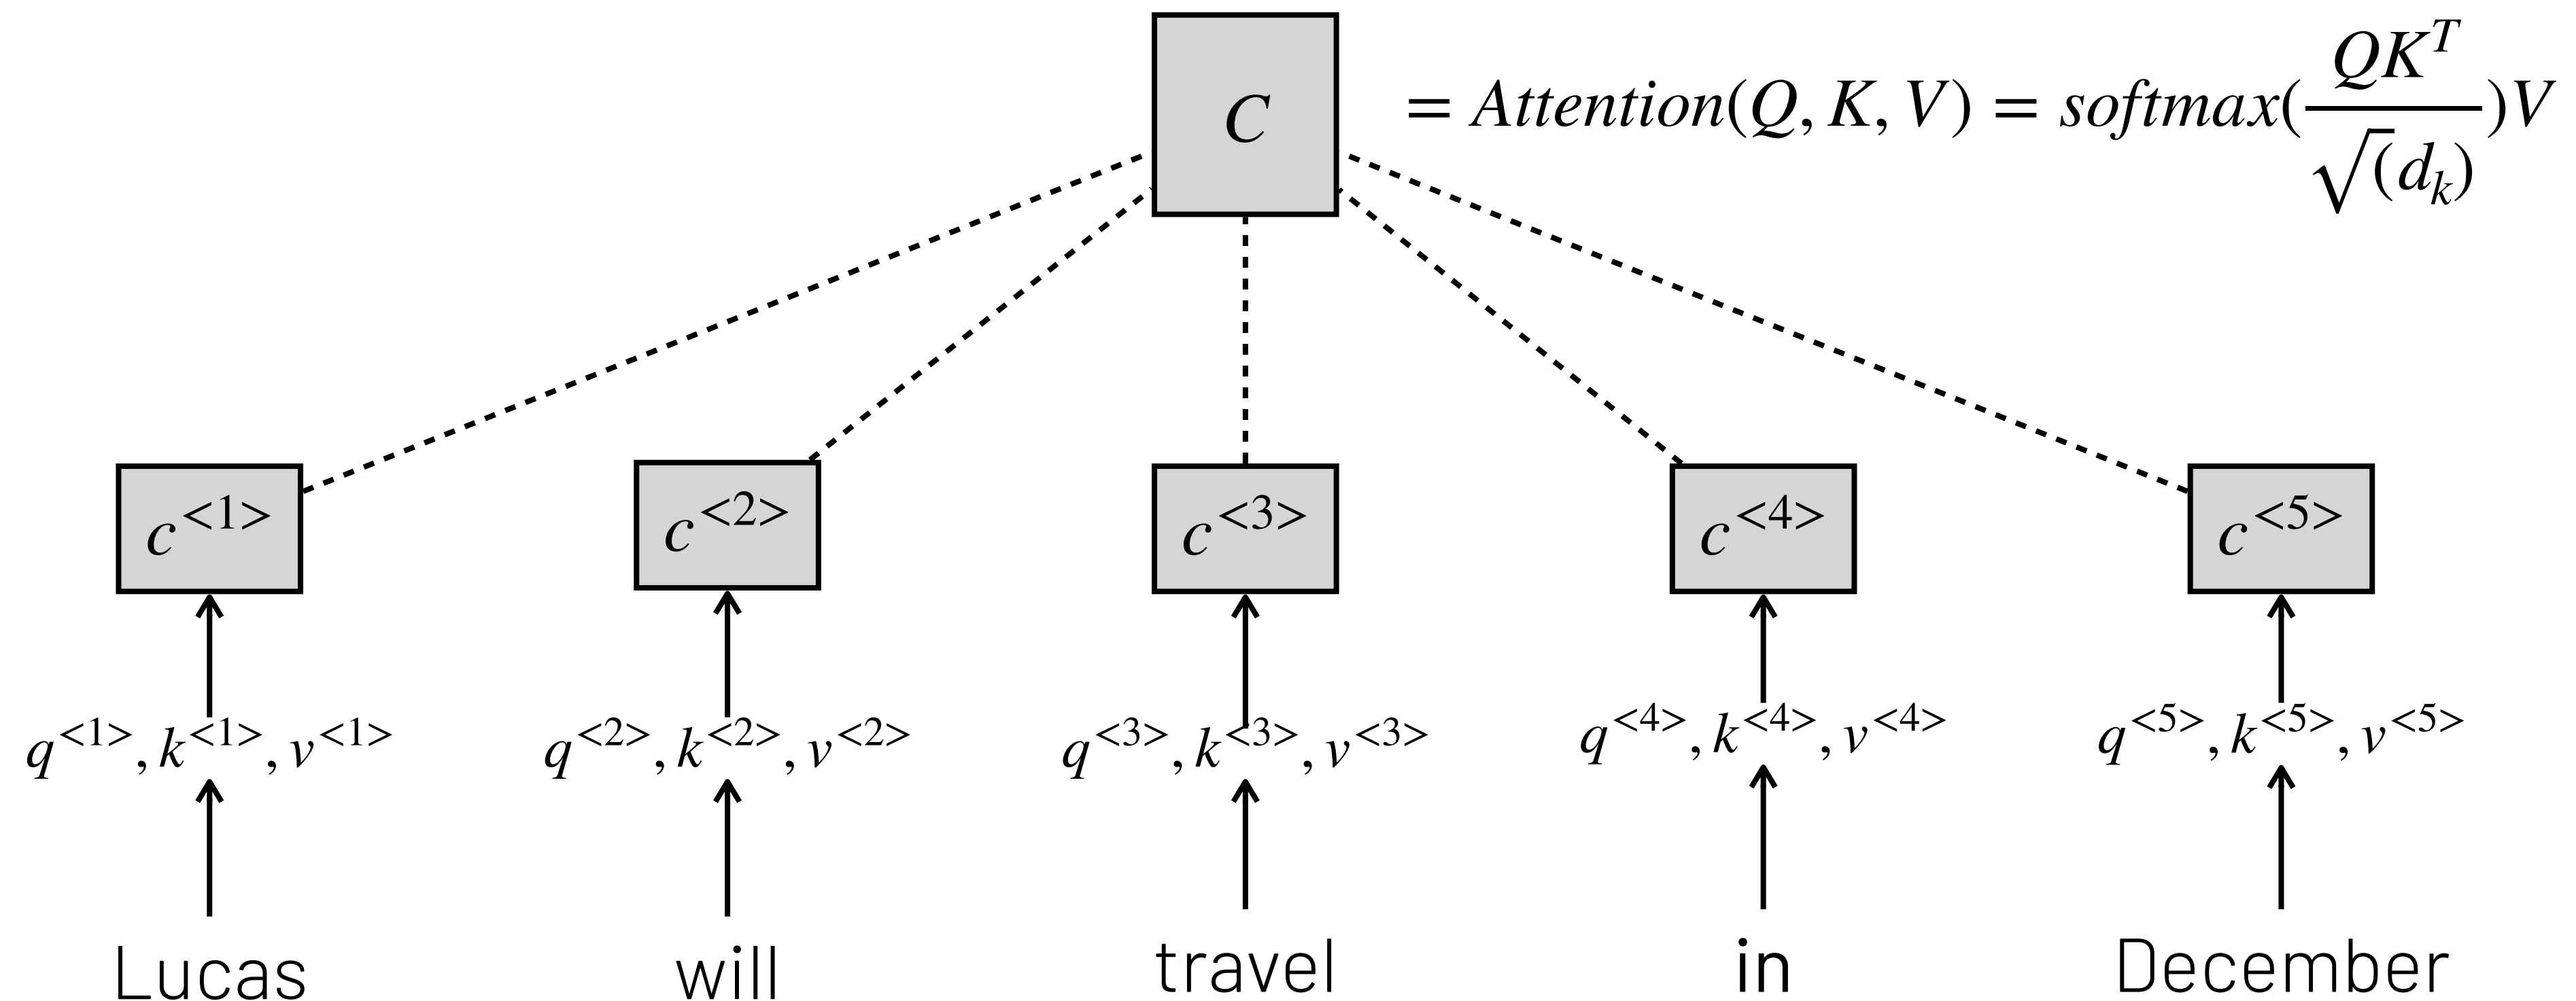
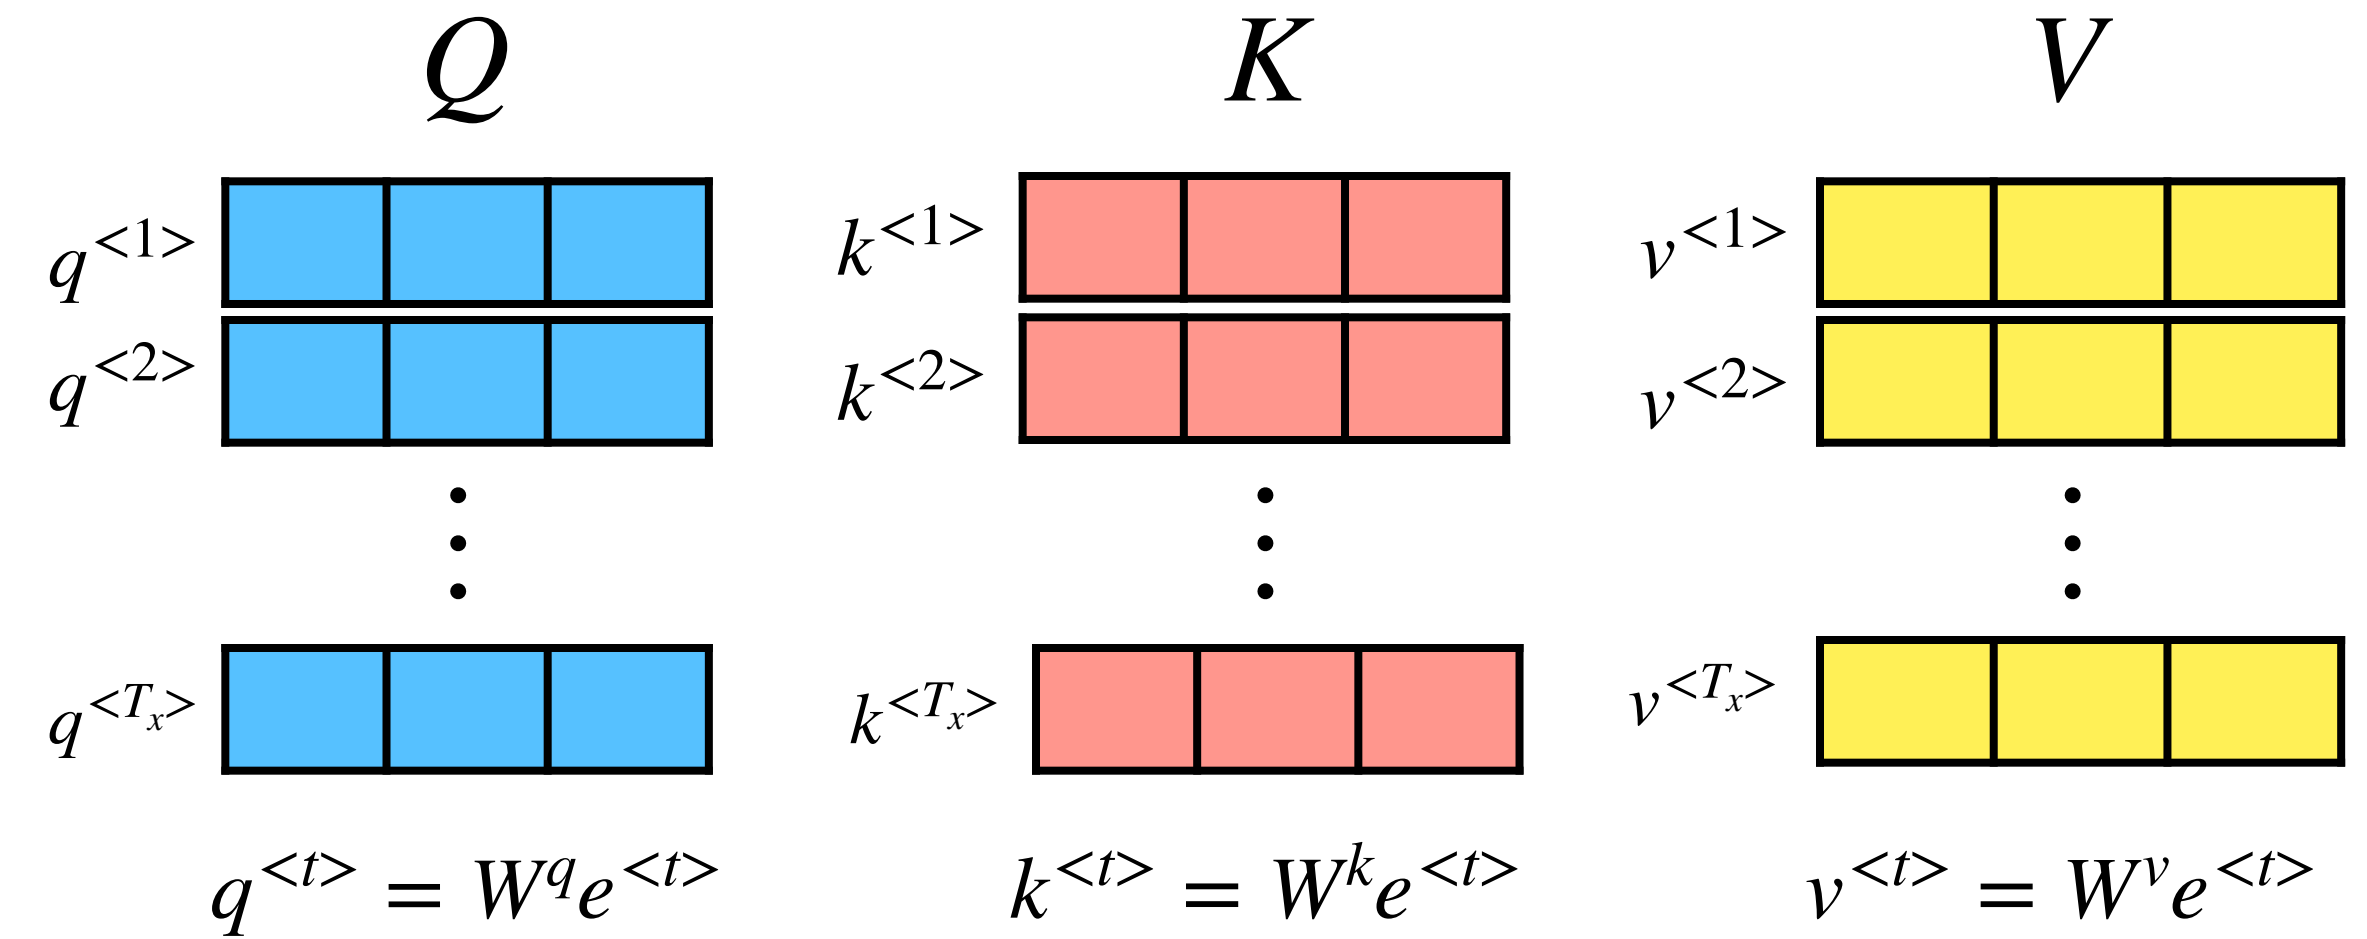
Self-Attention

The key idea behind the Transformer is the **self-attention mechanism**, which learns a context vector $c^{<t>}$ for each input element $x^{<t>}$ based on the input sequence x itself.



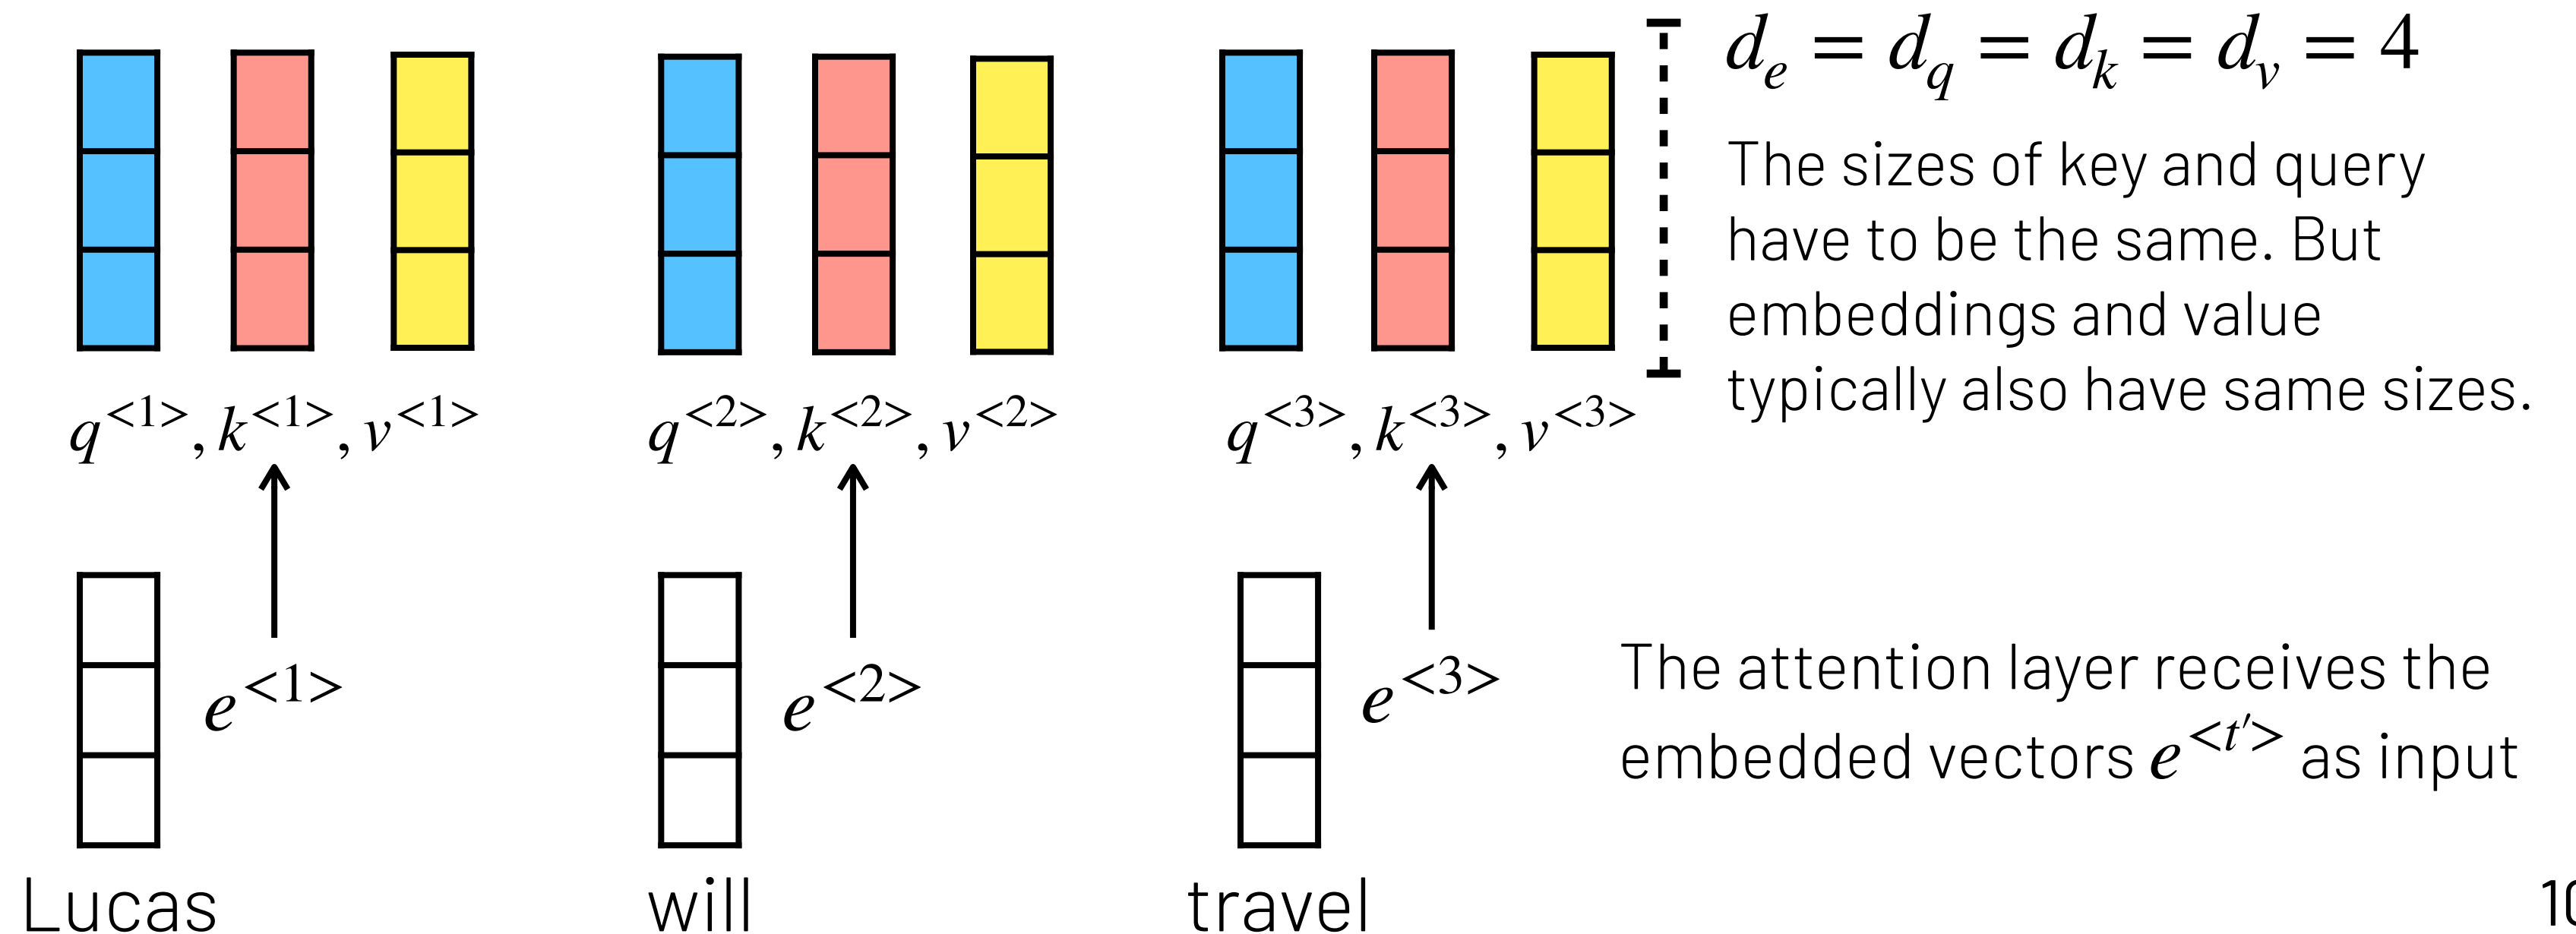
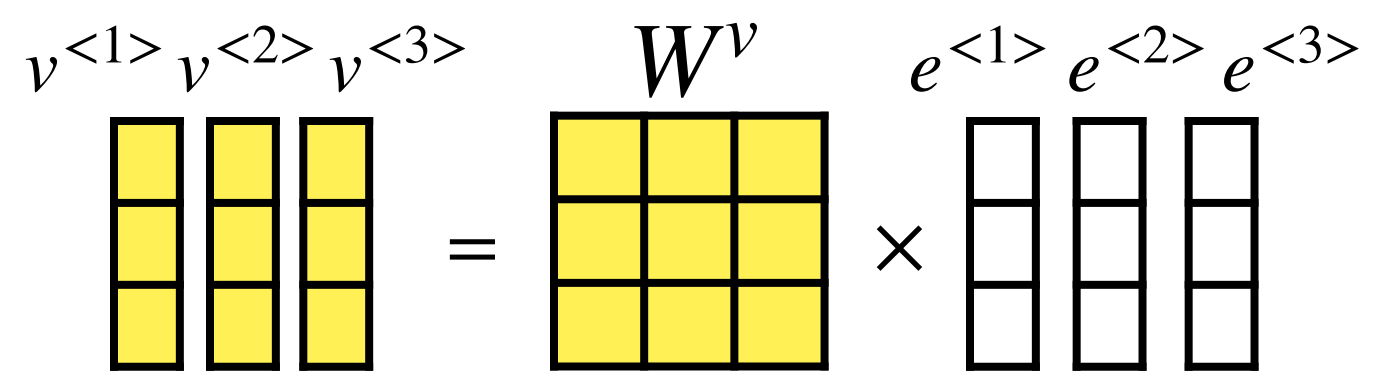
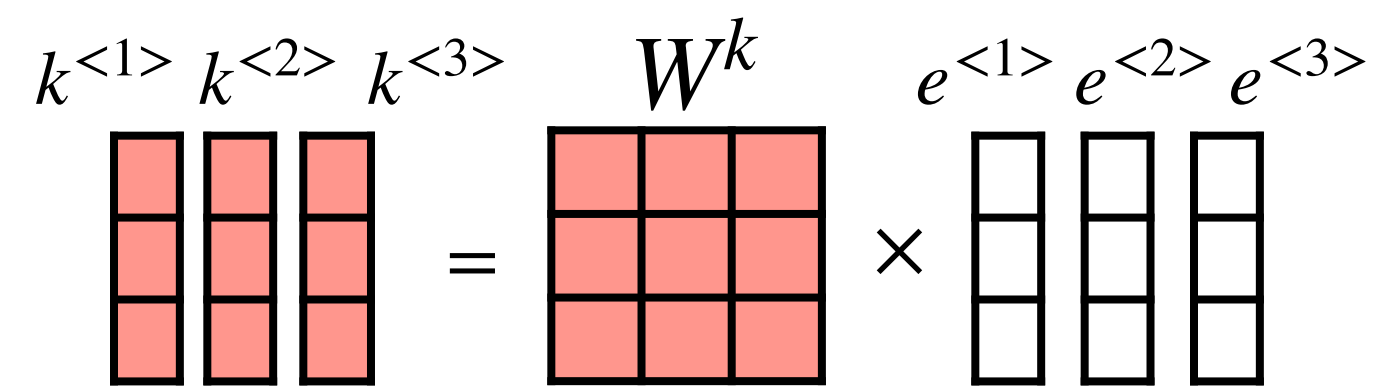
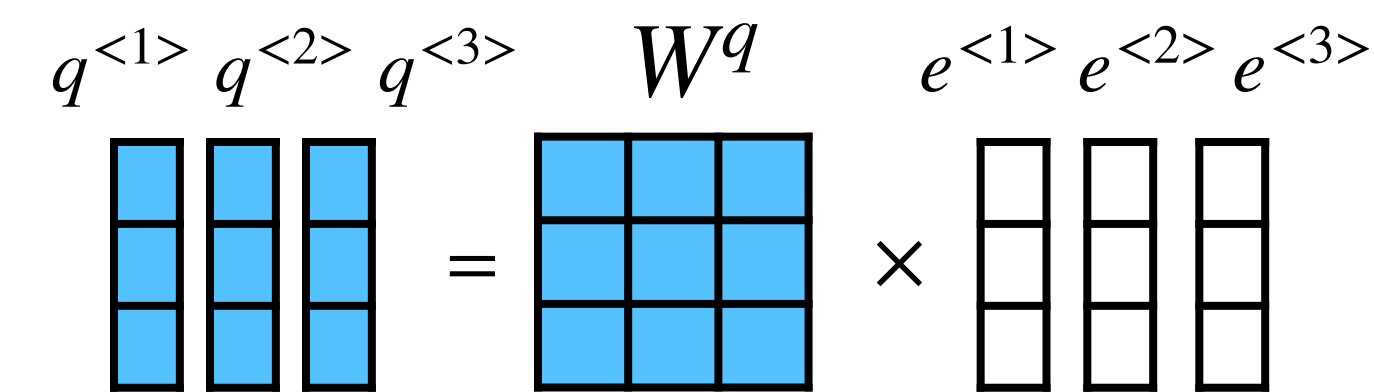
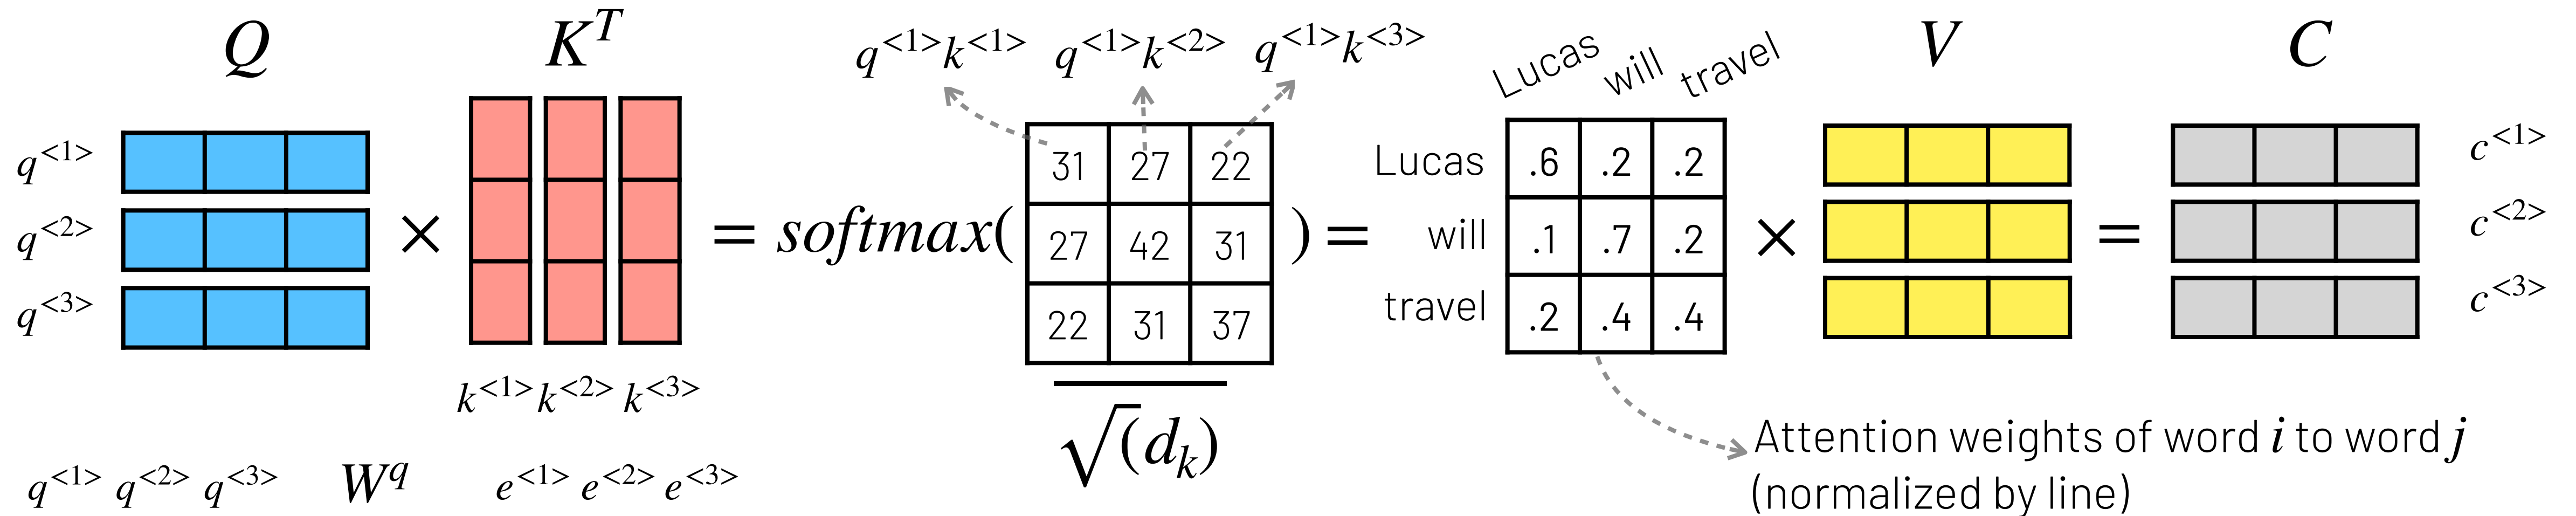
Self-Attention

The contextual representation $C = \{c^{<1>}, \dots, c^{<T_x>}\}$ of the entire input sequence $x = \{x^{<1>}, \dots, x^{<T_x>}\}$ can be computed in a vectorized way combining vectors $q^{<t>}, k^{<t>}, v^{<t>}$ in matrices Q, K e V



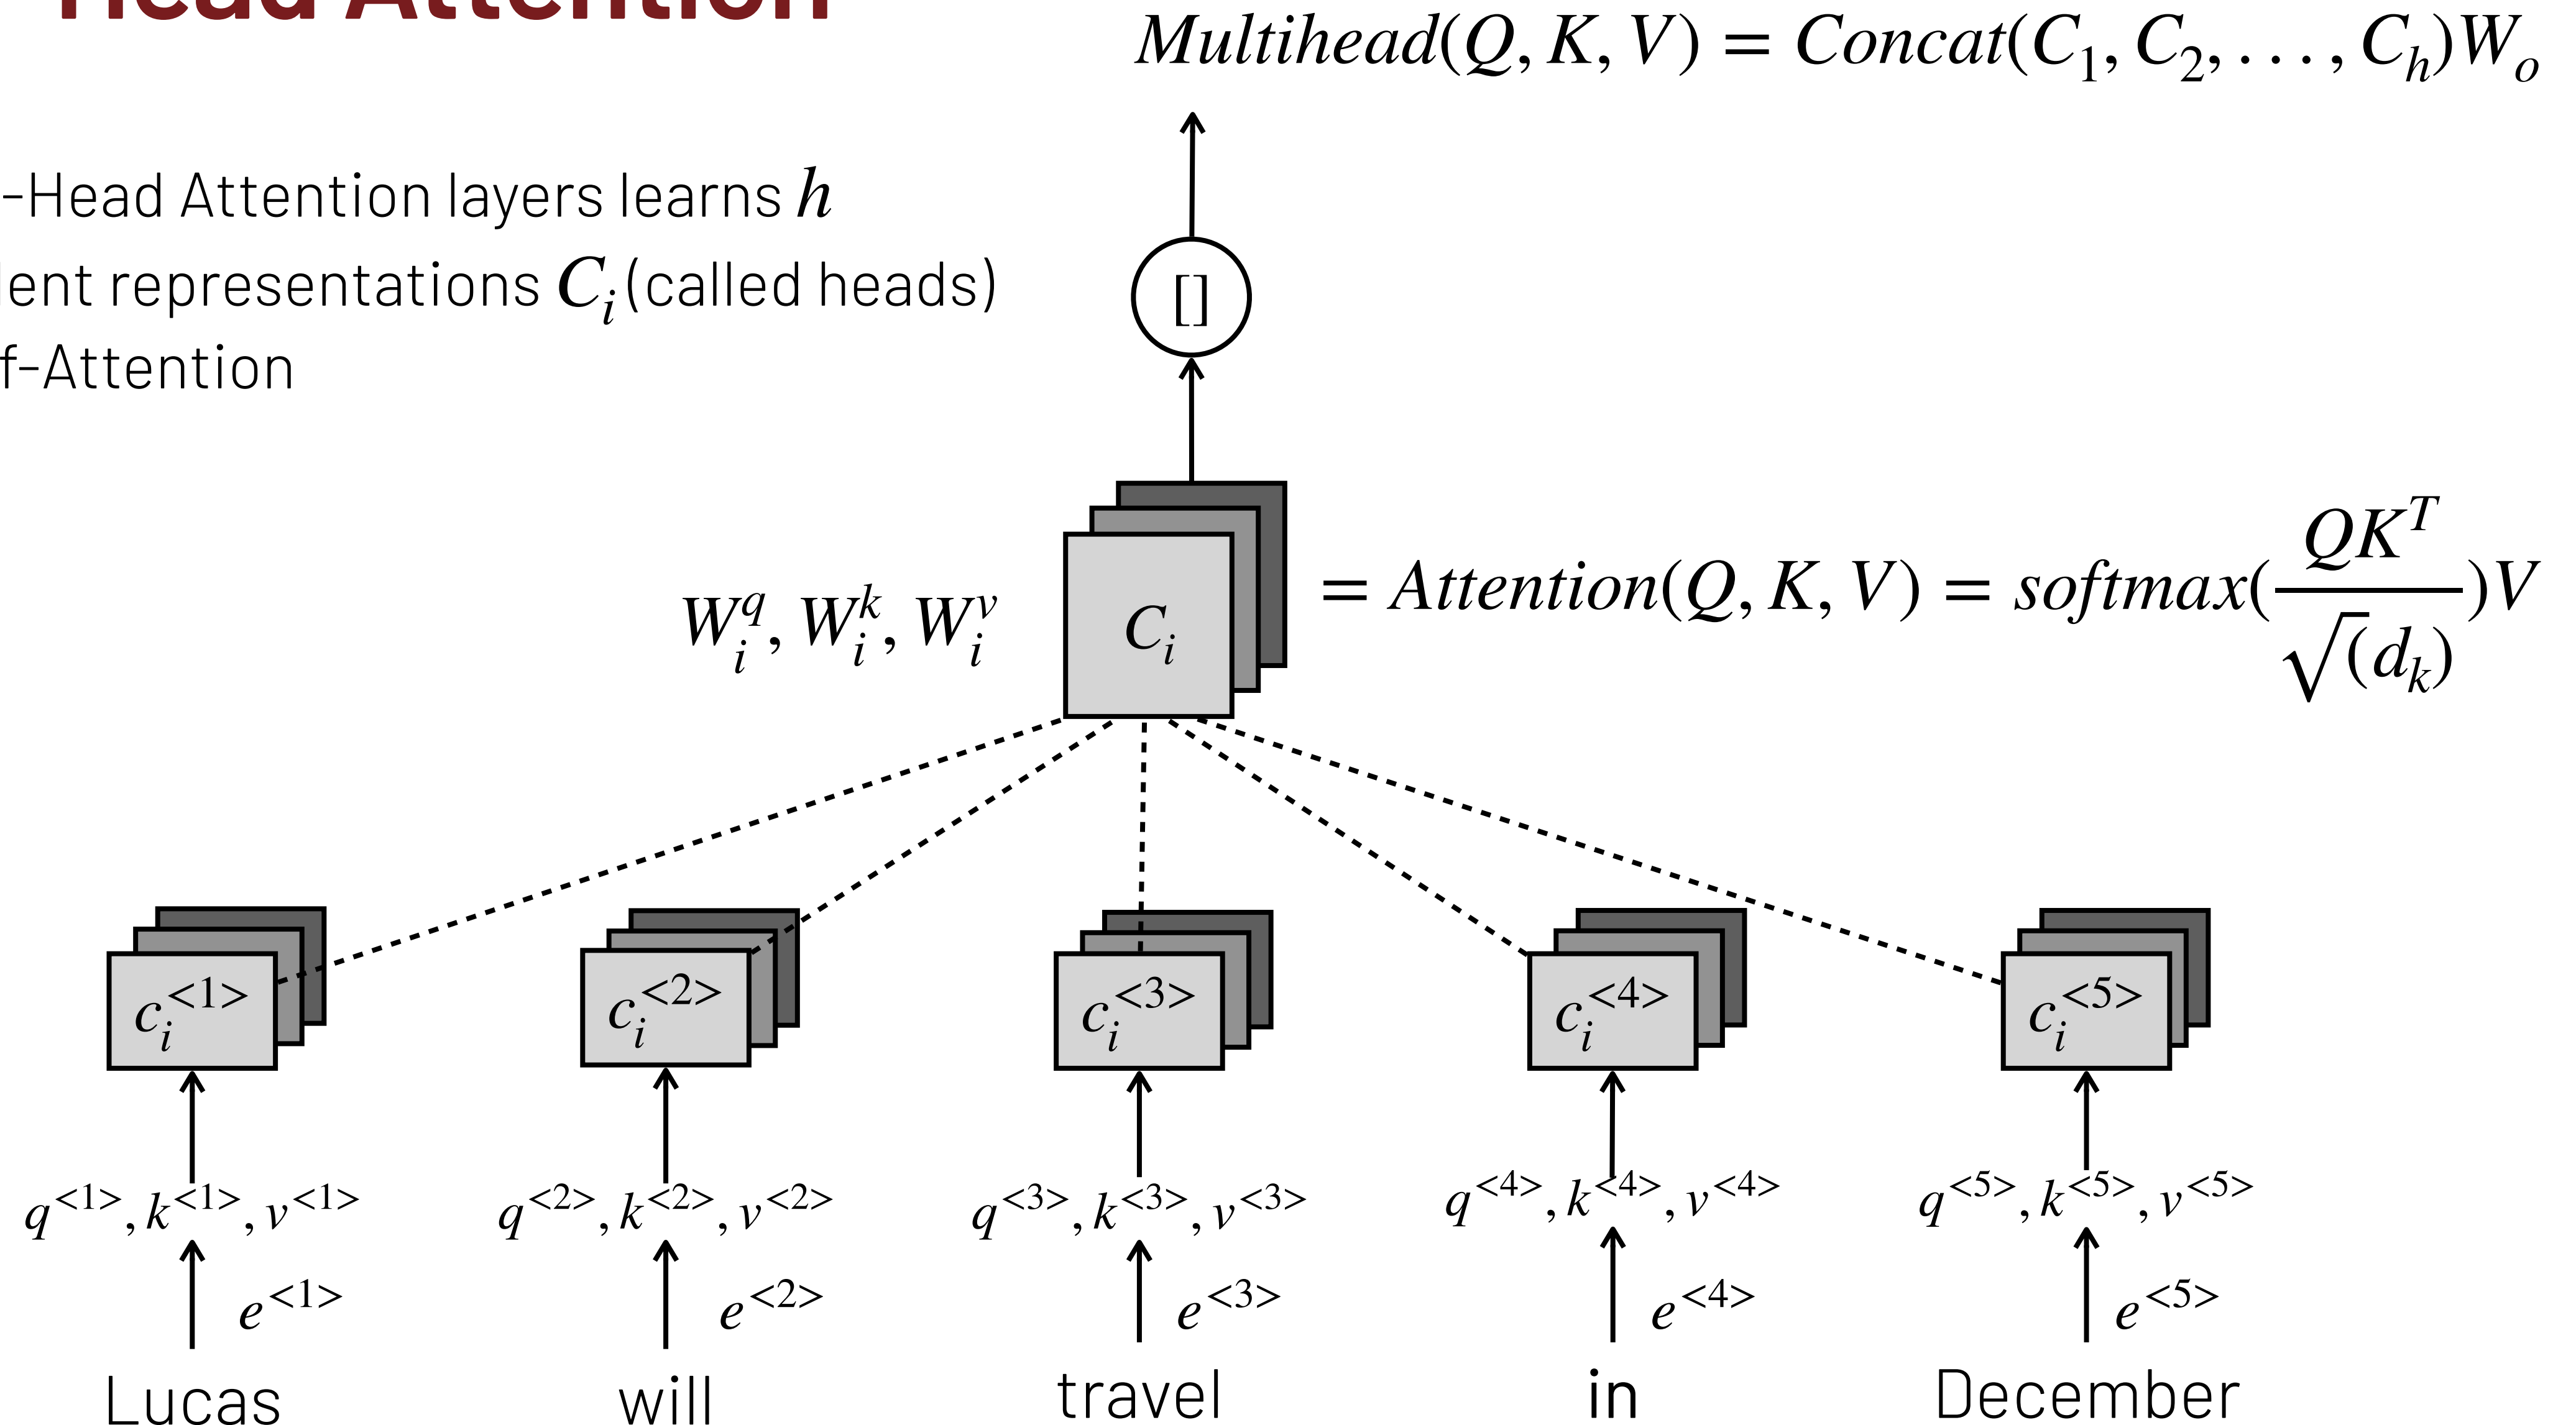
Self-Attention

$$C = \text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

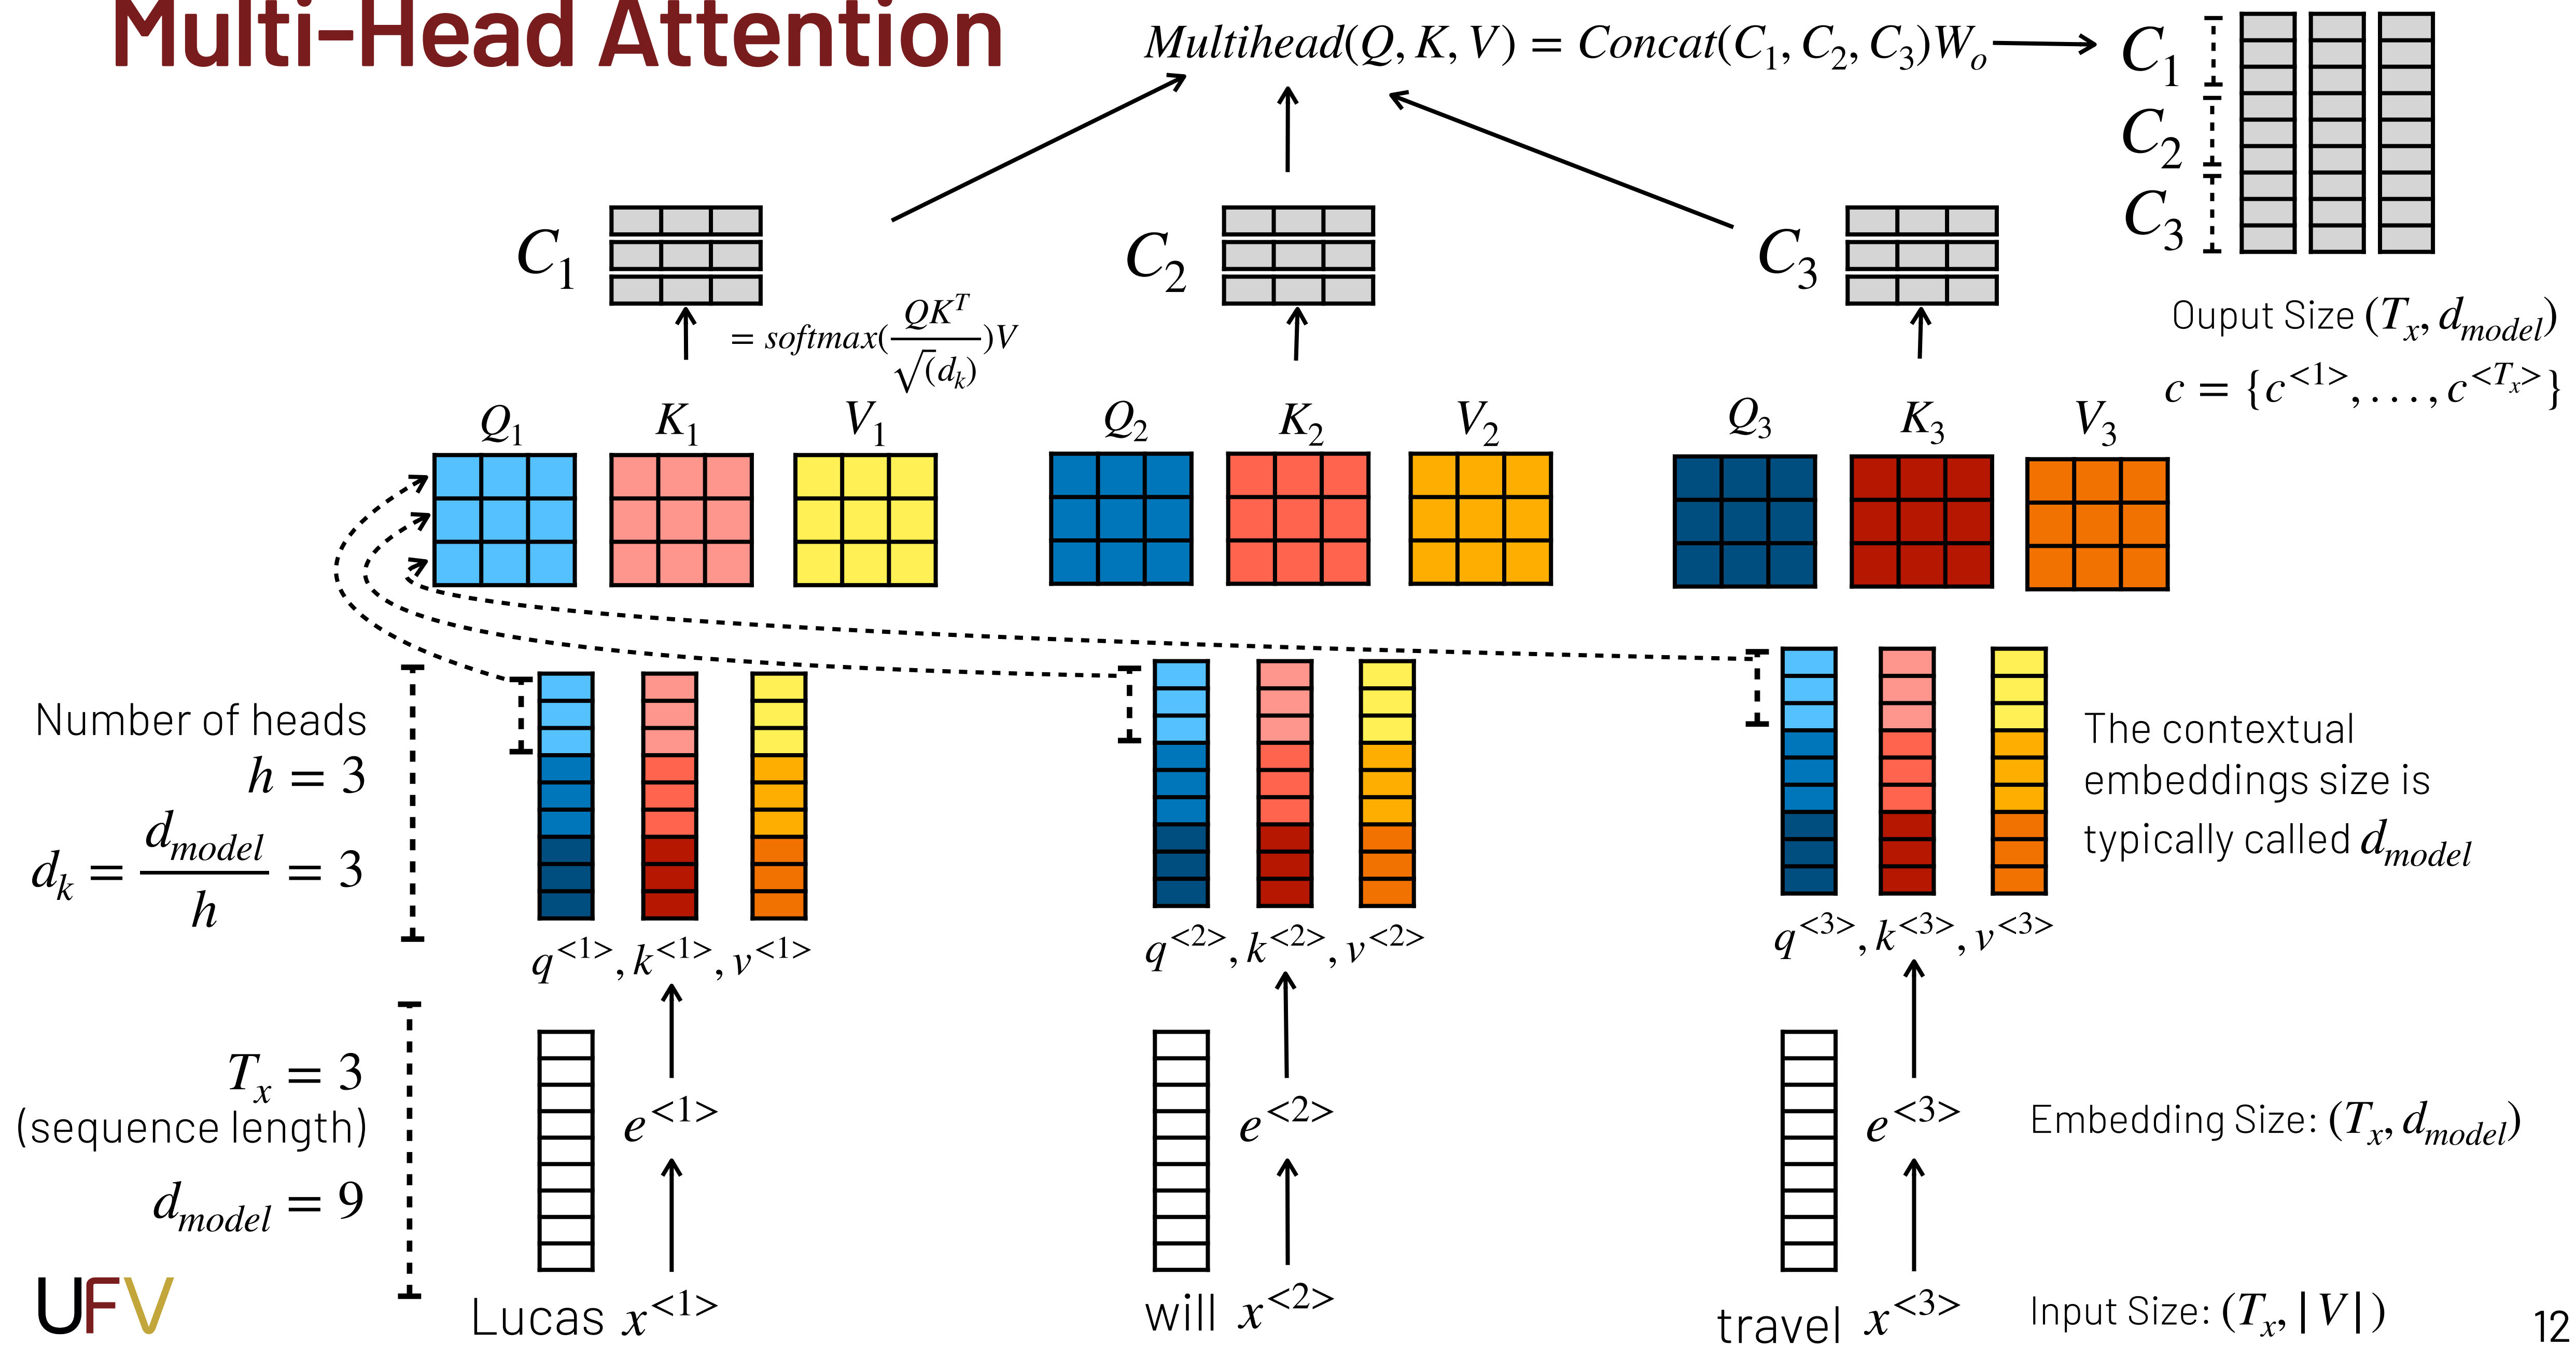


Multi-Head Attention

The Multi-Head Attention layers learns h independent representations C_i (called heads) using Self-Attention



Multi-Head Attention



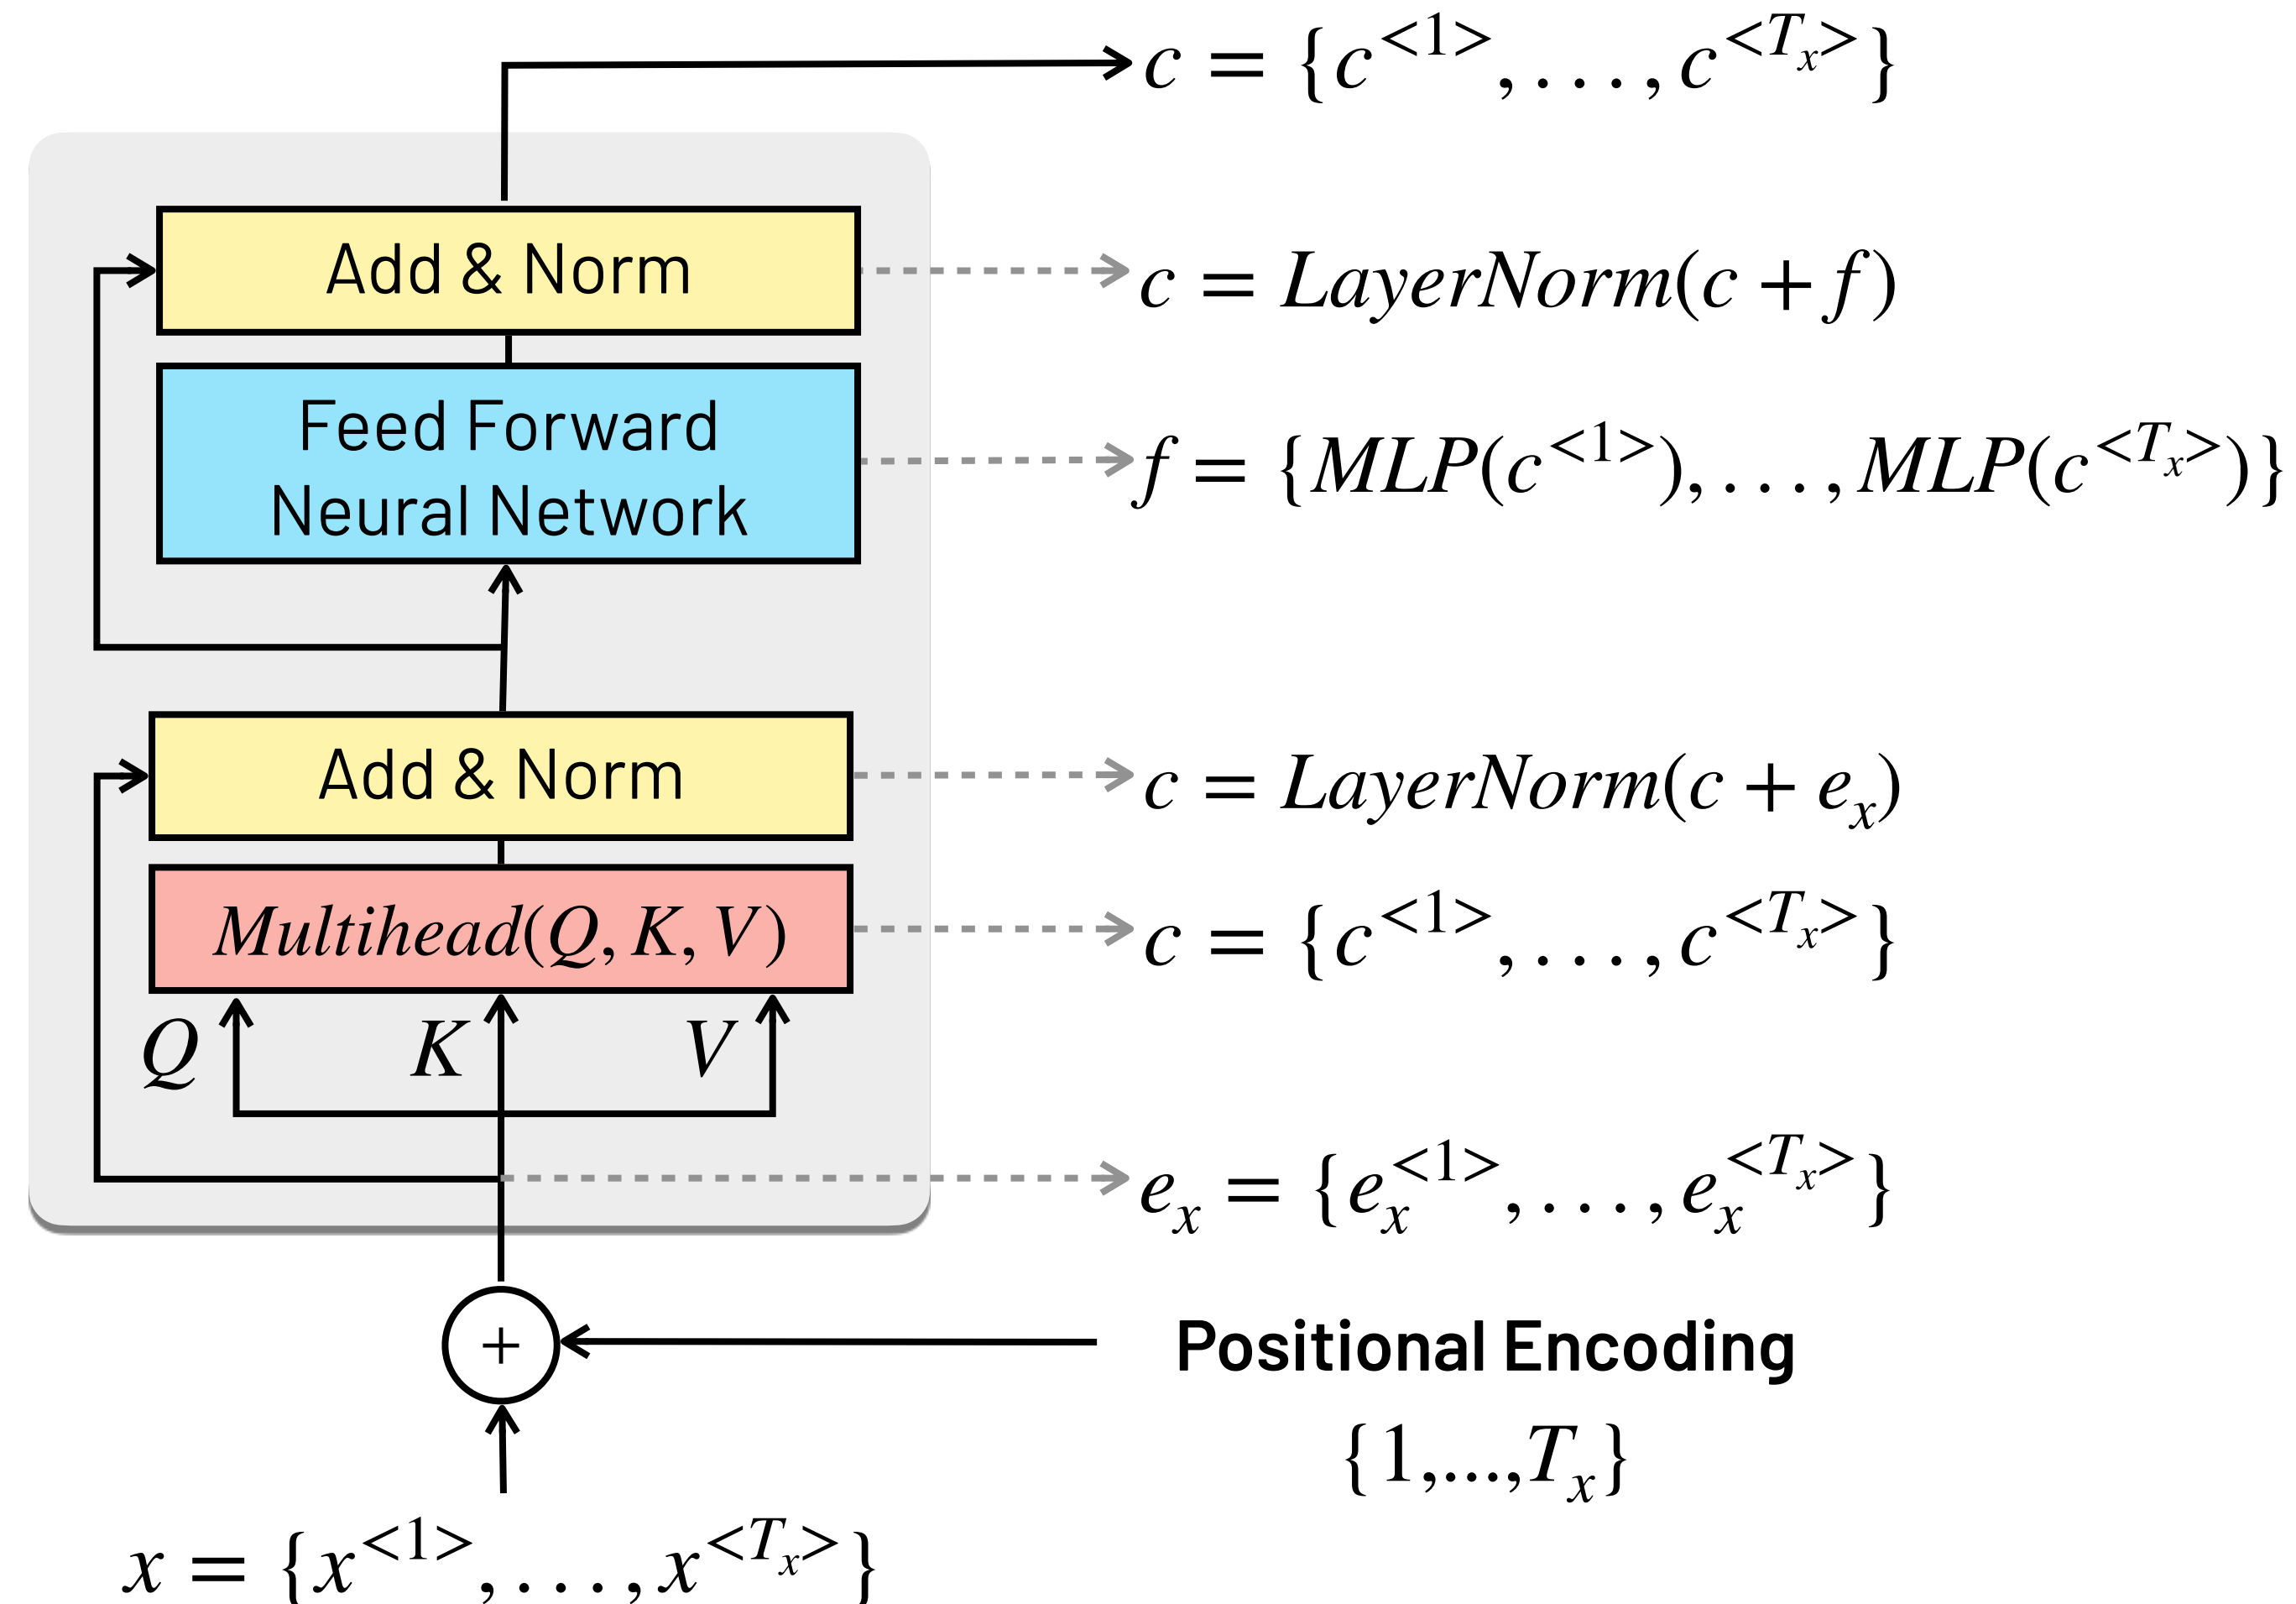
Encoder

Input: a sequence $x = \{x^{<1>}, \dots, x^{<T_x>}\}$

Output: a contextual representation $C = \{c^{<1>}, \dots, c^{<T_x>}\}$ of x

The encoder applies a **Multihead Layer** followed by a **Feed Forward Neural Network** (MLP).

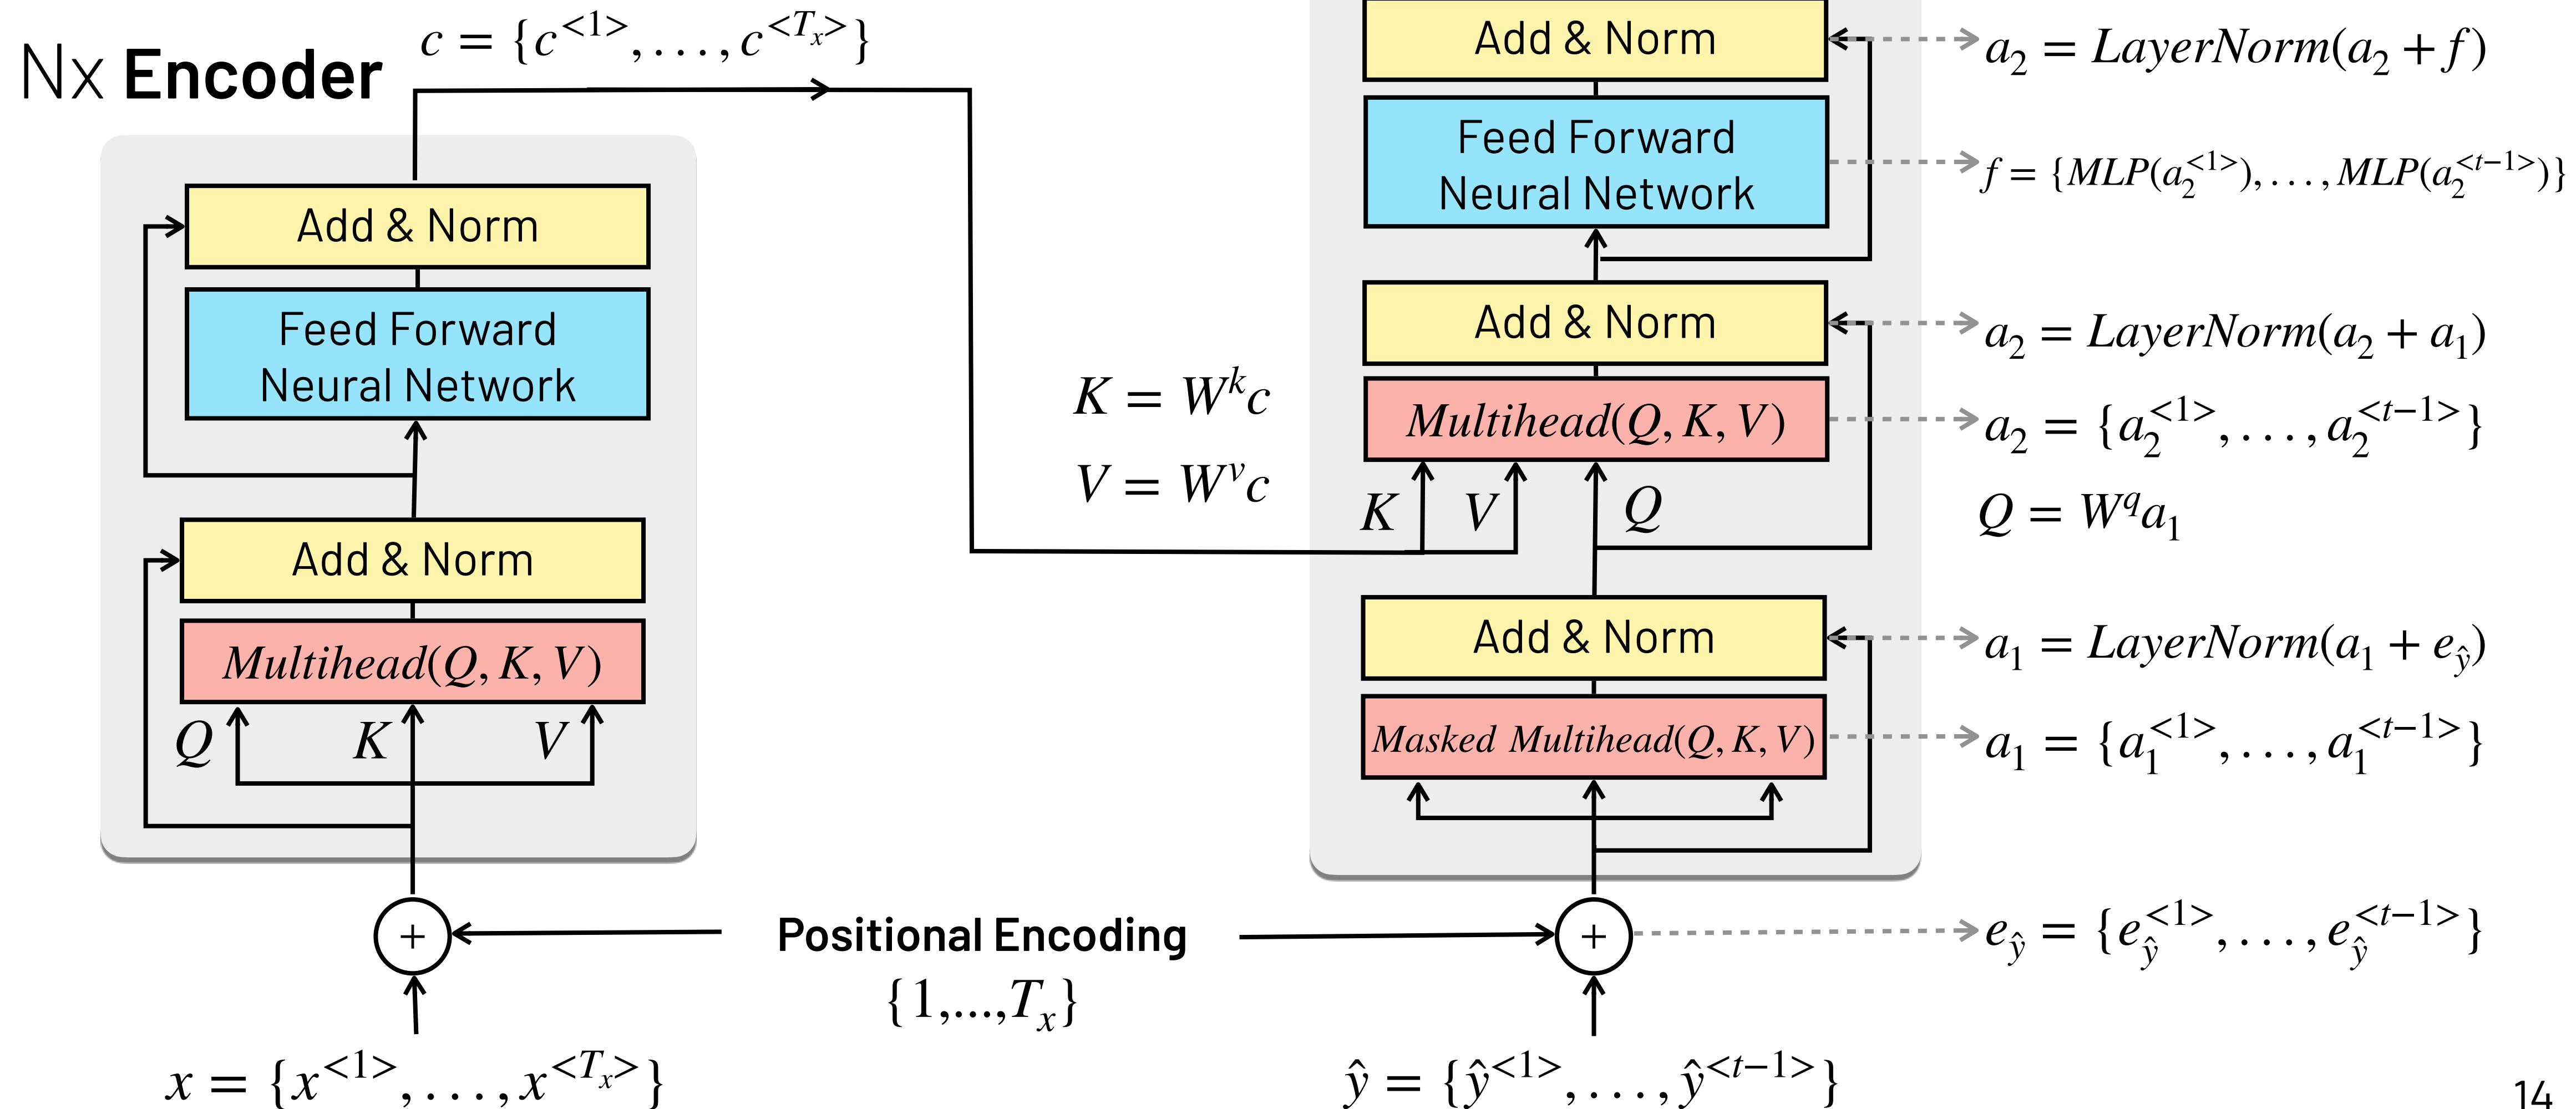
Both are normalized with Layer Norm (**Norm**) and connected with a residual connection (**Add**)



Decoder

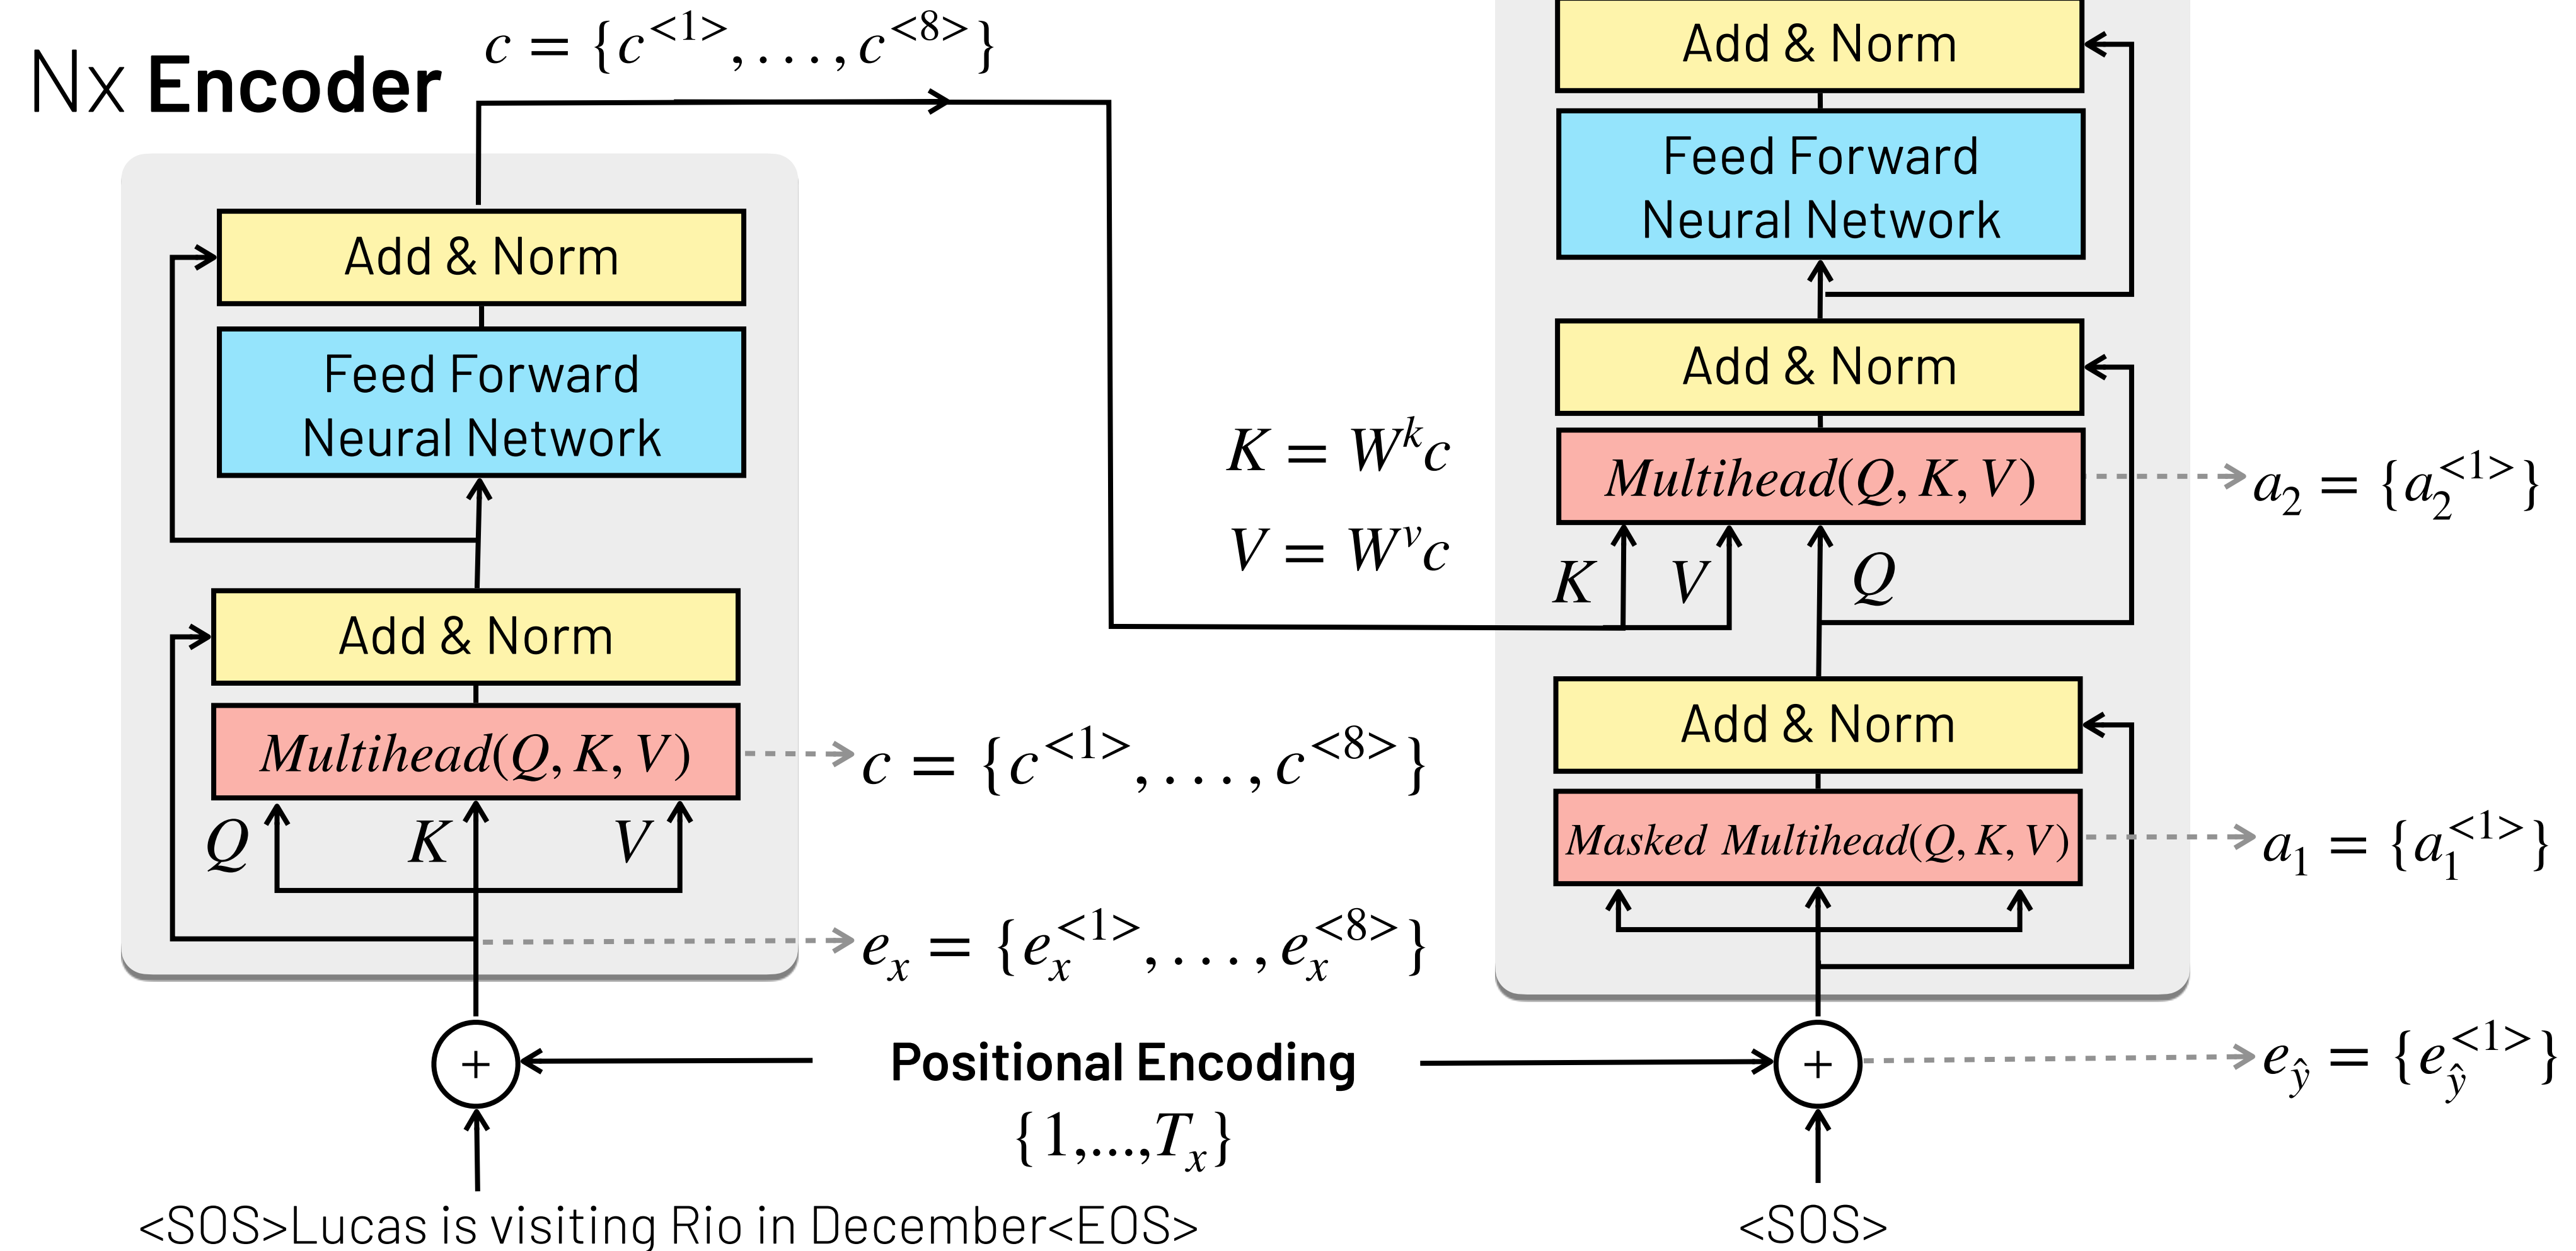
Input: Input context C and previous $\{\hat{y}^{<1>}, \dots, \hat{y}^{<t-1>}\}$ output tokens

Output: The next token $\hat{y}^{<t>}$



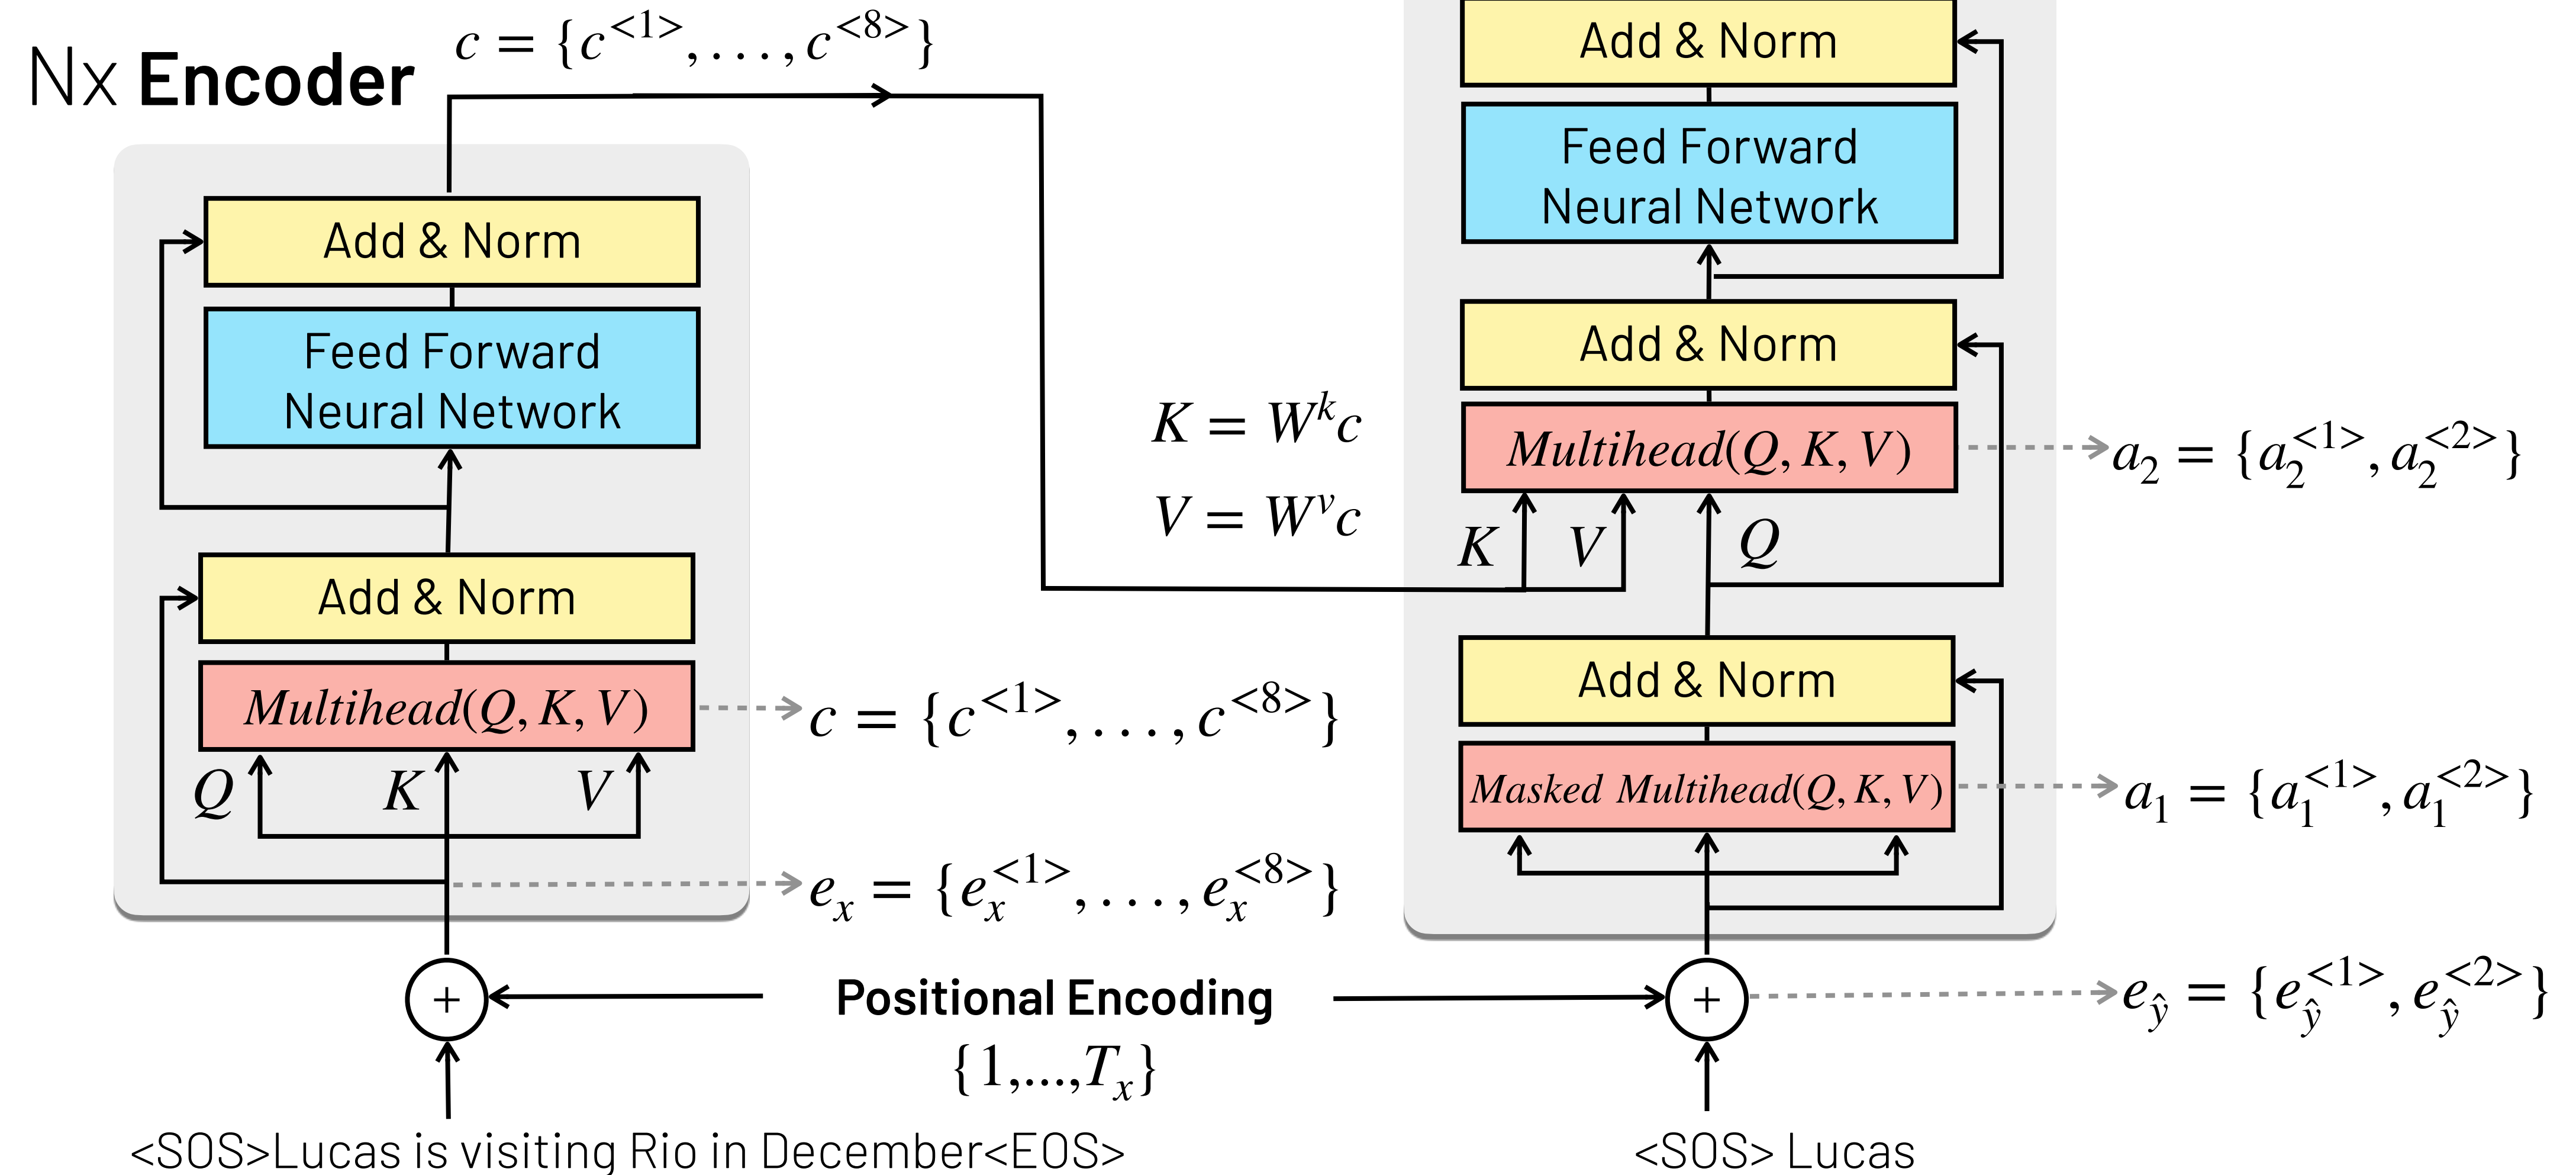
Inference Example

At inference time, the model is autoregressive, i.e., it generates one word at a time.



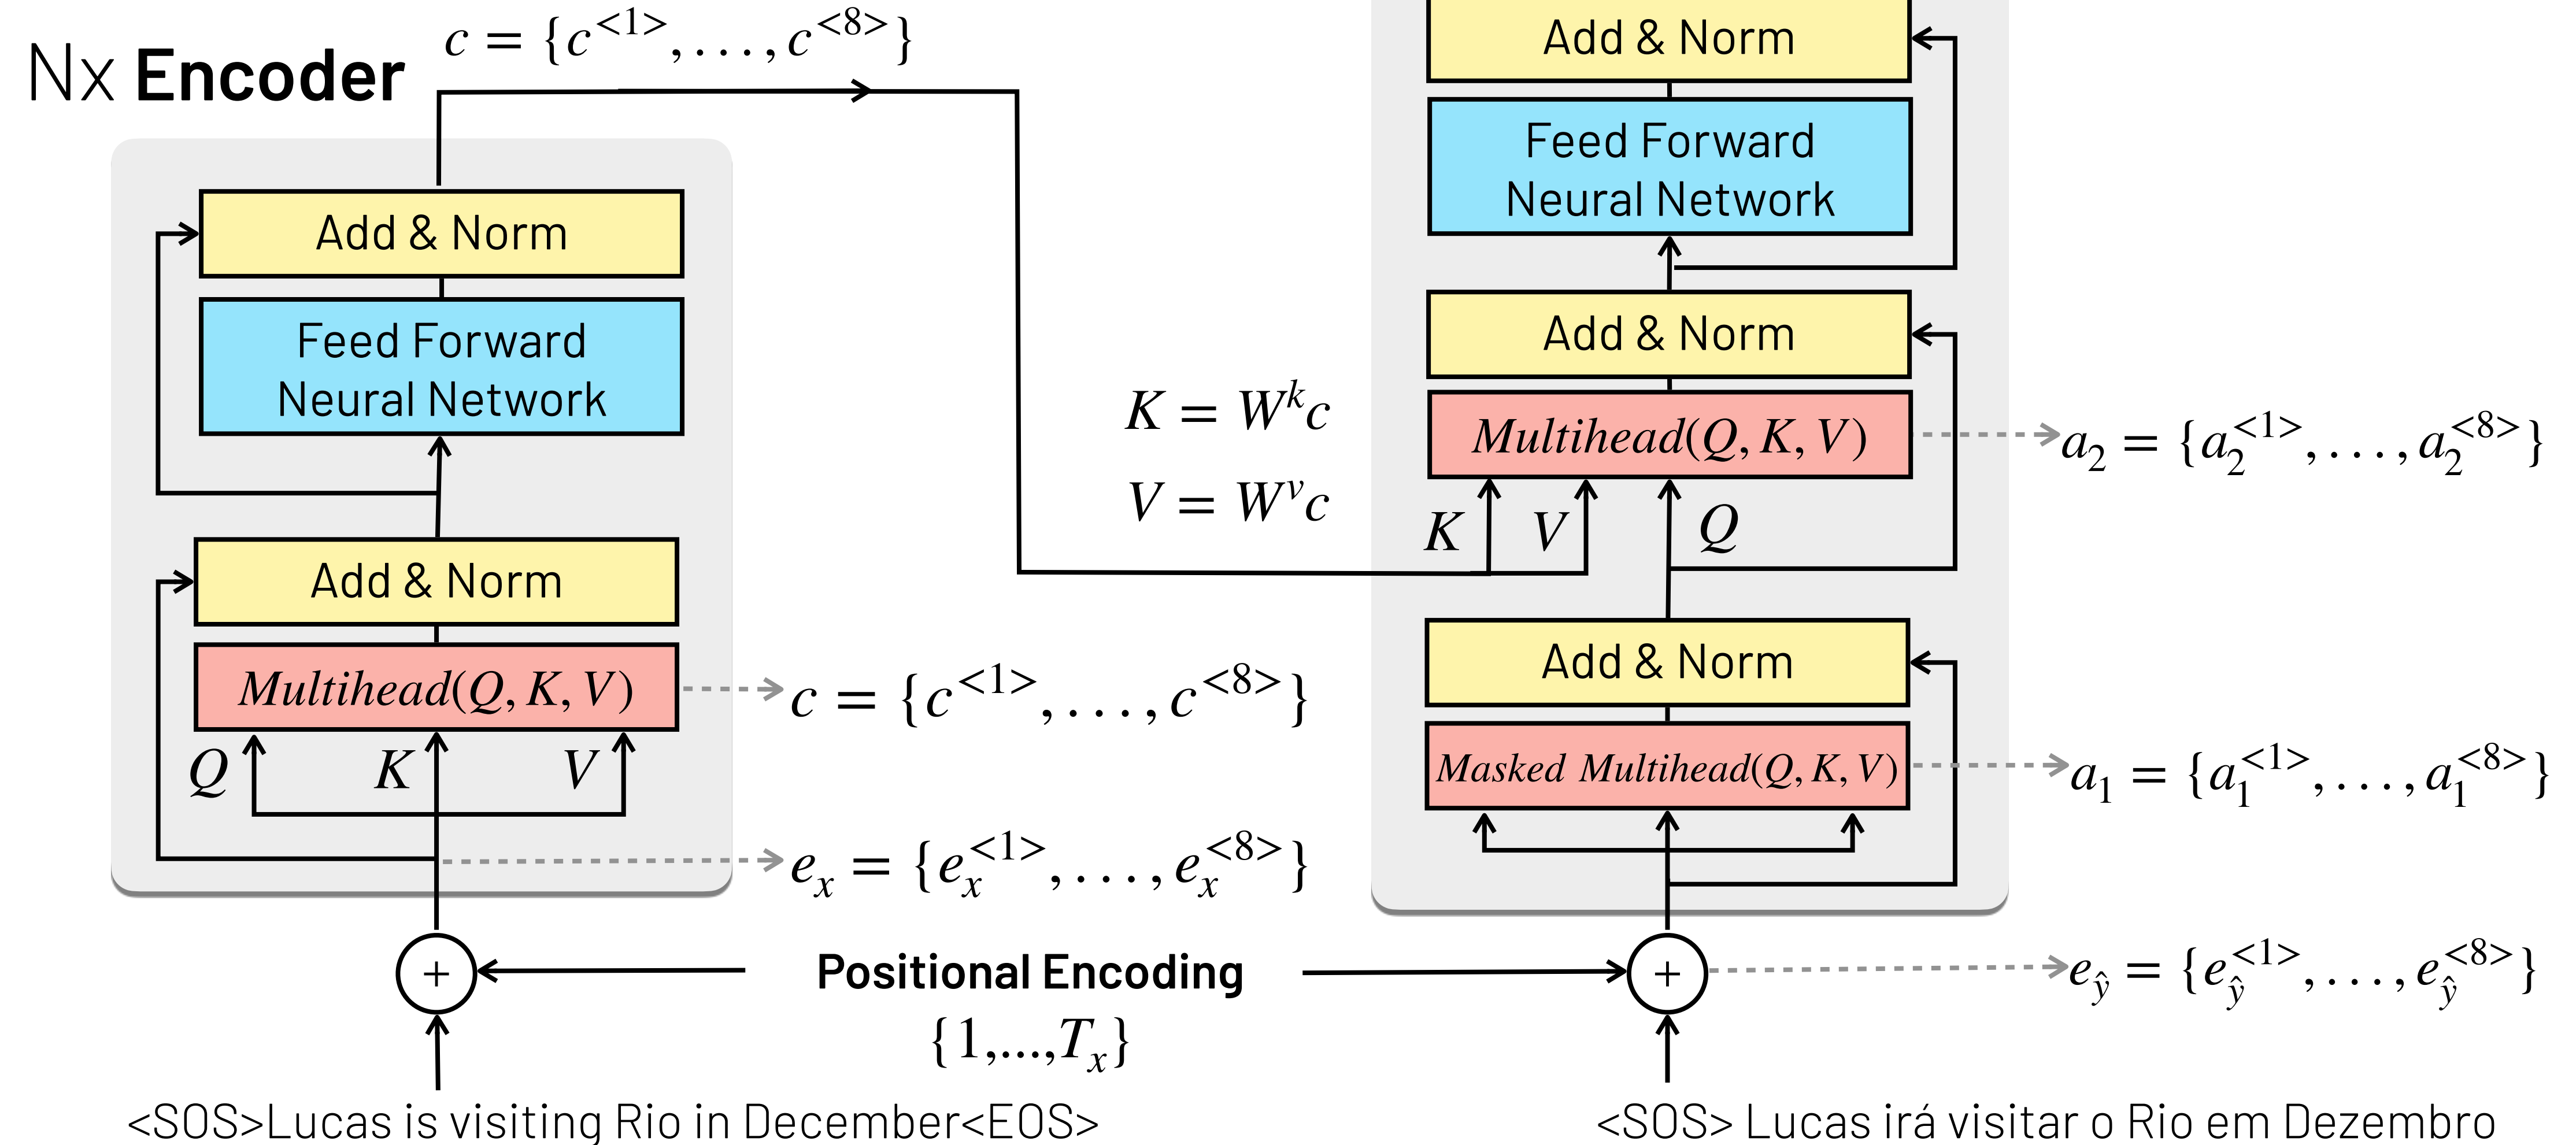
Inference Example

At inference time, the model is autoregressive, i.e., it generates one word at a time.



Inference Example

At inference time, the model is autoregressive, i.e., it generates one word at a time.



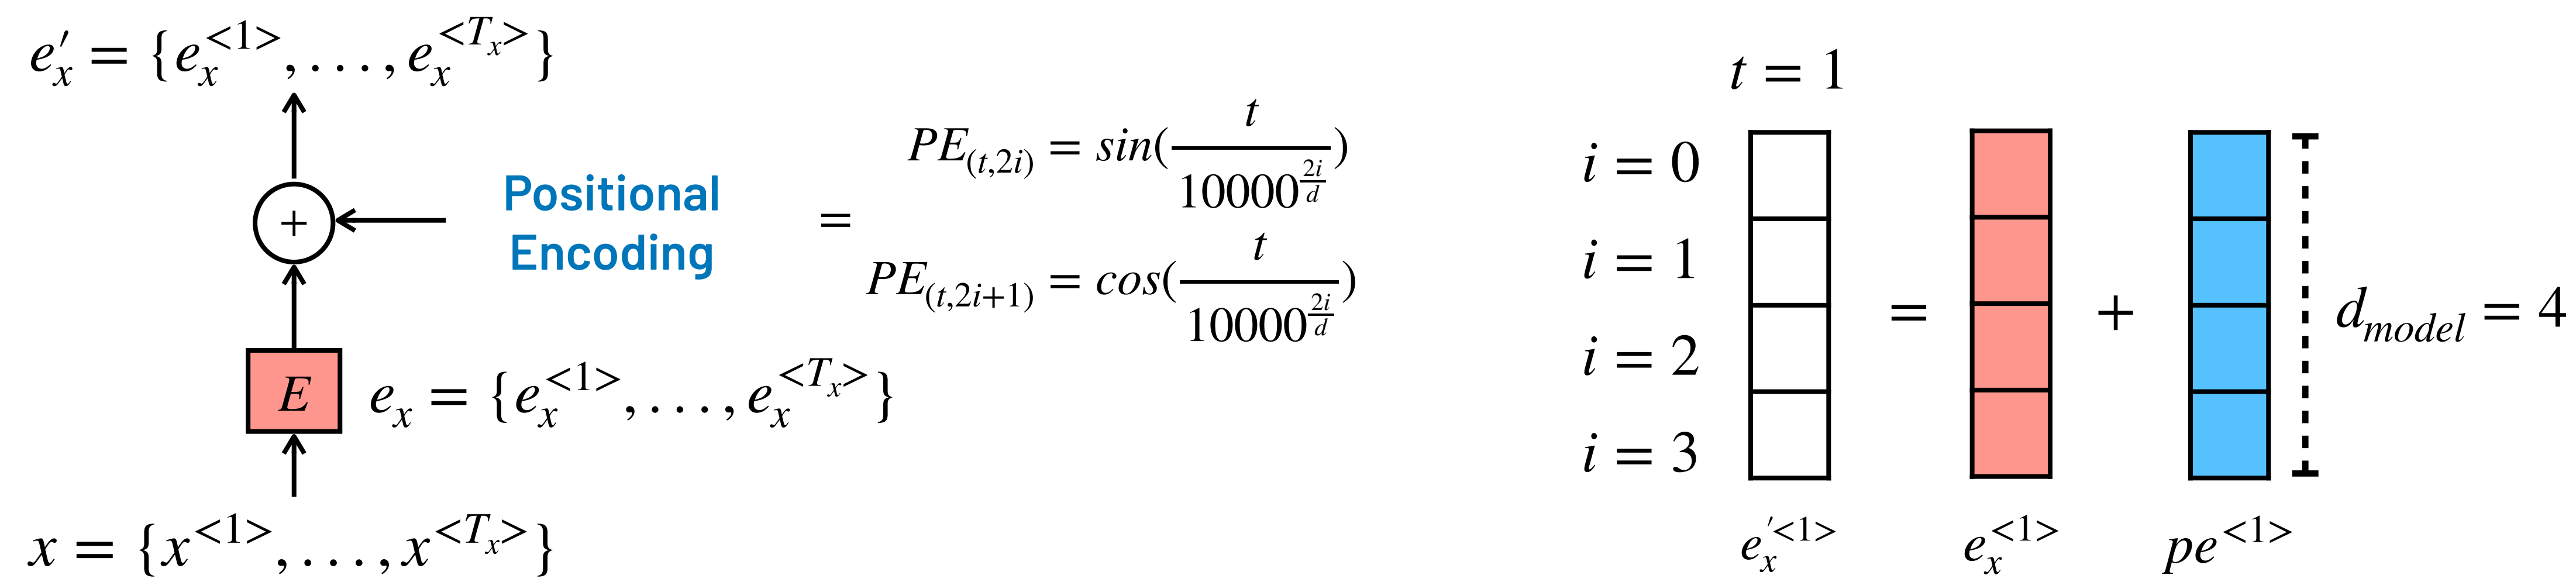
Positional Encoding

The self-attention mechanism does not consider the position of the words.

C

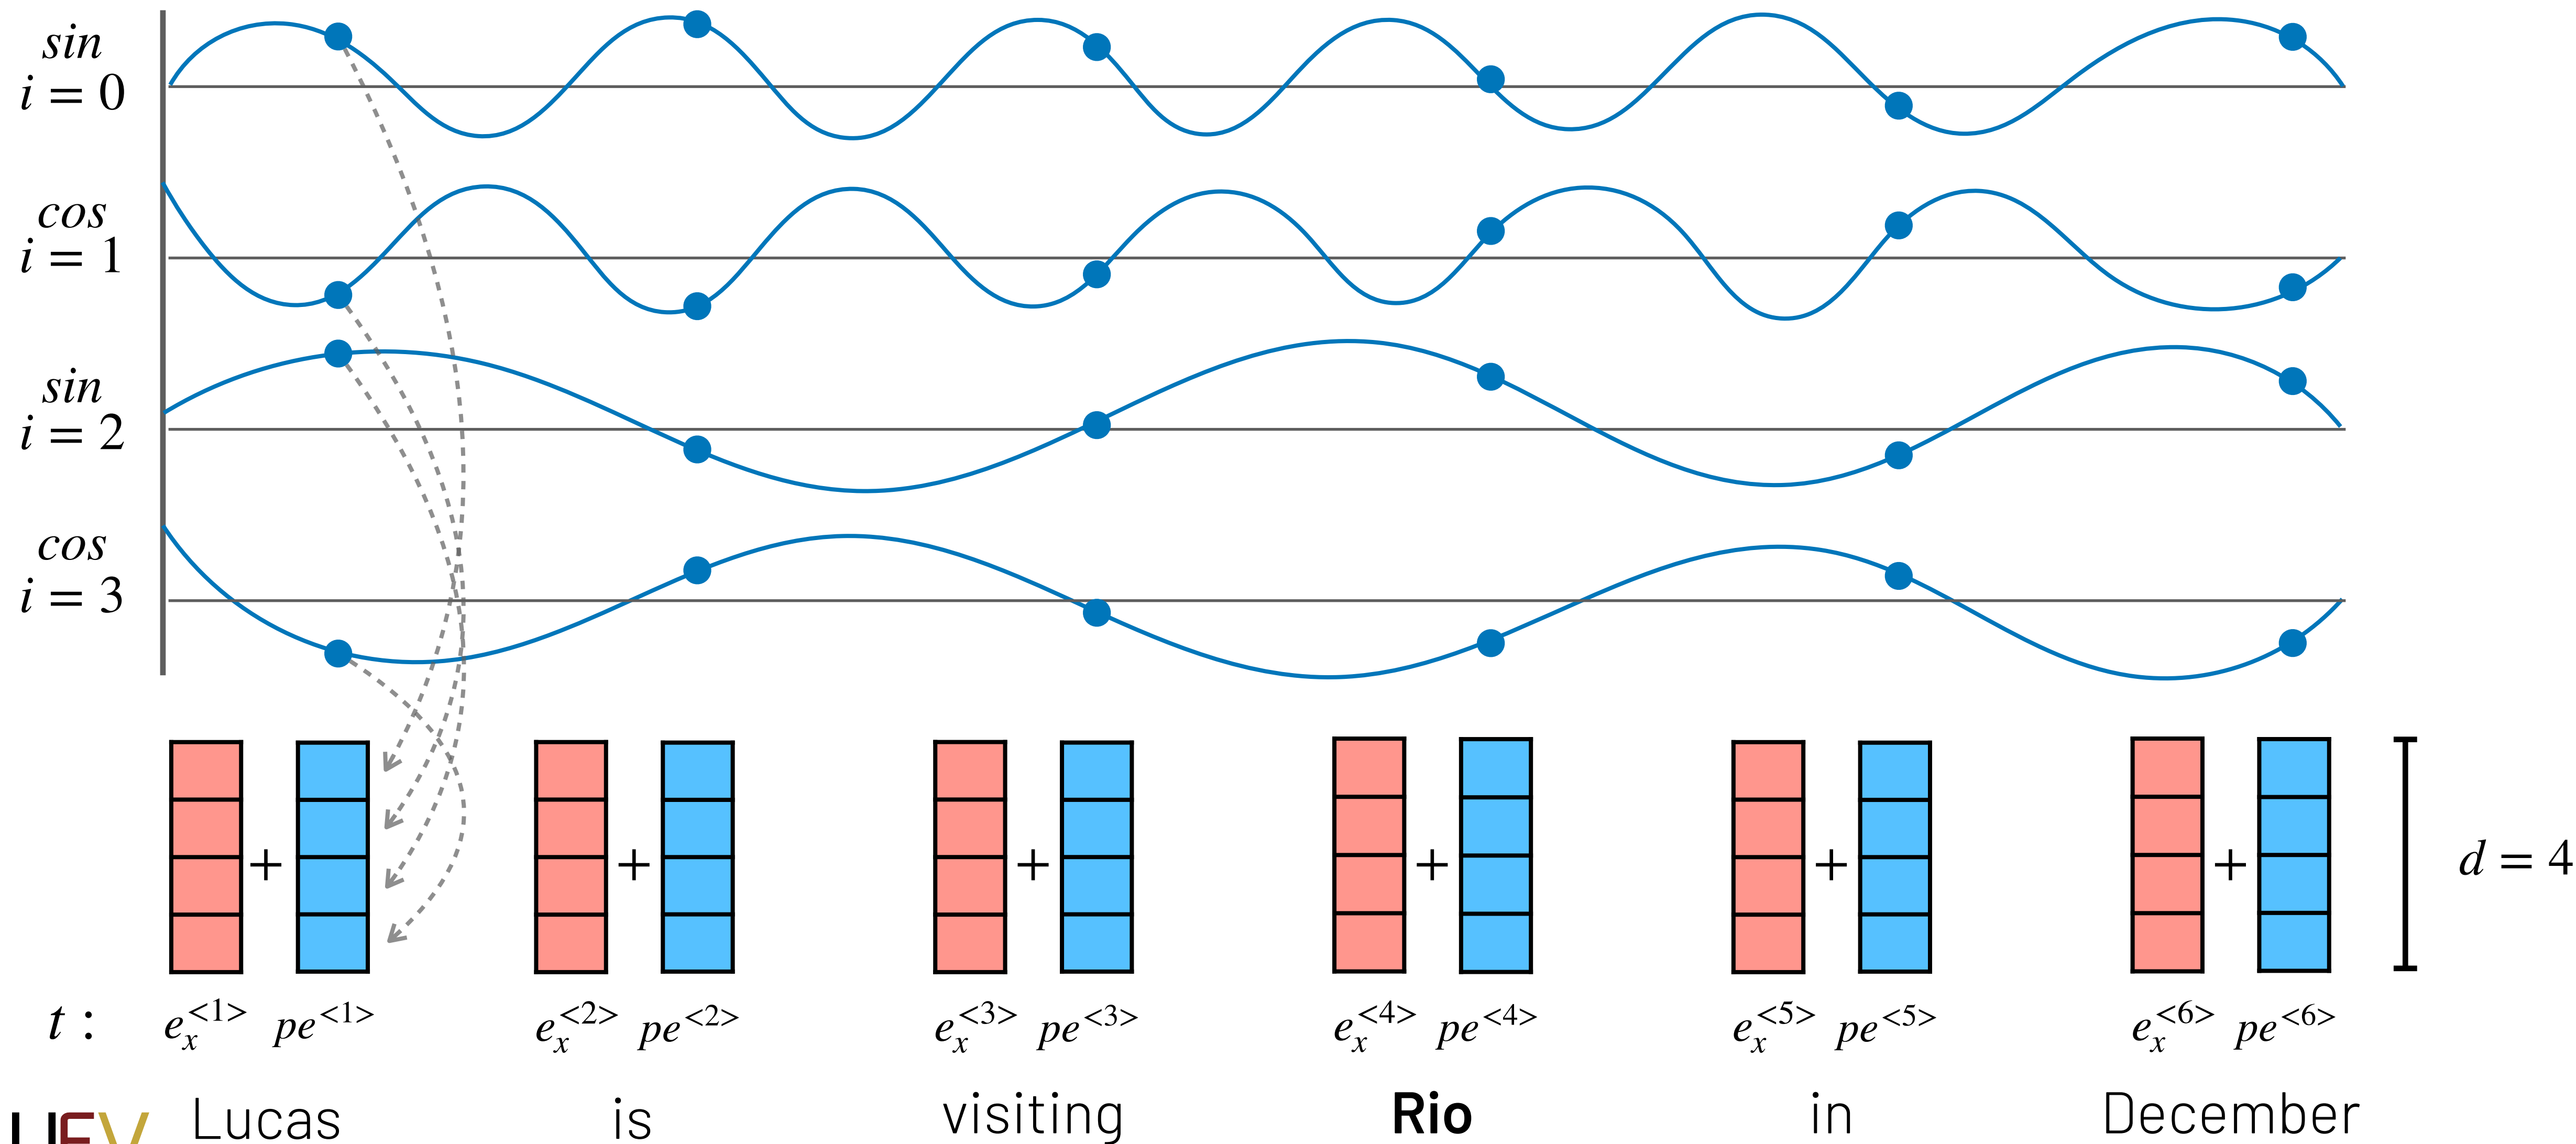
 $= \text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$

To add this information to the learned contextual representation \mathbf{C} , both encoder and decoder add a positional information to each element $x^{<t>}$ of the input $x = \{x^{<1>}, \dots, x^{<T_x>}\}$

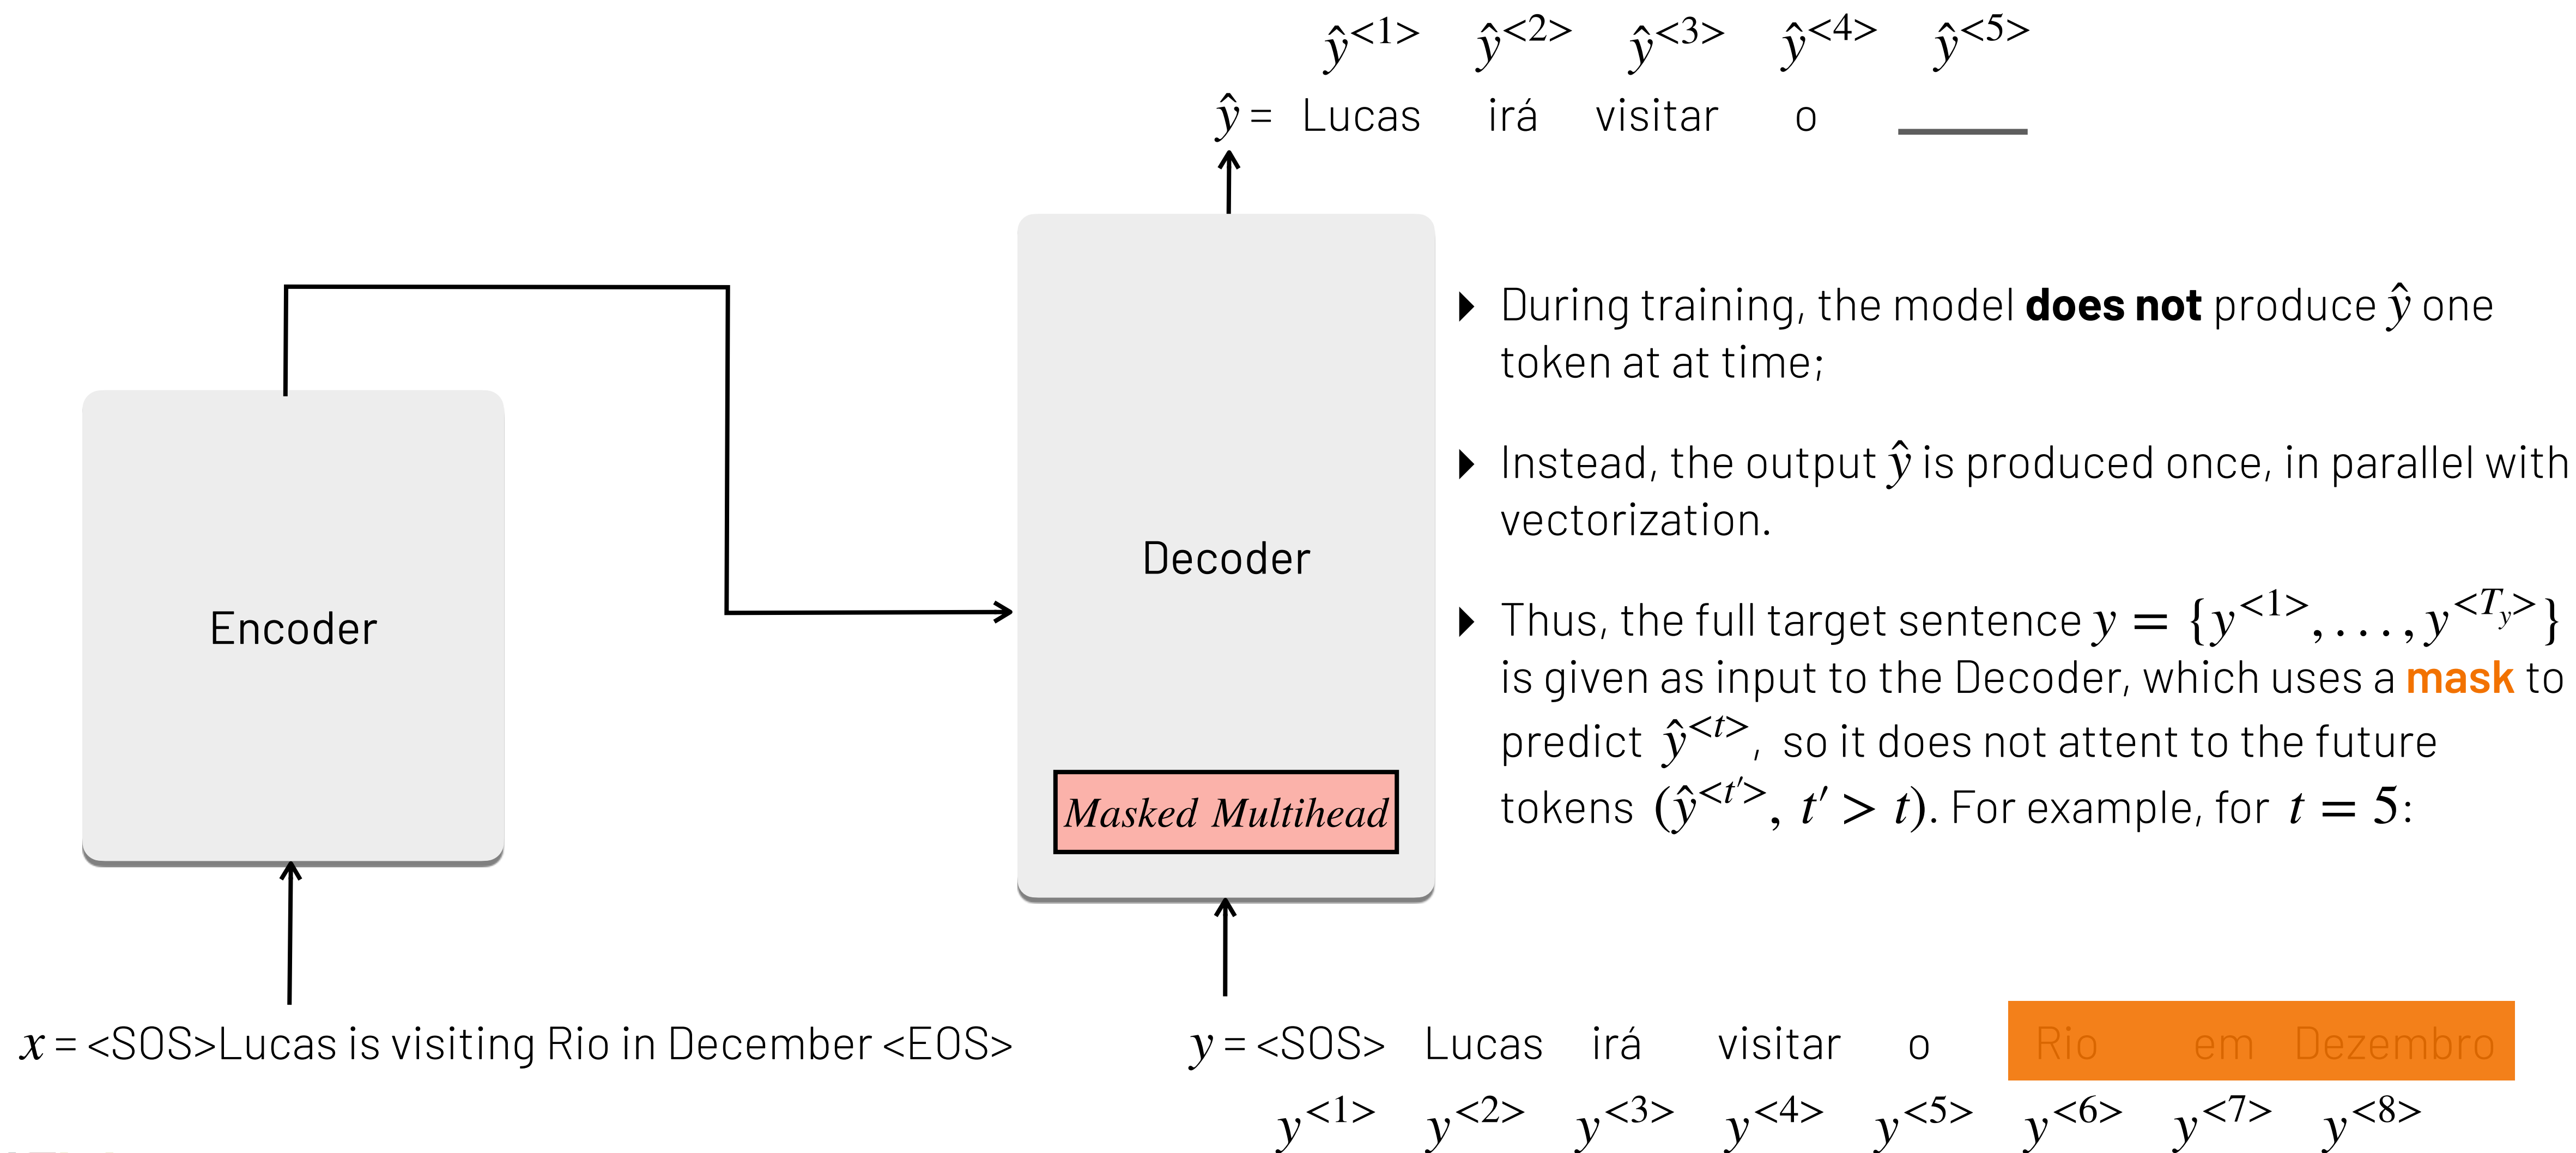


Positional Encoding

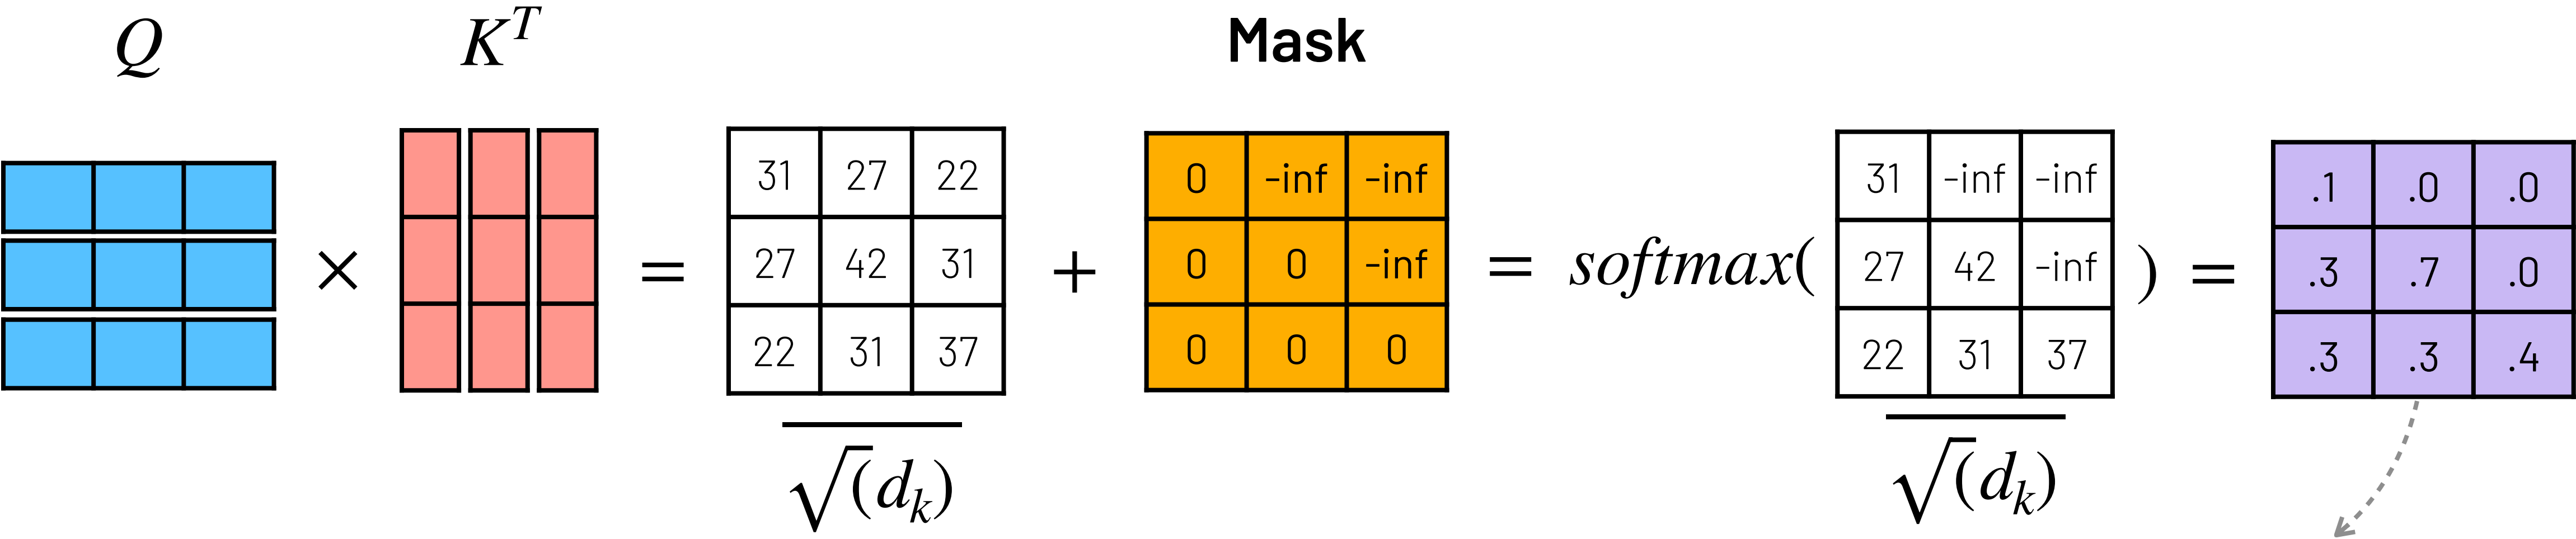
$$PE_{(t,2i)} = \sin\left(\frac{t}{10000^{\frac{2i}{d}}}\right) \quad \text{if } i \text{ is even}$$
$$PE_{(t,2i+1)} = \cos\left(\frac{t}{10000^{\frac{2i}{d}}}\right) \quad \text{if } i \text{ is odd}$$



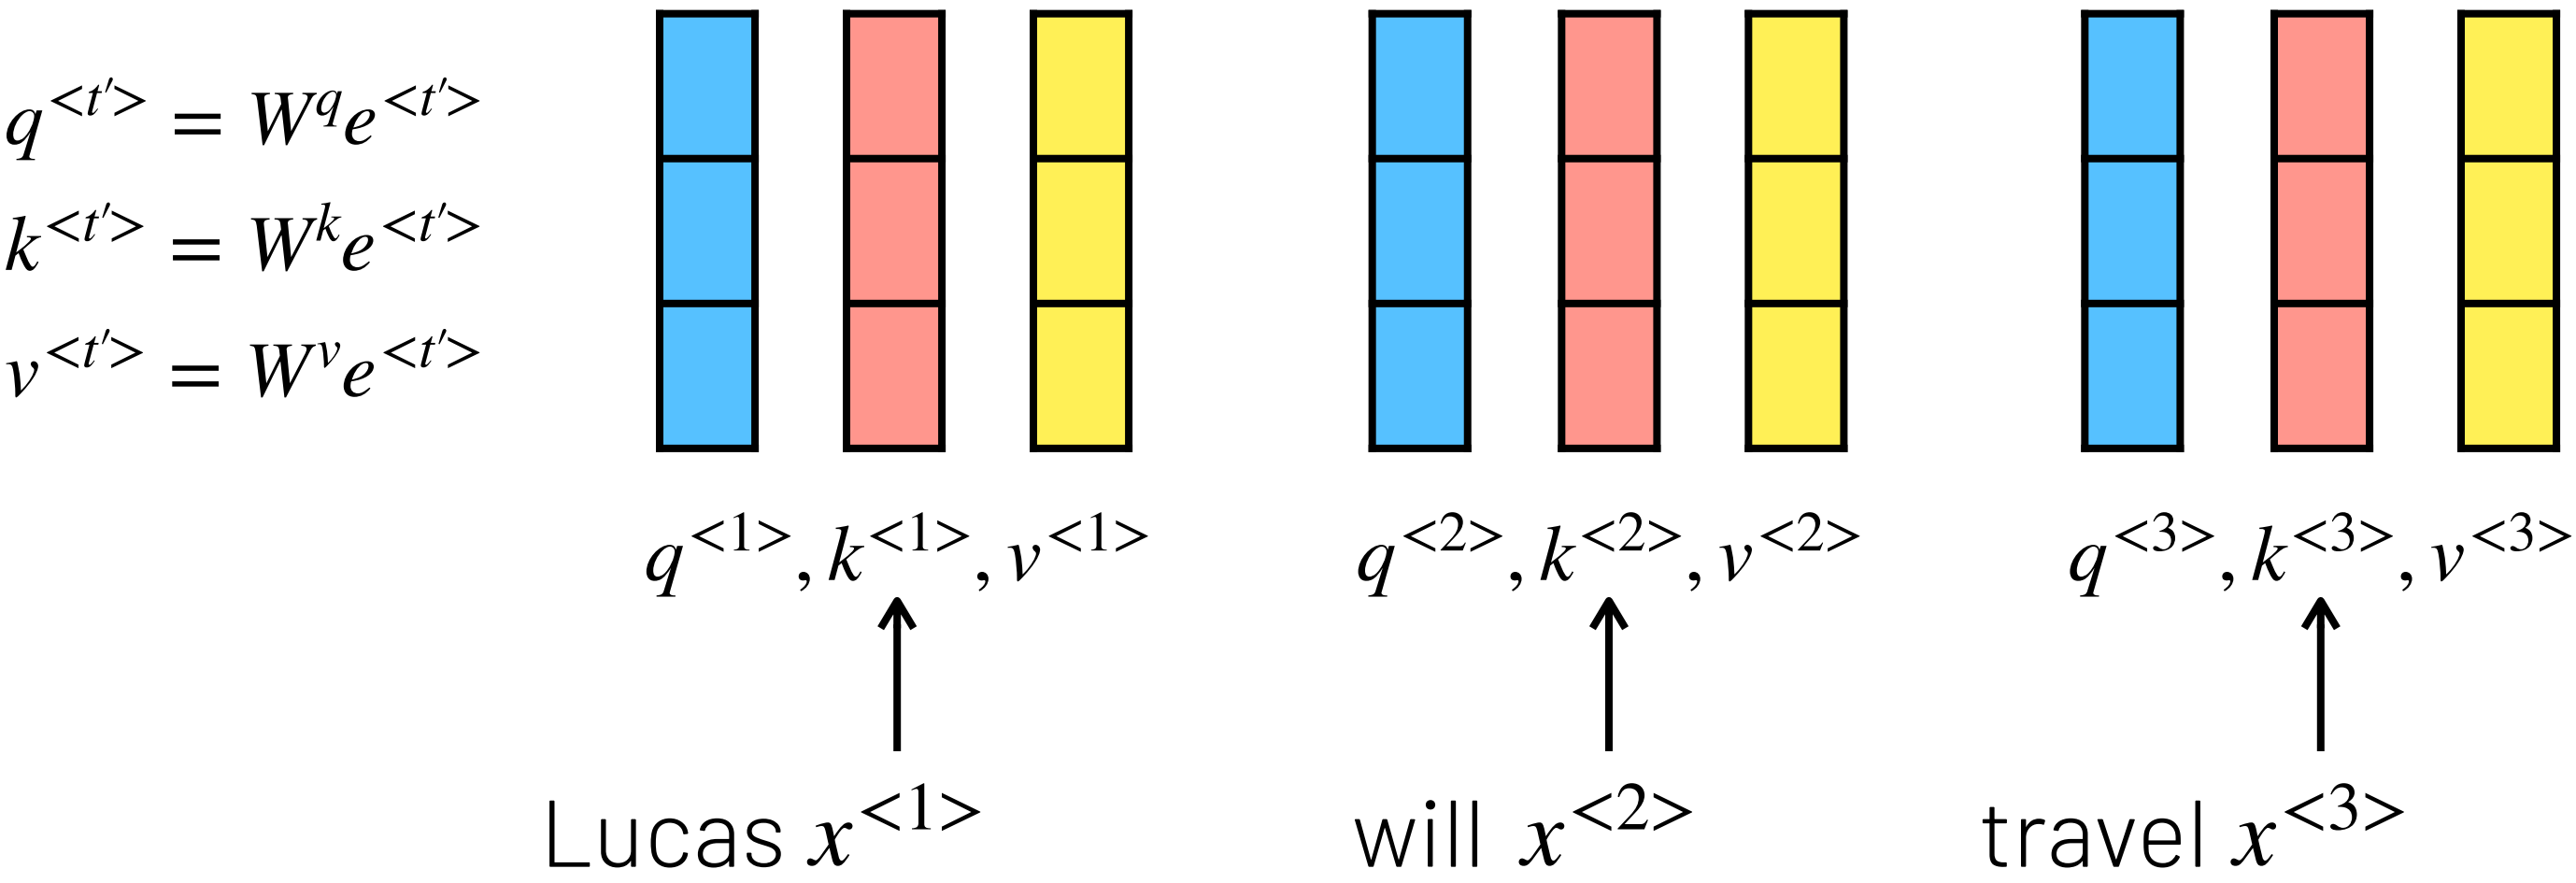
Training with Masked Multi-head Attention



Training with Masked Multi-head Attention



Attention weights (normalized by row) of word i to word j after mask application



Next Lecture

L17: Transformers (Part II)

Case studies of transformers: BERT and GPT