

**INF721**

2024/2

**UFV**

# Deep Learning

## L20: Variational Autoencoders

# Logistics

## Last Lecture

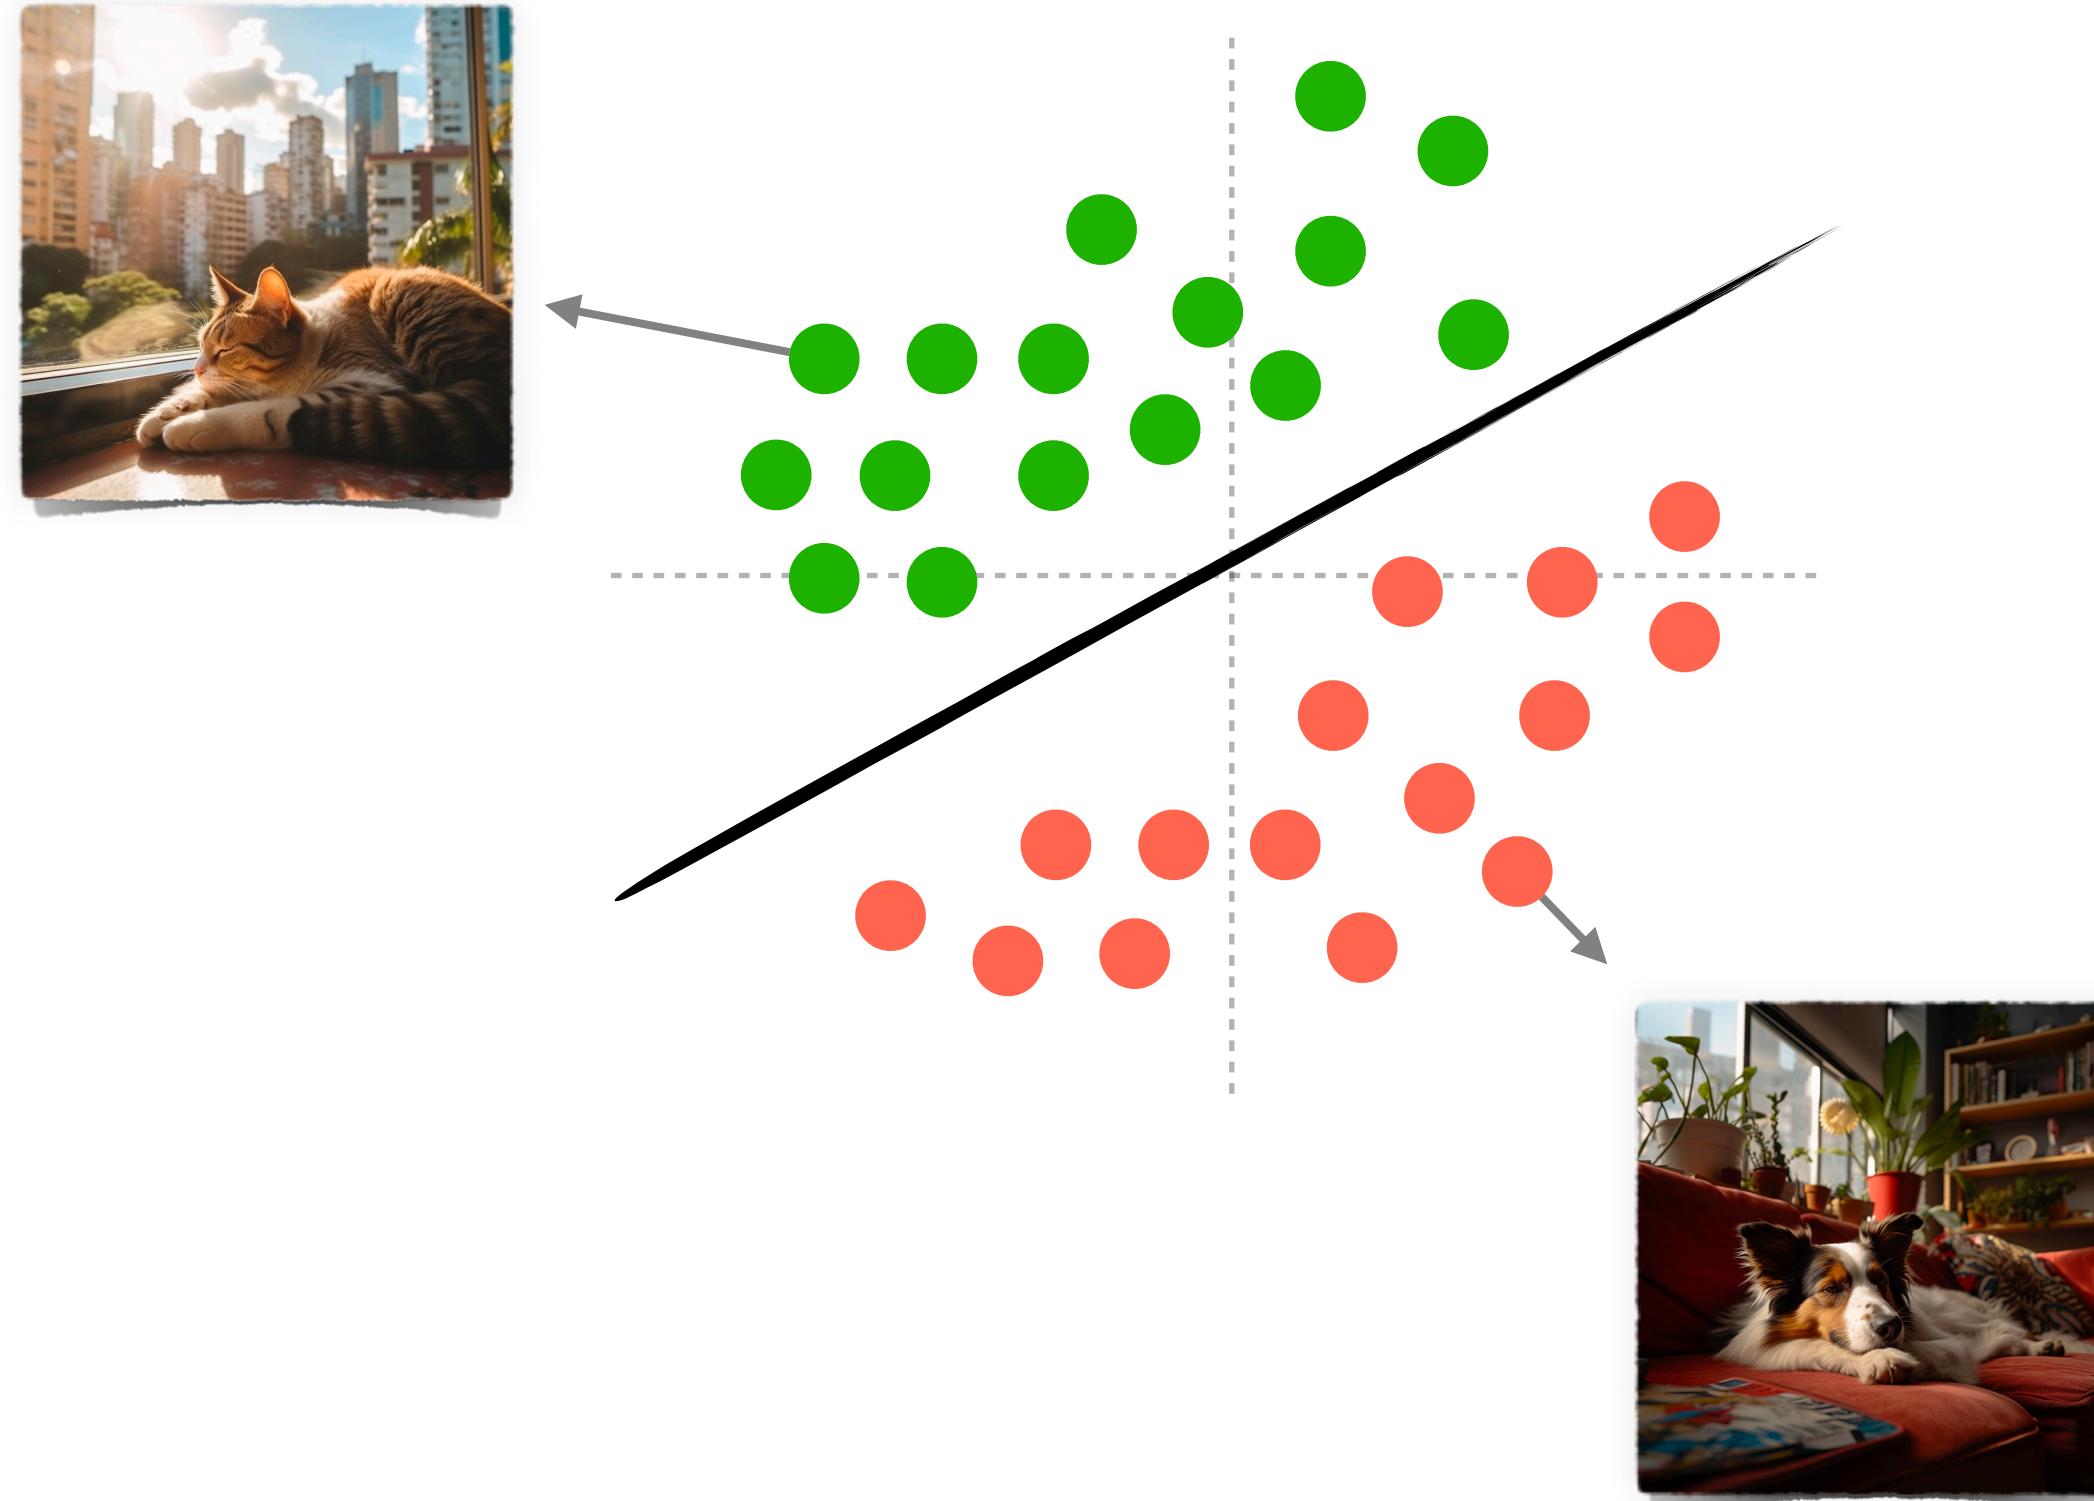
- ▶ Multimodal Learning
- ▶ Visual Transformers (ViT)
- ▶ Text-to-Speech
- ▶ Text-to-Image
- ▶ CLIP

# Lecture Outline

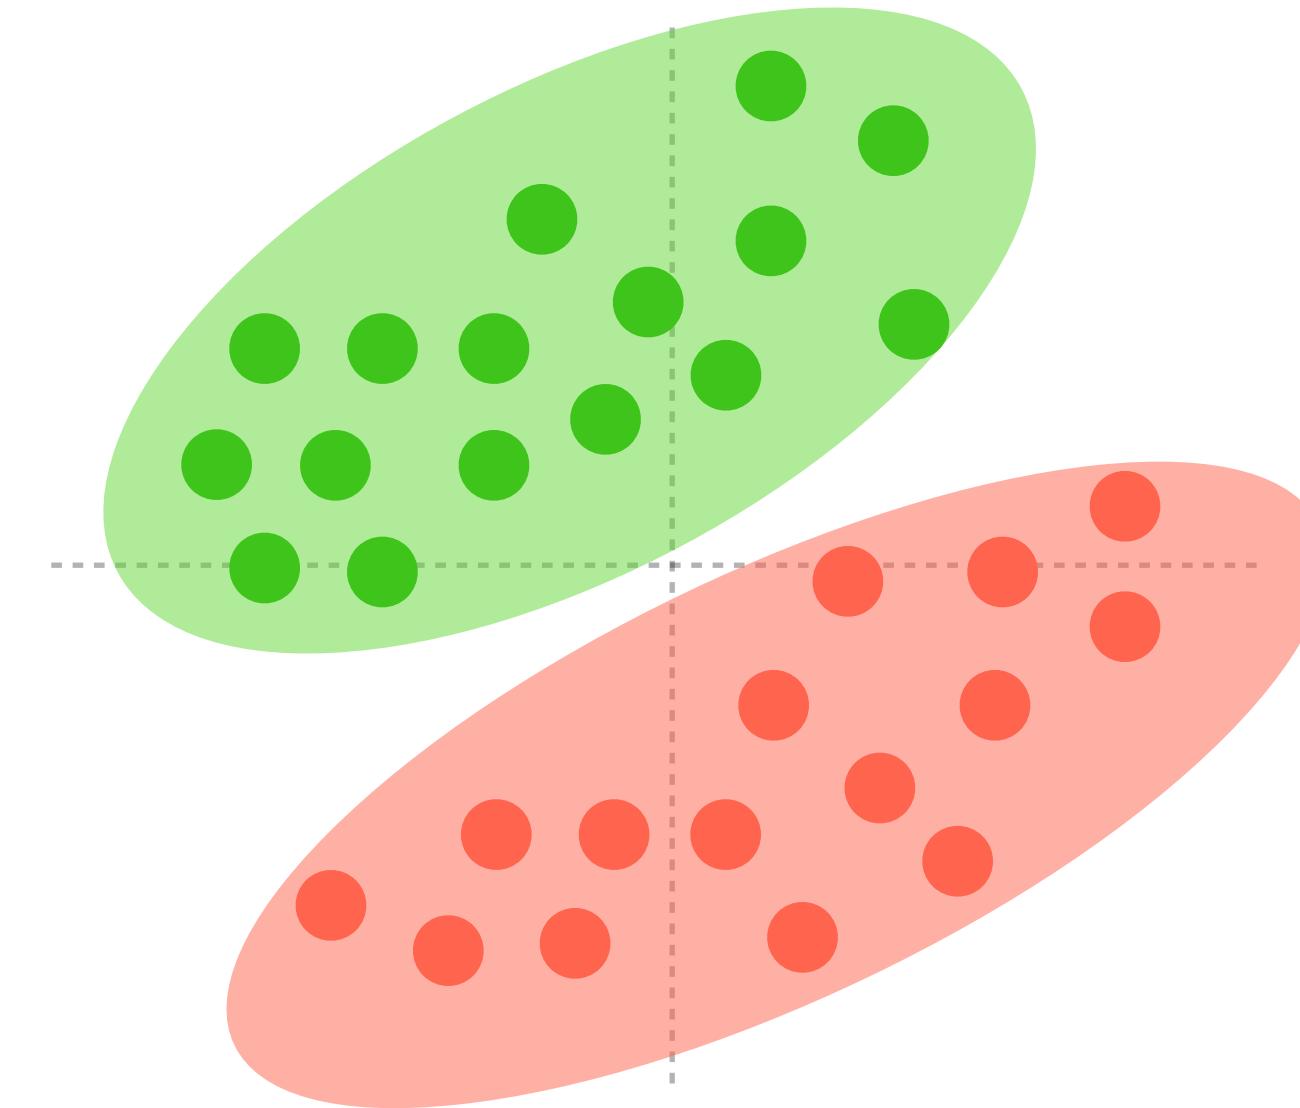
- ▶ Generative Modeling
- ▶ Text Generation
- ▶ Image Generation
- ▶ Autoencoders
  - ▶ Reconstruction Loss
  - ▶ Variational Autoencoders
  - ▶ Kullback-Leibler Divergence

# Discriminative vs. Generative Models

Learn a function  $y = h(\mathbf{x})$  that maps an input feature vector  $\mathbf{x}$  into a label  $y$



Learn a function  $h(\mathbf{x})$  that approximates the distribution that generated the samples  $\mathbf{x}$



- ▶ Autoregressive Models
- ▶ Variational Autoencoders
- ▶ Generative Adversarial Networks
- ▶ Diffusion Models

# Autoregressive Models

In previous lectures, we've discussed RNNs and Transformers as autoregressive models  $P(x^{<t>} | x^{<t-1>}, \dots, x^{<1>})$  to generate sequences:

- ▶ Text in Natural Language
- ▶ Source code in Programming Language
- ▶ Image Generation (as a sequence of pixels or patches)
- ▶ Music Generation
- ▶ Speech
- ▶ ...

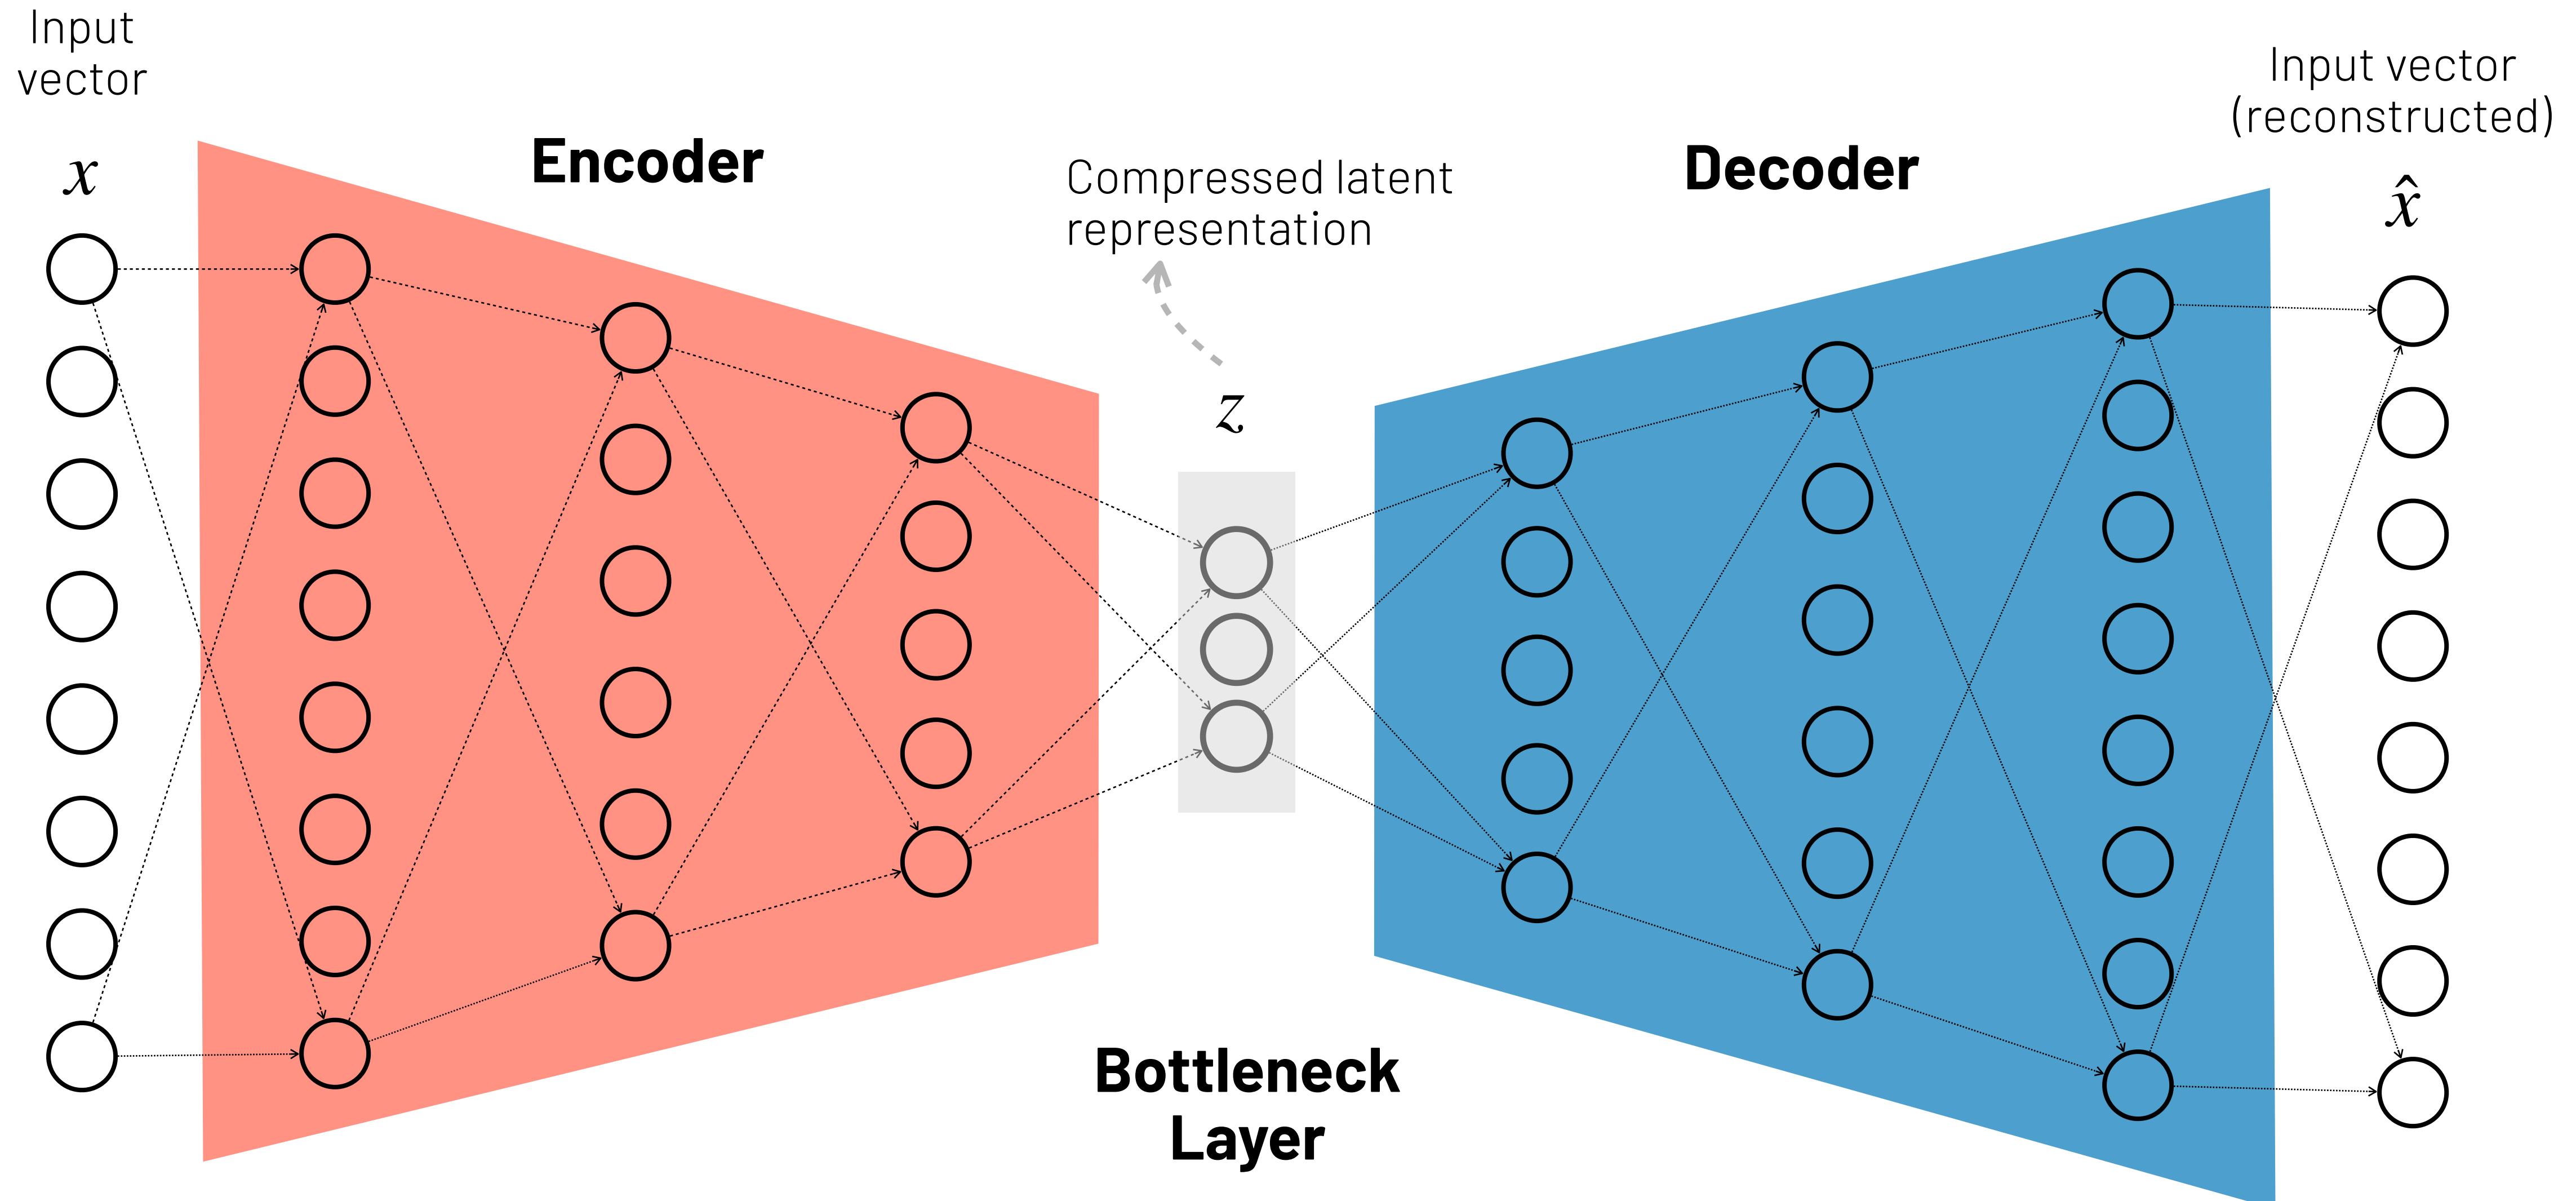
# Image Generation

In next lectures, we will discuss image generation models that learn the probability distribution of pixels without treating them as a sequence:

- ▶ Variational Autoencoders
- ▶ Generative Adversarial Networks
- ▶ Diffusion Models

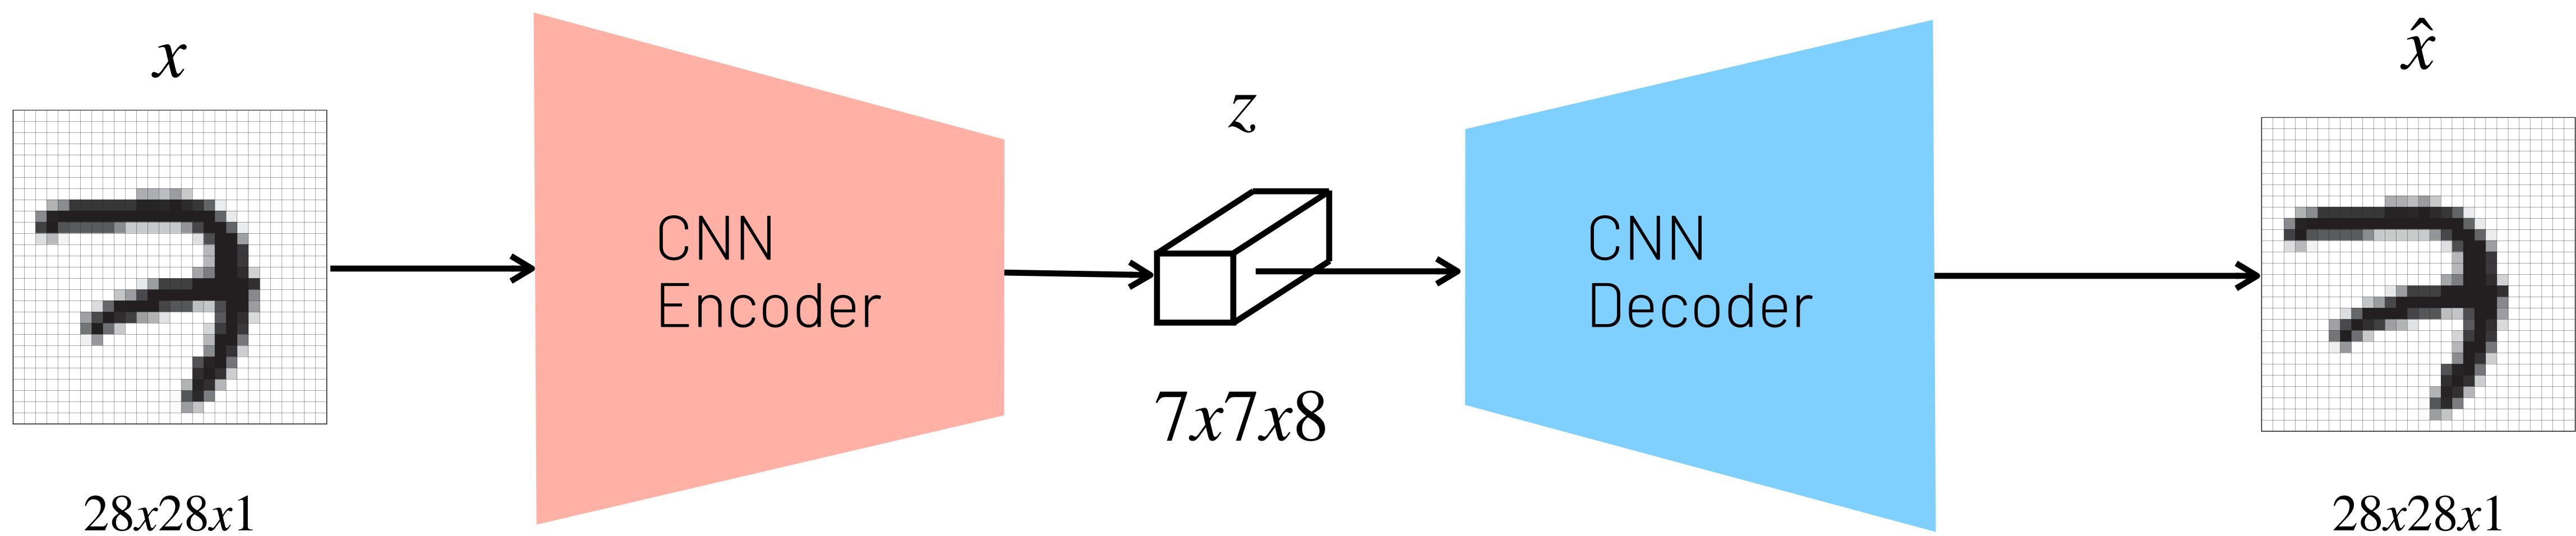
# Autoencoder

**Autoencoder** is an encoder-decoder model designed to learn to compress and reconstruct data



# Example: Image Compression

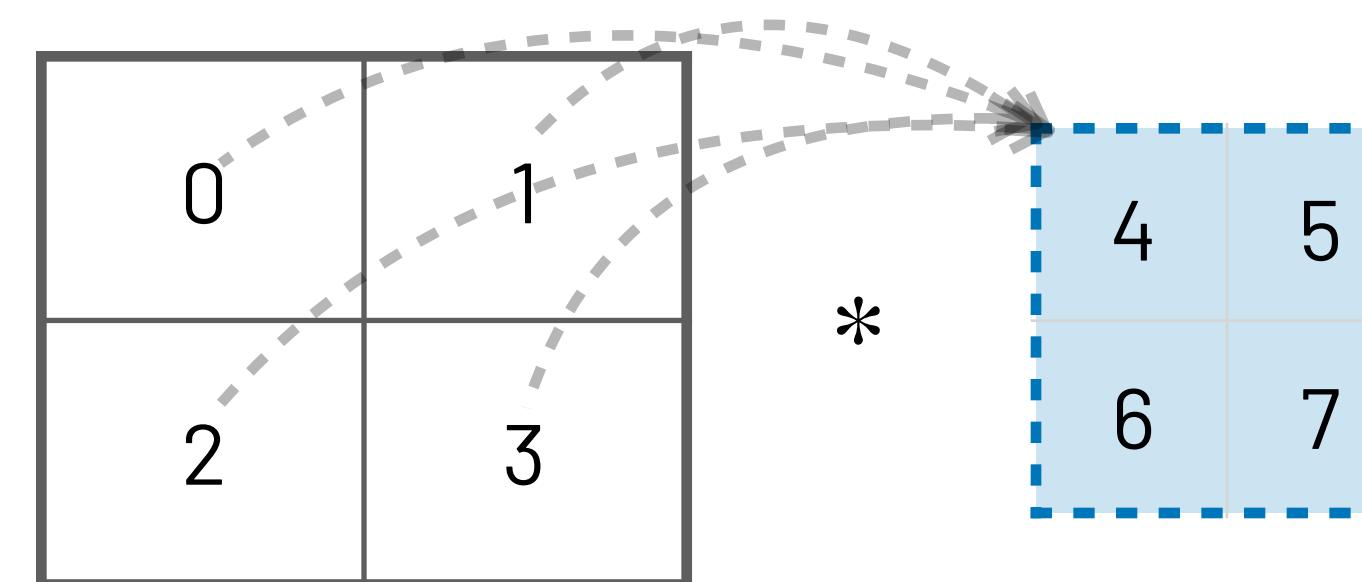
We can apply autoencoders to learn compressed representations of images using a CNN as an encoder and decoder



How to map a feature map back to an image?  
**Transposed Convolution!**

# Transposed Convolutions

A transposed convolutional layer slides the input over the kernel and performs element-wise multiplication and summation:



Input  $X$   
( $2 \times 2$ )

Filter  $f$   
( $2 \times 2$ )

$$\begin{array}{c} \begin{array}{|c|c|c|} \hline 0 & 0 & \\ \hline 0 & 0 & \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline & 4 & 5 \\ \hline & 6 & 7 \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline & & \\ \hline 8 & 10 & \\ \hline 12 & 14 & \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline & & \\ \hline & 12 & 15 \\ \hline & 18 & 21 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 0 & 4 & 5 \\ \hline 0 & 28 & 22 \\ \hline 12 & 32 & 21 \\ \hline \end{array} \end{array}$$

$$X_{0,0} * f = 0 * f$$

$$X_{0,1} * f = 1 * f$$

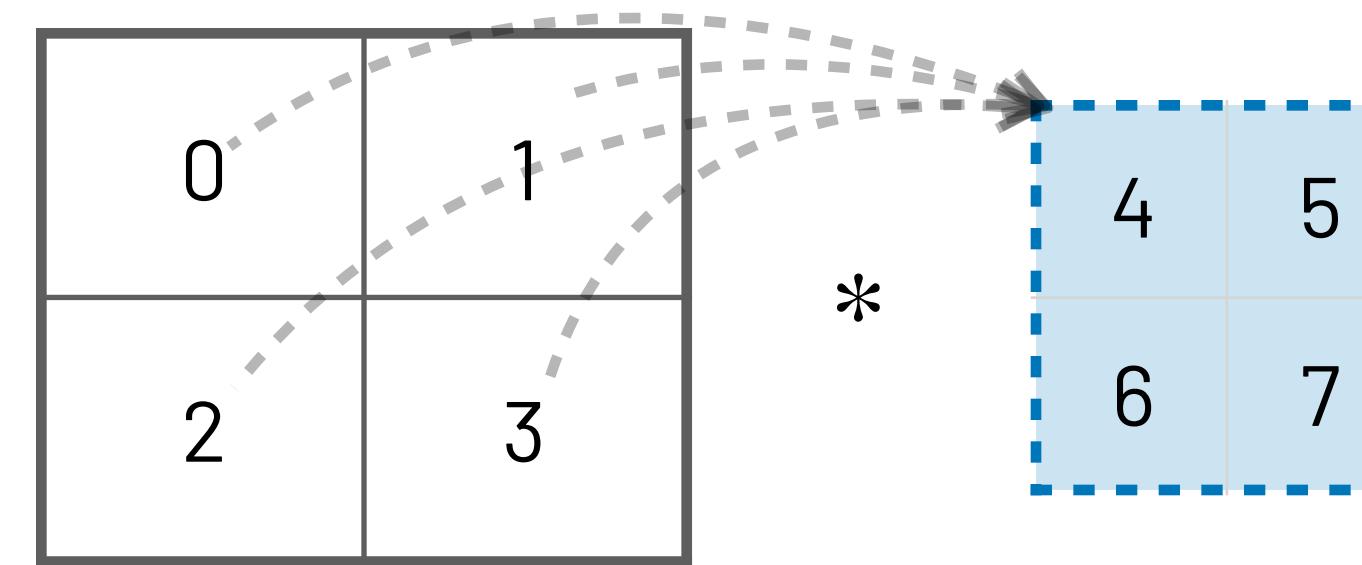
$$X_{1,0} * f = 2 * f$$

$$X_{1,1} * f = 3 * f$$

Output  
( $3 \times 3$ )

# Transposed Convolutions with Stride

In the transposed convolution, strides are specified for the the output, not for input. Let's performed the previous transposed convolution with stide of 2



$$\begin{array}{c} \begin{array}{|c|c|c|c|} \hline 0 & 0 & & \\ \hline 0 & 0 & & \\ \hline & & & \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline & & 4 & 5 \\ \hline & & 6 & 7 \\ \hline & & & \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & 8 & 10 \\ \hline & & 12 & 14 \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & 12 & 15 \\ \hline & & 18 & 21 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 0 & 4 & 5 \\ \hline 0 & 0 & 6 & 7 \\ \hline 8 & 10 & 12 & 15 \\ \hline 12 & 14 & 18 & 21 \\ \hline \end{array} \end{array}$$

$$X_{0,0} * f = 0 * f$$

$$X_{0,1} * f = 1 * f$$

$$X_{1,0} * f = 2 * f$$

$$X_{1,1} * f = 3 * f$$

$$\text{Output} \quad (4 \times 4)$$

# Transposed Convolutions with Padding

Different from in the regular convolution where padding is applied to input, it is applied to output in the transposed convolution.

0	1
2	3

Input  
( $3 \times 3$ )

4	5
6	7

Filter  
( $2 \times 2$ )

We simply remove the first  
and last row and column

0	0	
0	0	

	4	5
	6	7

+

8	10	
12	14	

+

	12	15
	18	21

+

0	4	5
0	4	6
4	12	9

=

$$X_{0,0} * f = 0 * f$$

$$X_{0,1} * f = 1 * f$$

$$X_{1,0} * f = 2 * f$$

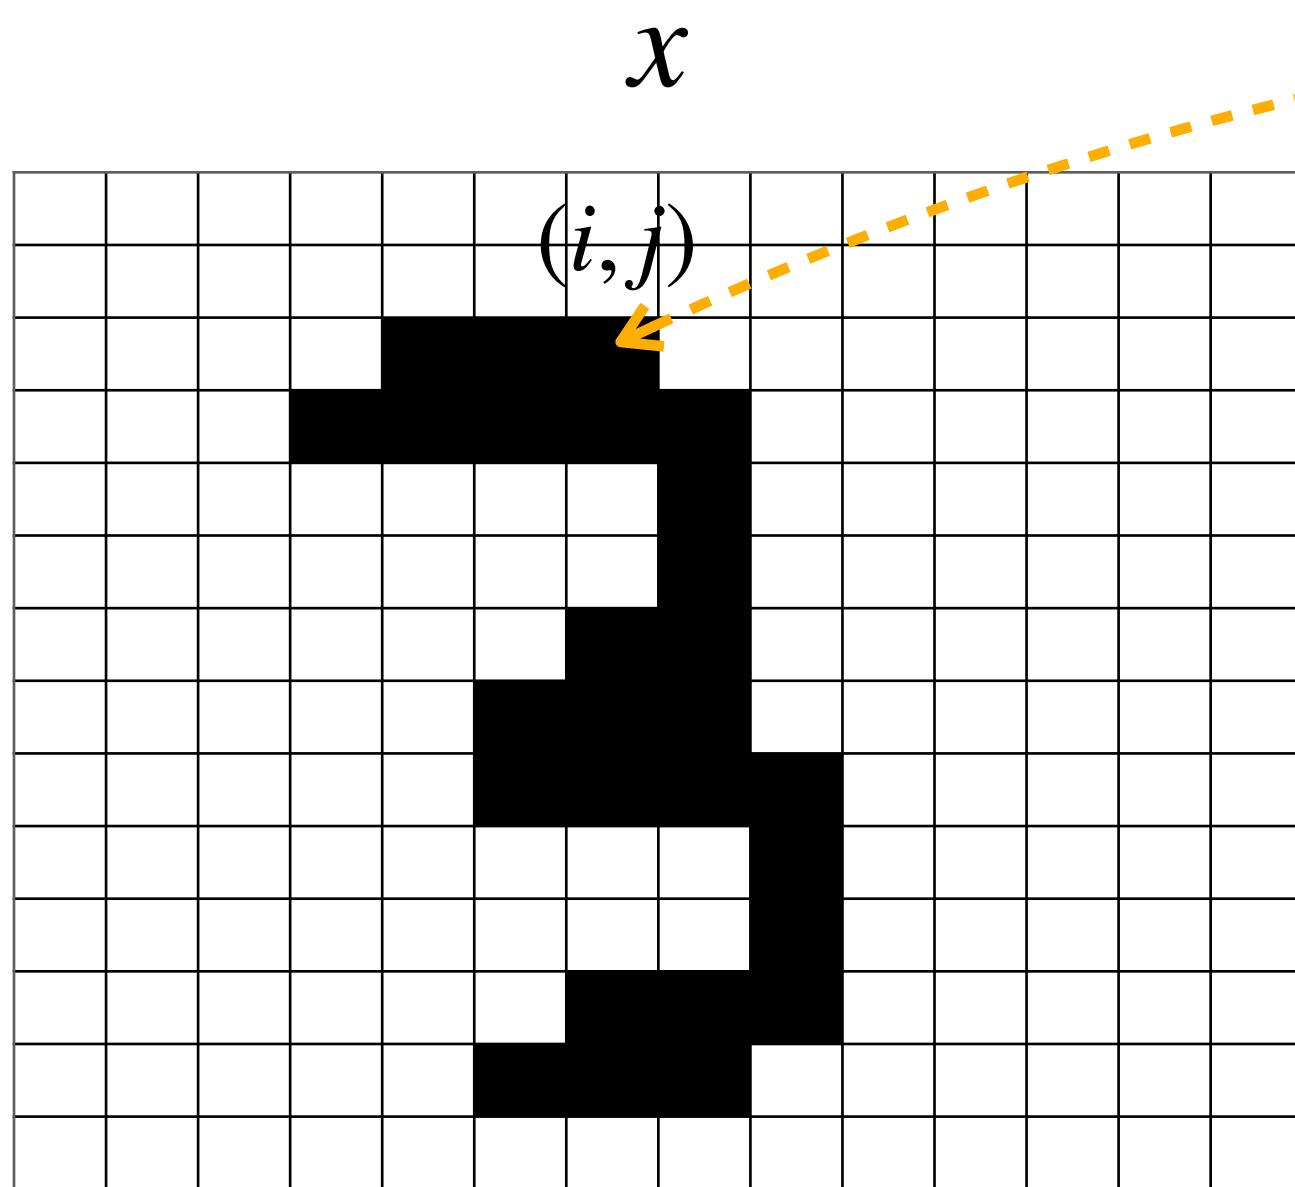
$$X_{1,1} * f = 3 * f$$

Output  
( $3 \times 3$ )

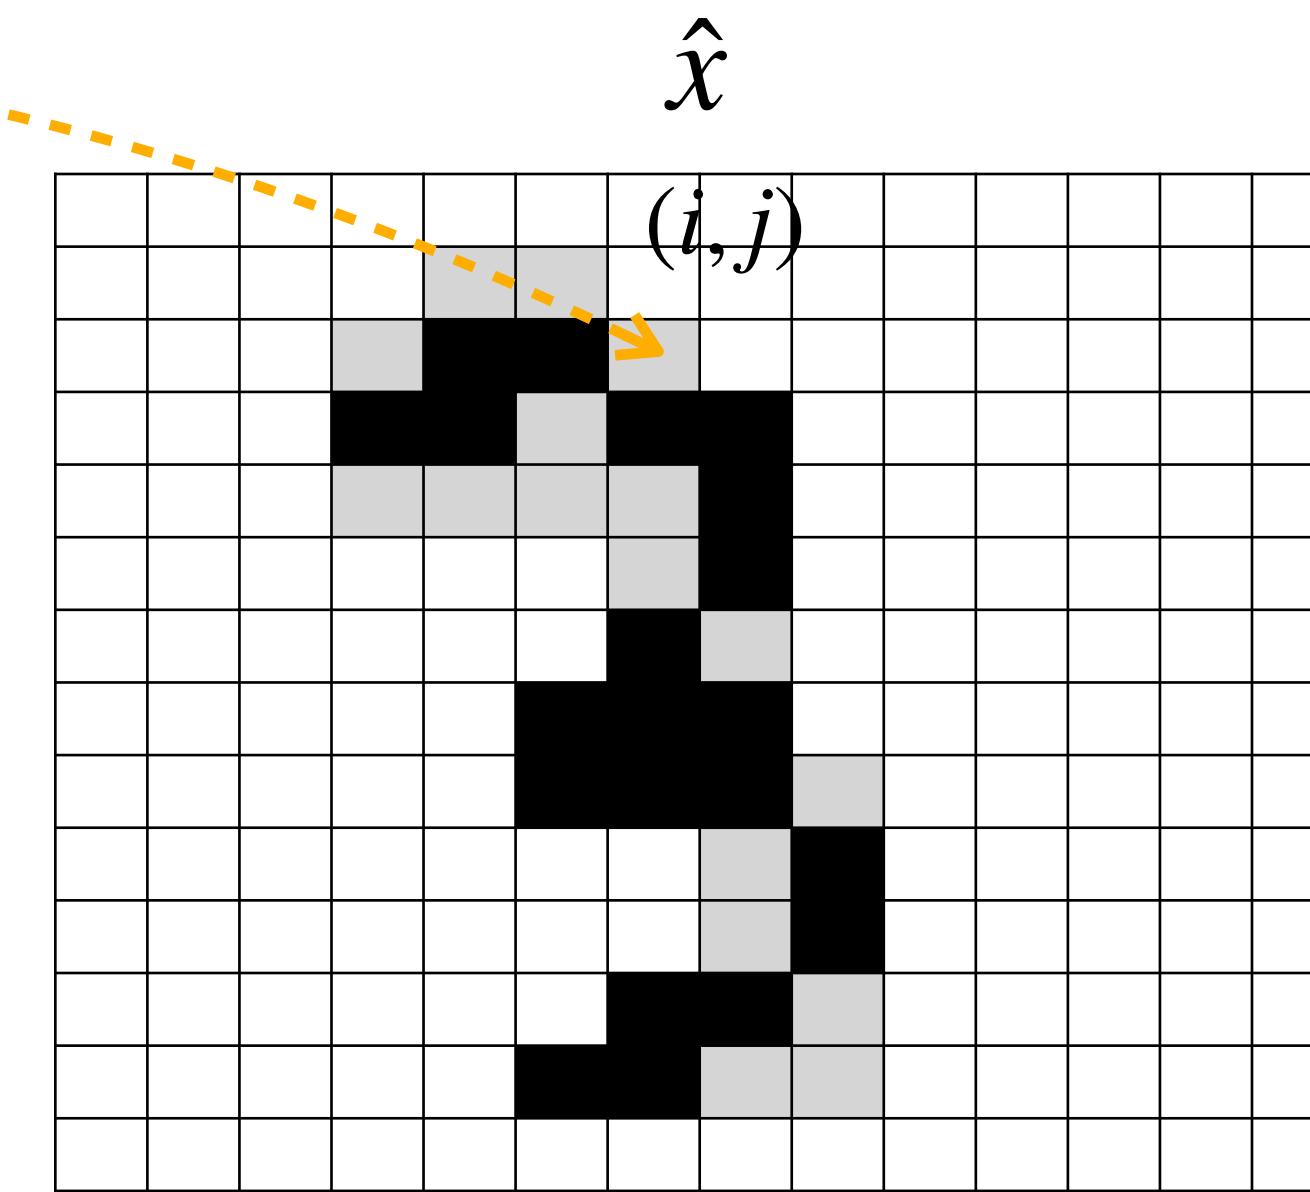
# Resconstruction Loss

We train an autoencoder with a reconstruction loss, which measures the error between the original input image  $x$  from the reconstructed output  $\hat{x}$

Error is computed pixel-by-pixel with MSE (or Binary Cross-Entropy)



$$L(h) = \frac{1}{2wh} \sum_{i=1}^h \sum_{j=1}^w (x_{i,j} - \hat{x}_{i,j})^2$$

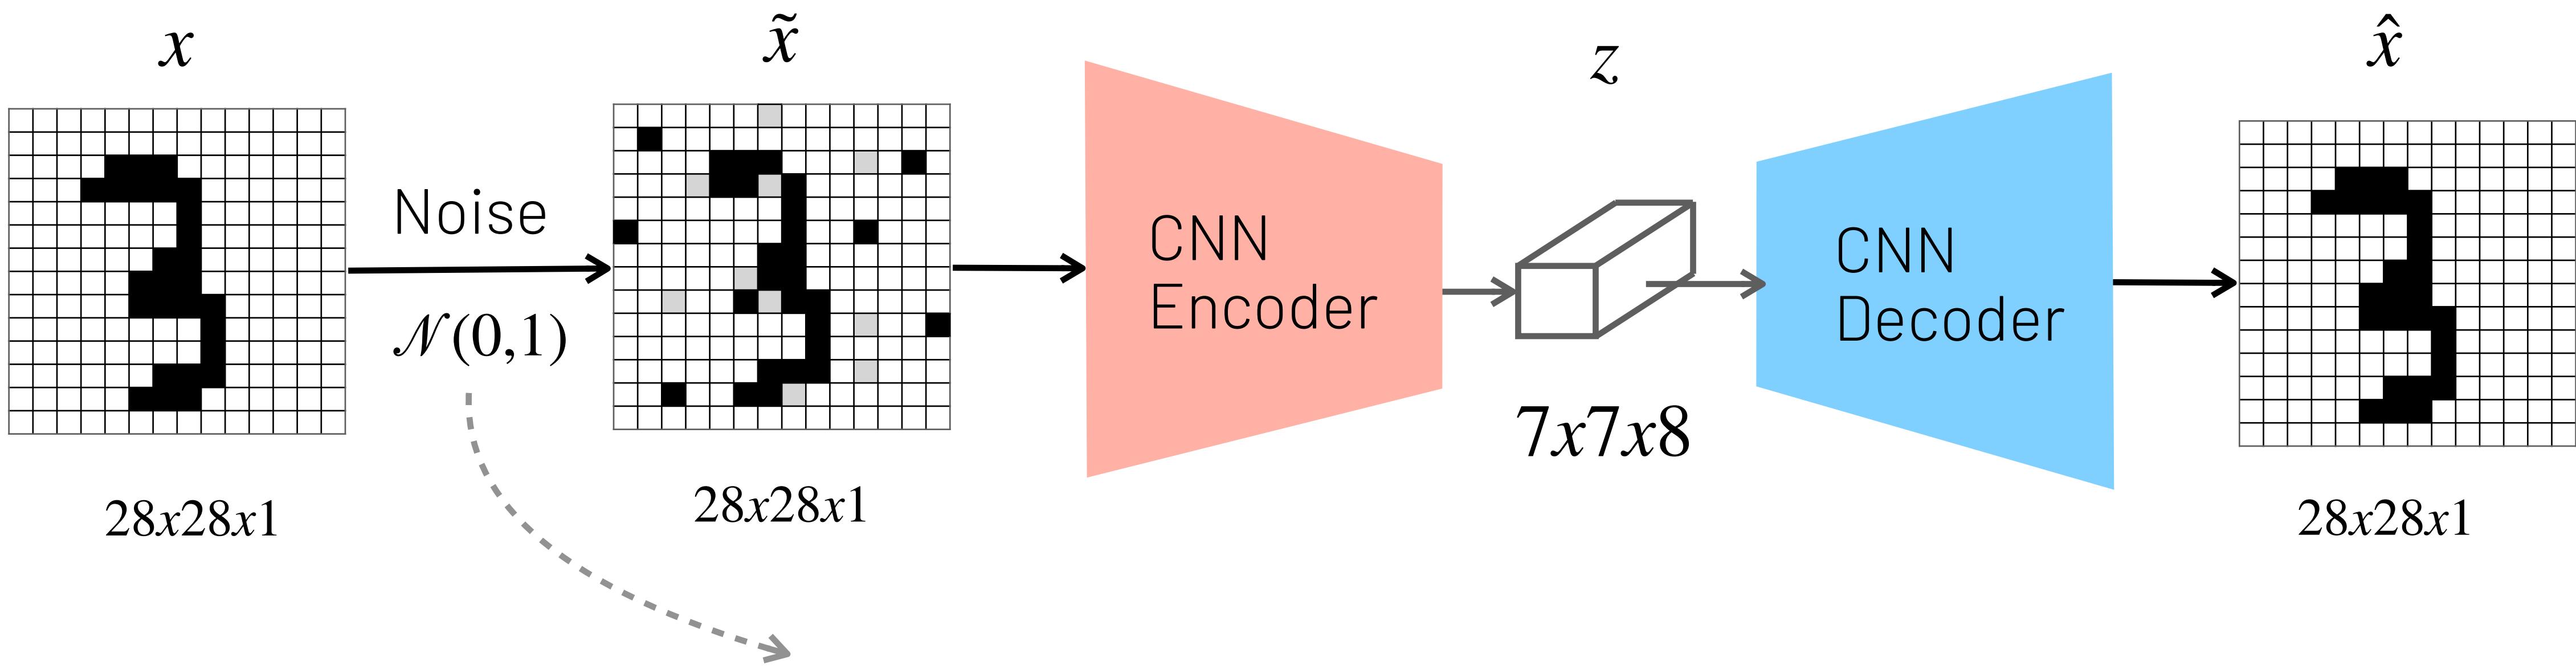


14x14x1

14x14x1

# Denoising Autoencoders

Denoising Autoencoders are a type of autoencoder where we train the model to remove noise from the original data (e.g., image)

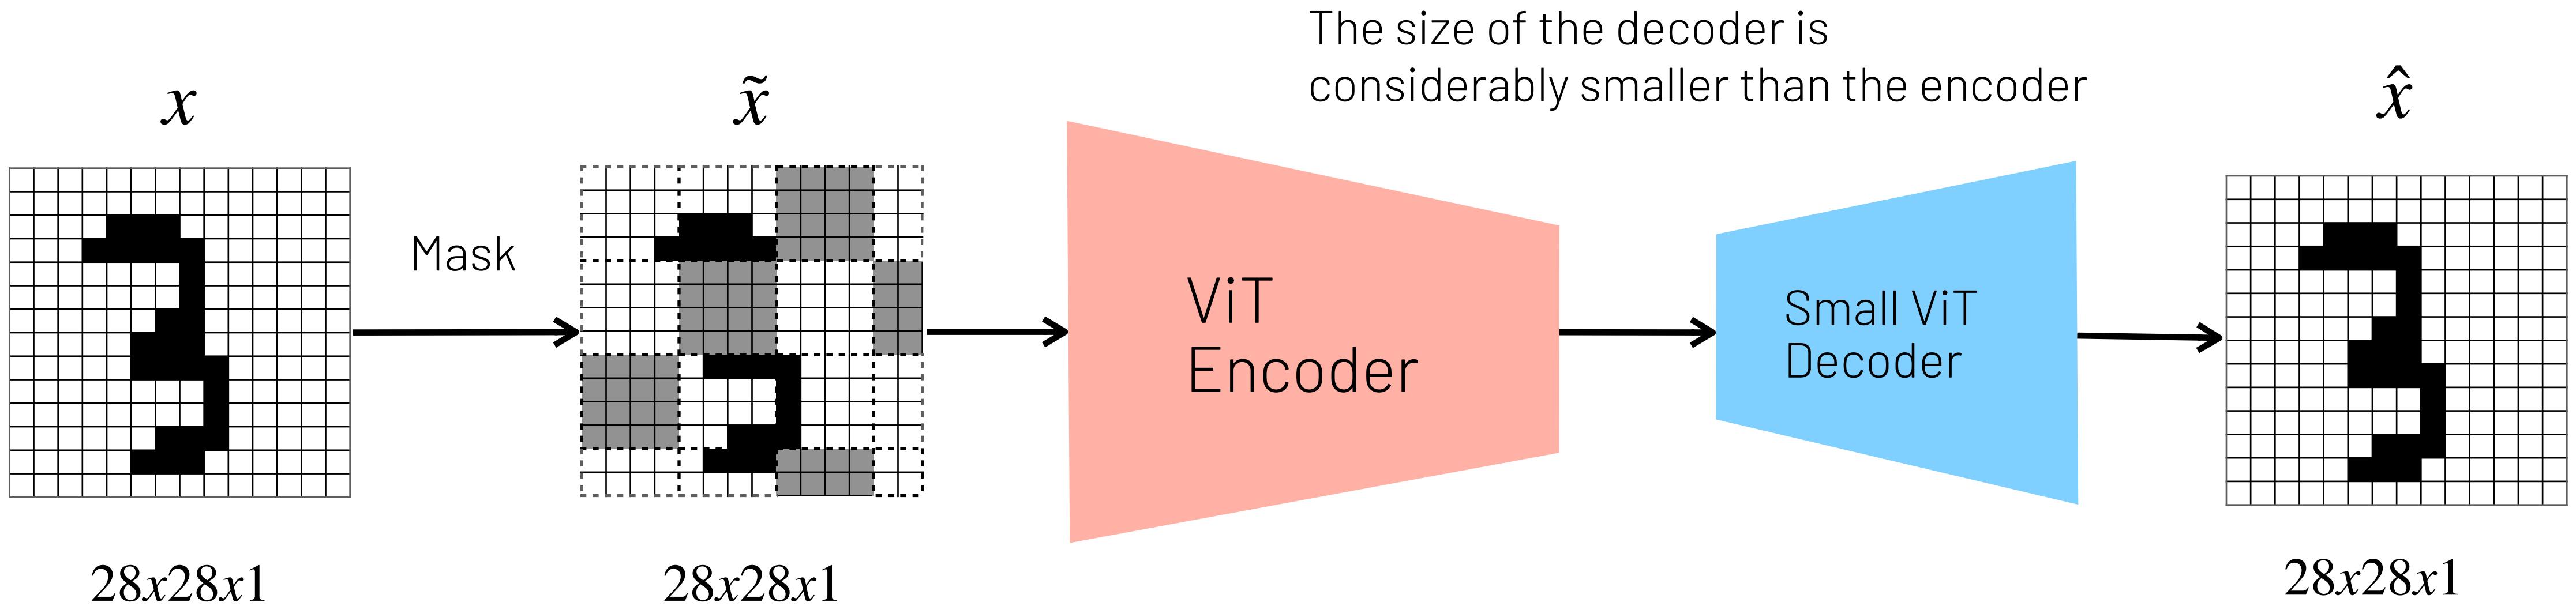


Adding noise forces the model to learn the internal structure of the data. To correct a given pixel, the model has to learn the correlation between its nearby pixels.

1. This task is also called inpainting

# Masked Autoencoders

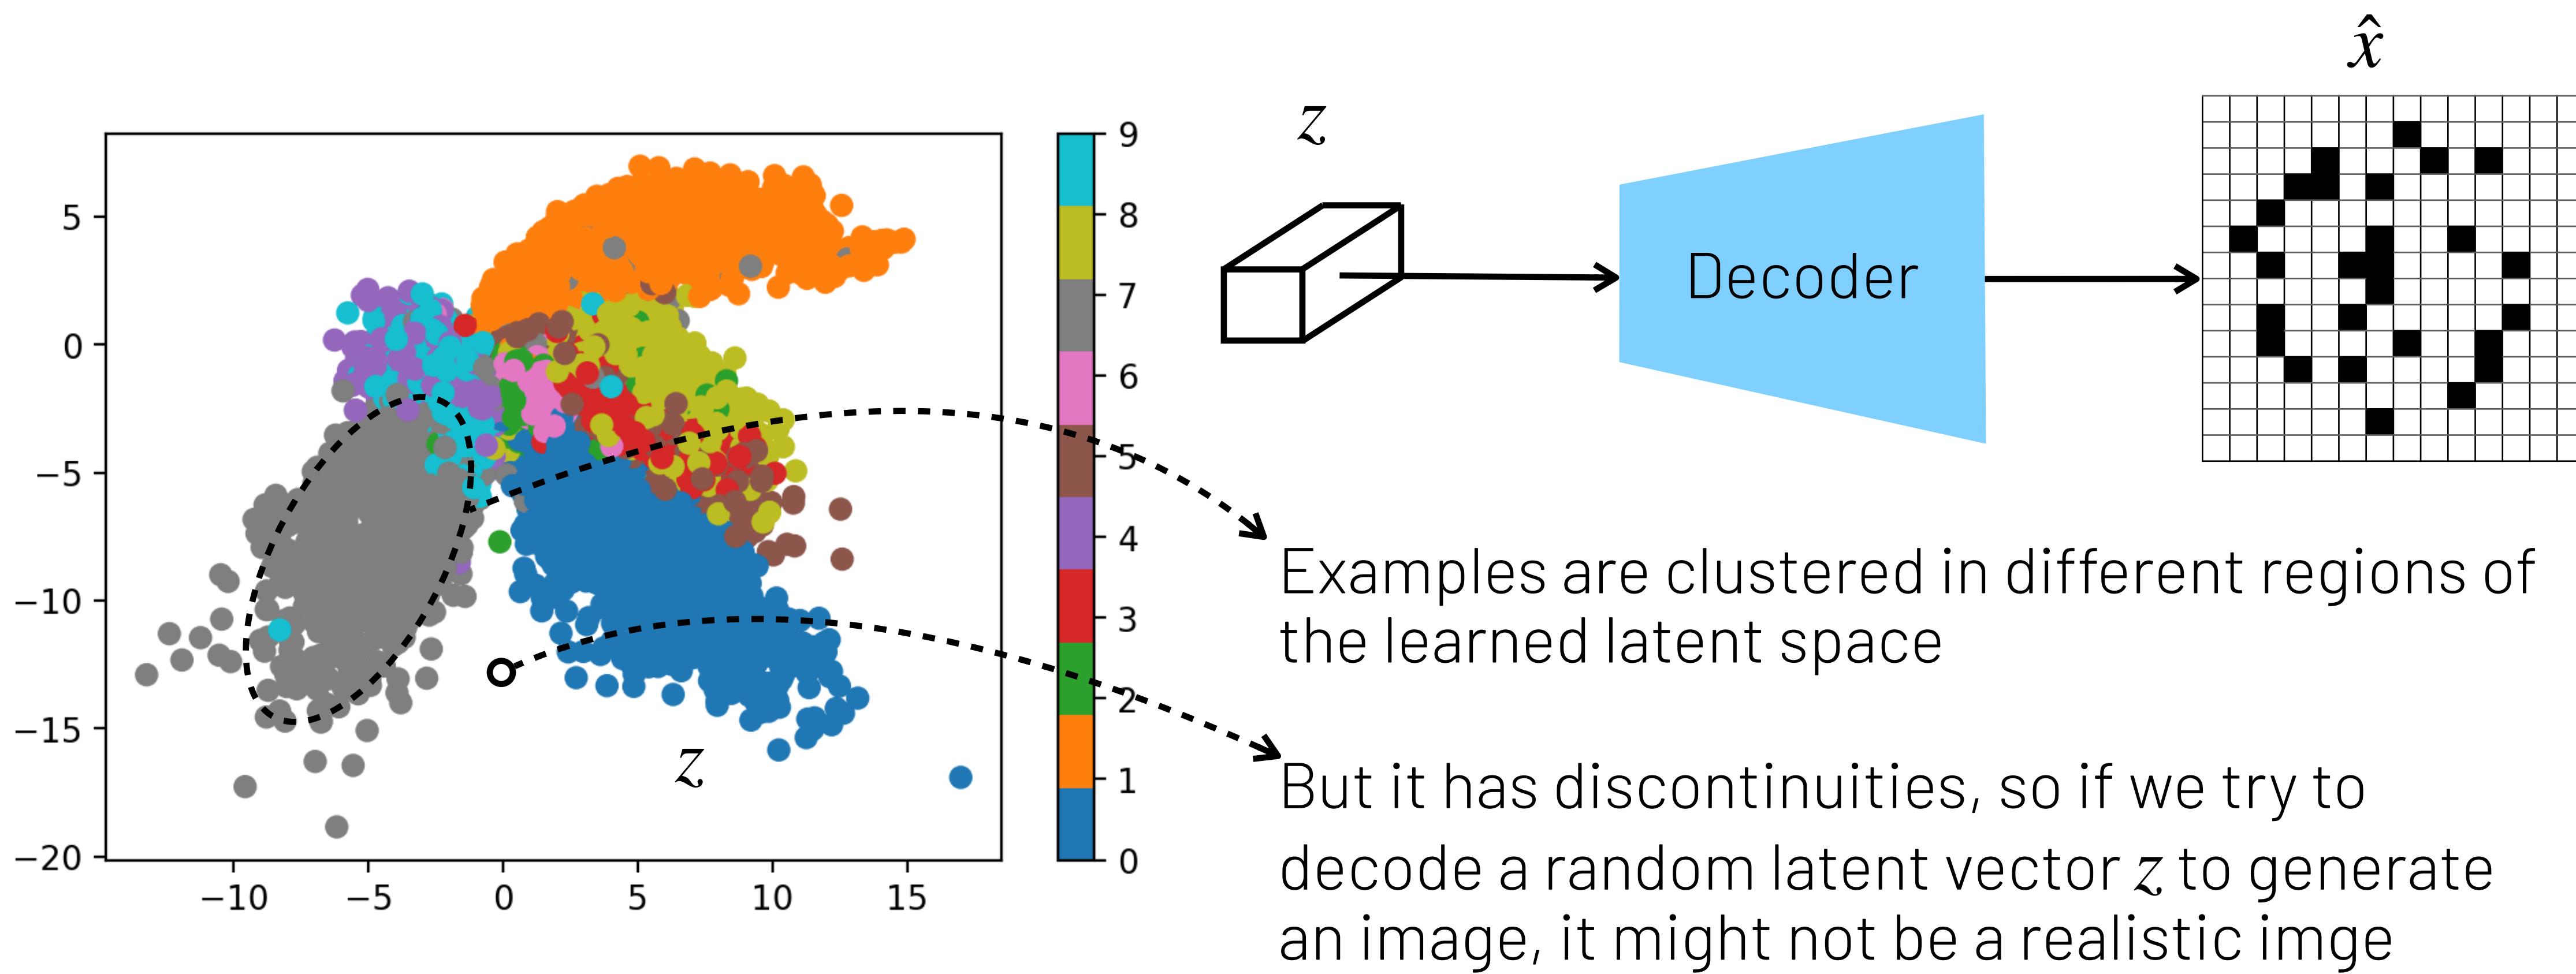
Masked Autoencoders are another type of autoencoder where we train the model to complete<sup>1</sup> a masked patch of the original image (similar to BERT).



Like adding noise, masking is a form of imposing a constraint to avoid the model learning the identity function.

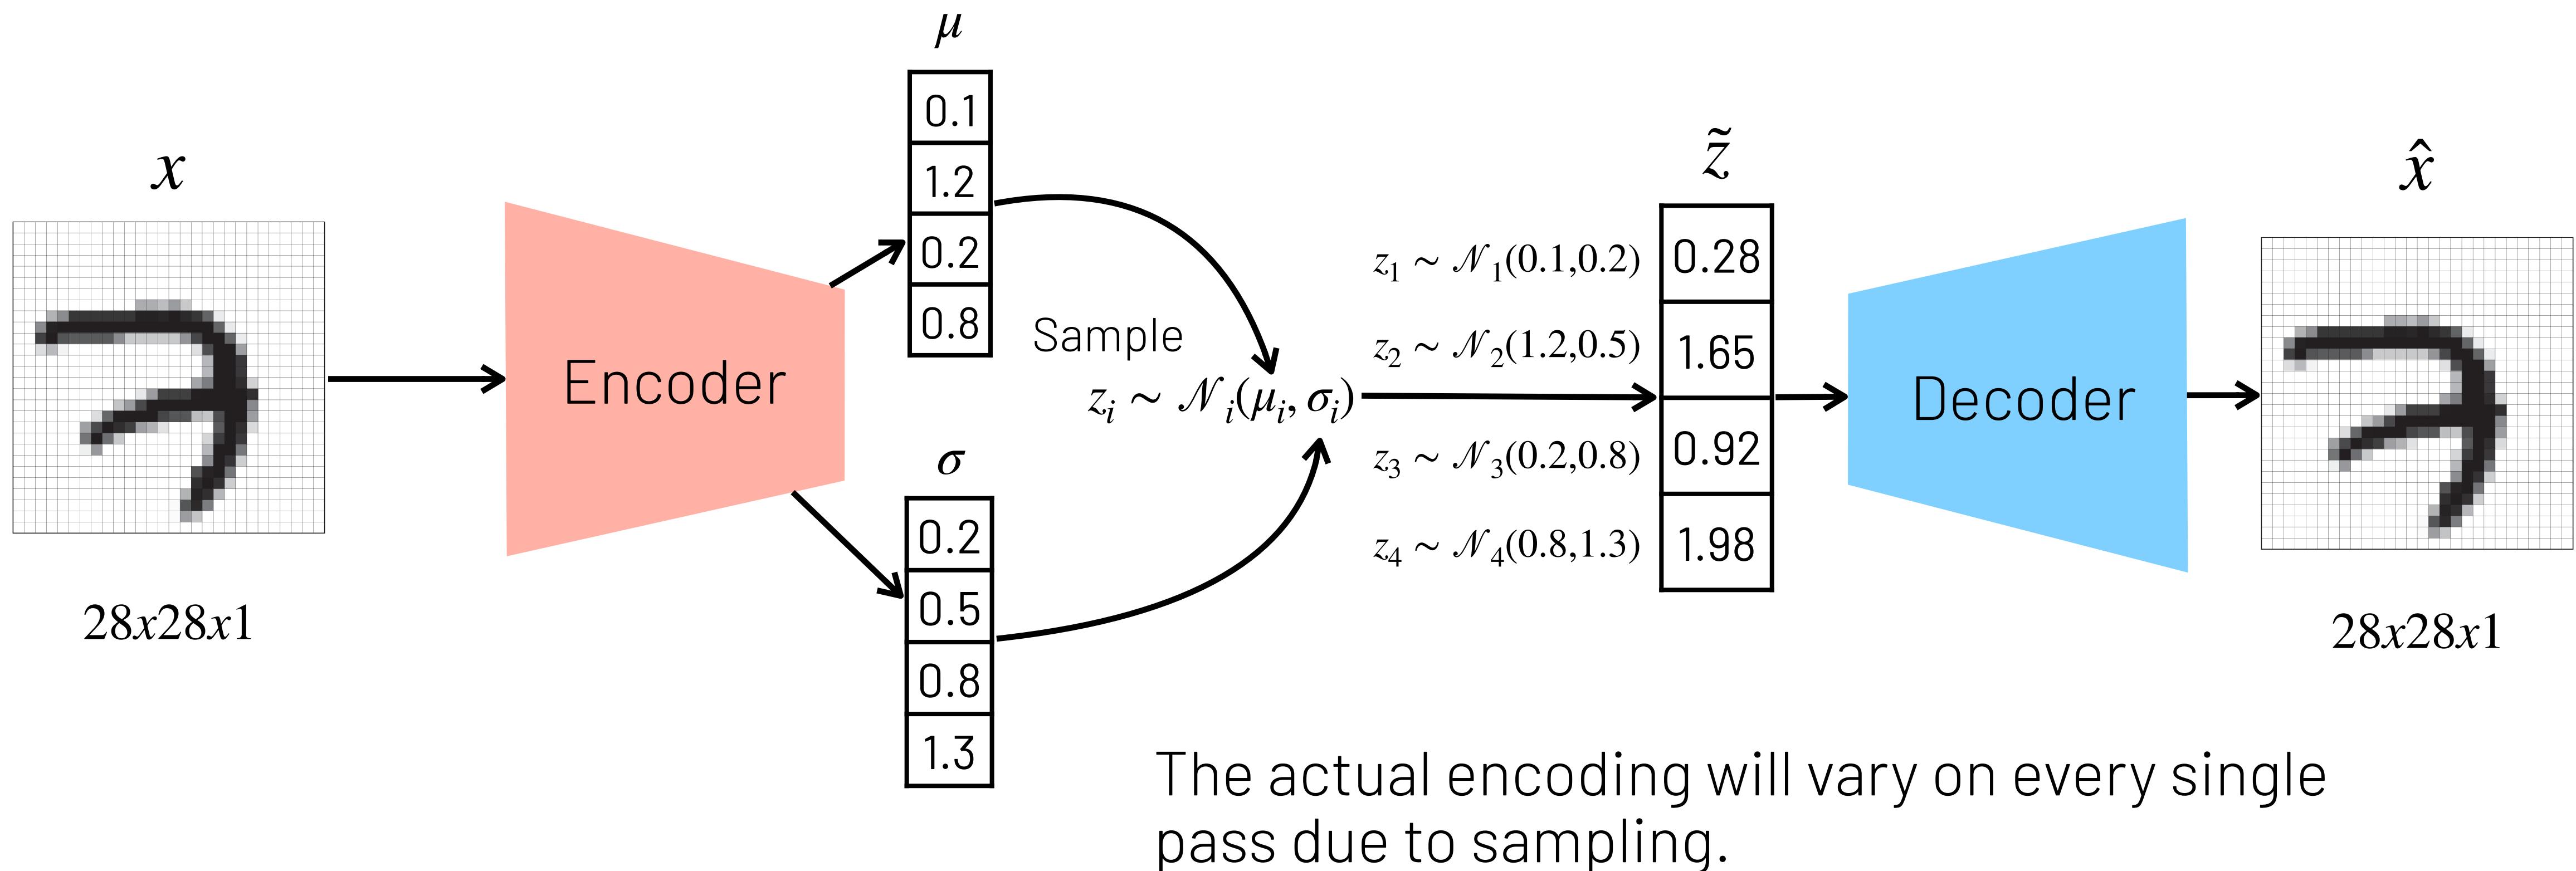
# Decoding random latent images

After training an autoencoder, we could try to decode random latent features maps  $z$  to generate new images, however the learned latent space might not be continuous!



# Variational autoencoders (VAEs)

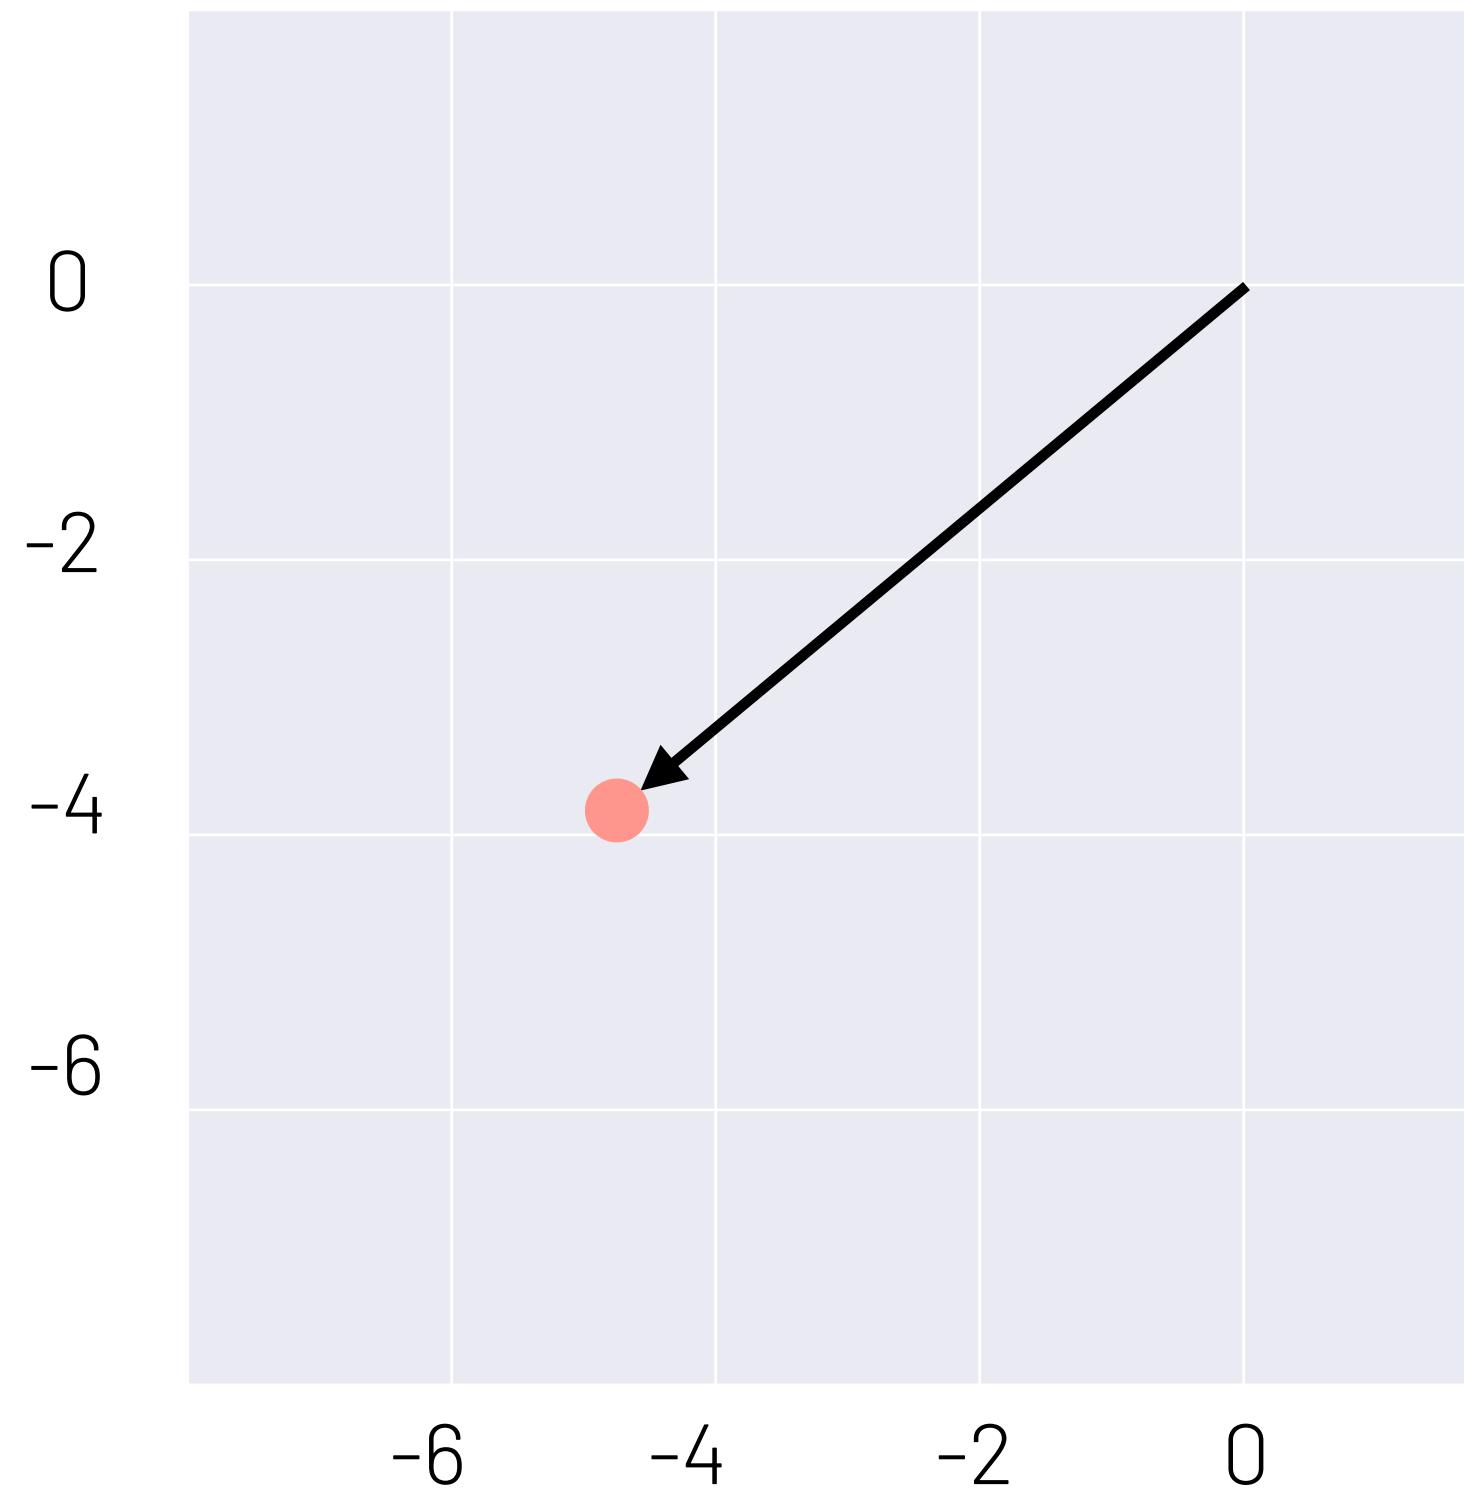
To solve that discontinuity problem, Variational Autoencoders (VAE) use the encoder to learn a probability distribution  $P$  explicitly, and sample the latent representation  $z$  from  $P$



# Autoencoders vs. VAEs

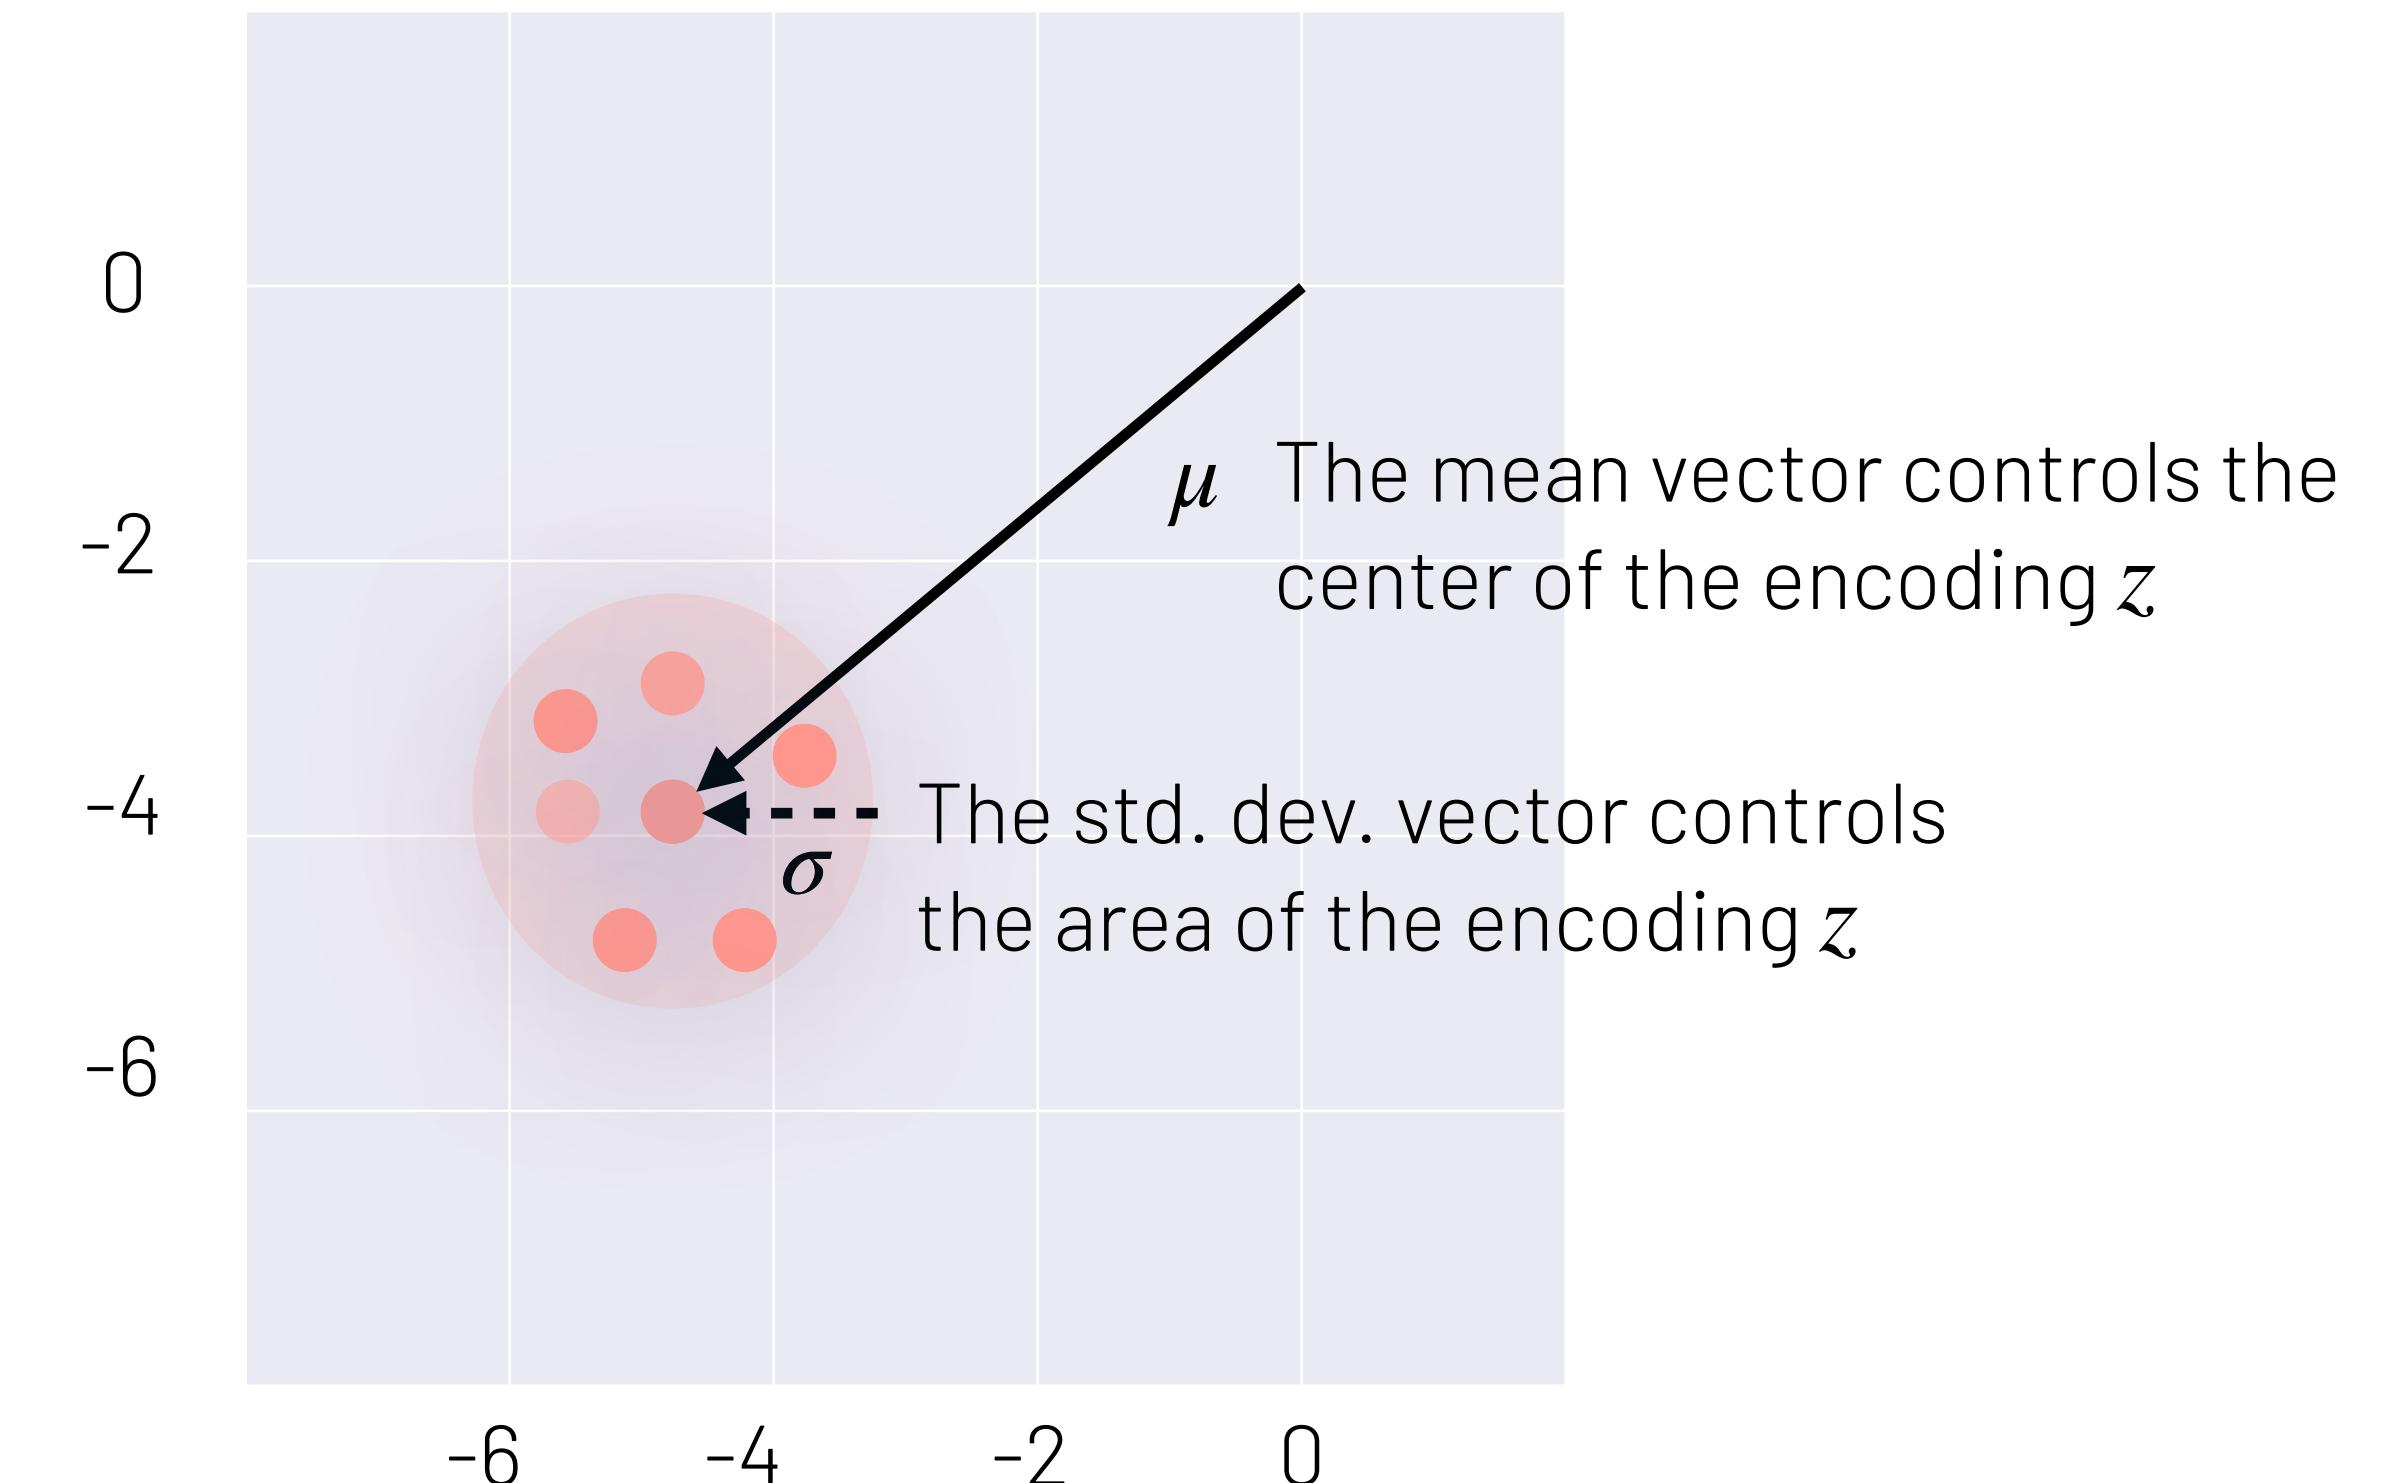
## Autoencoders

(direct encoding coordinates)



## VAEs

( $\mu$  and  $\sigma$  define a probability distribution)

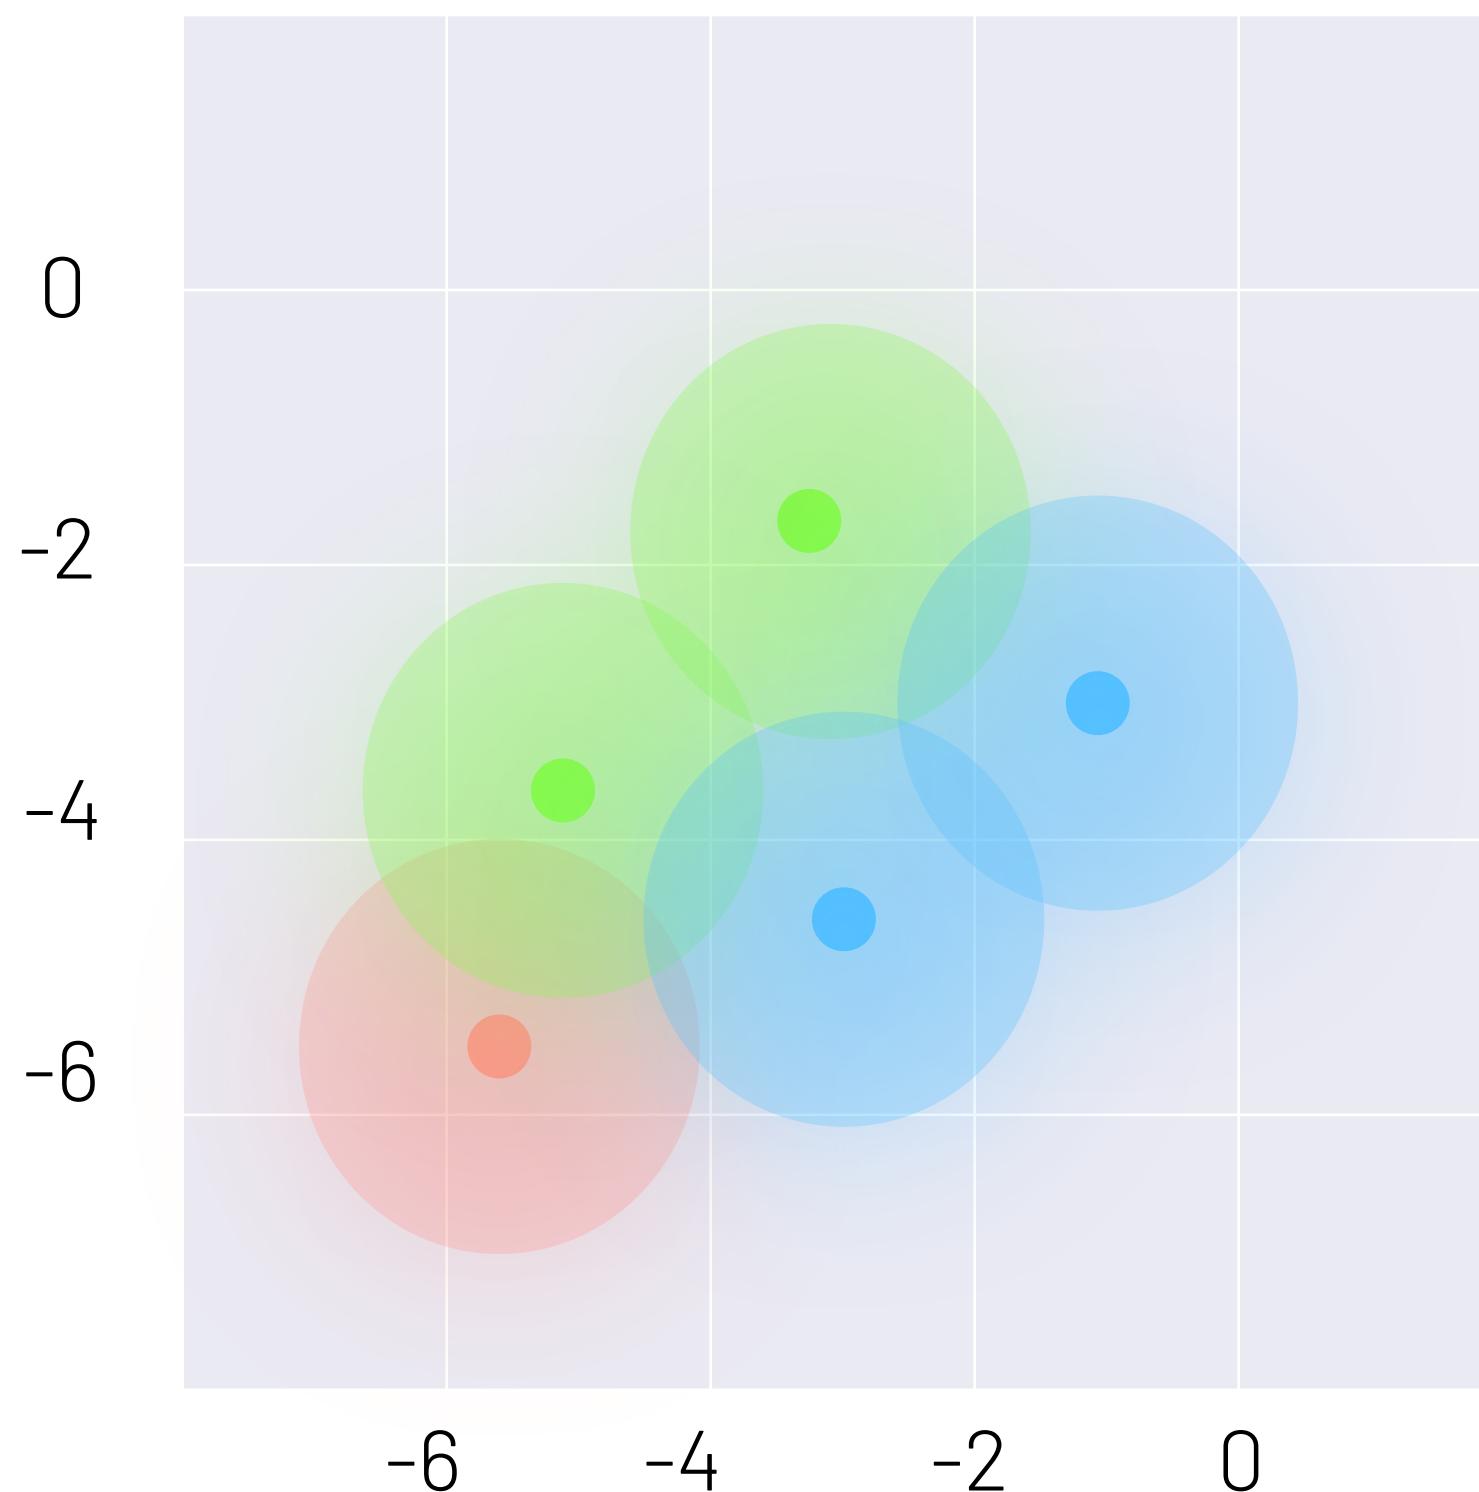


- ▶ Only a single point in latent space refers to a sample of that class (discontinuity)

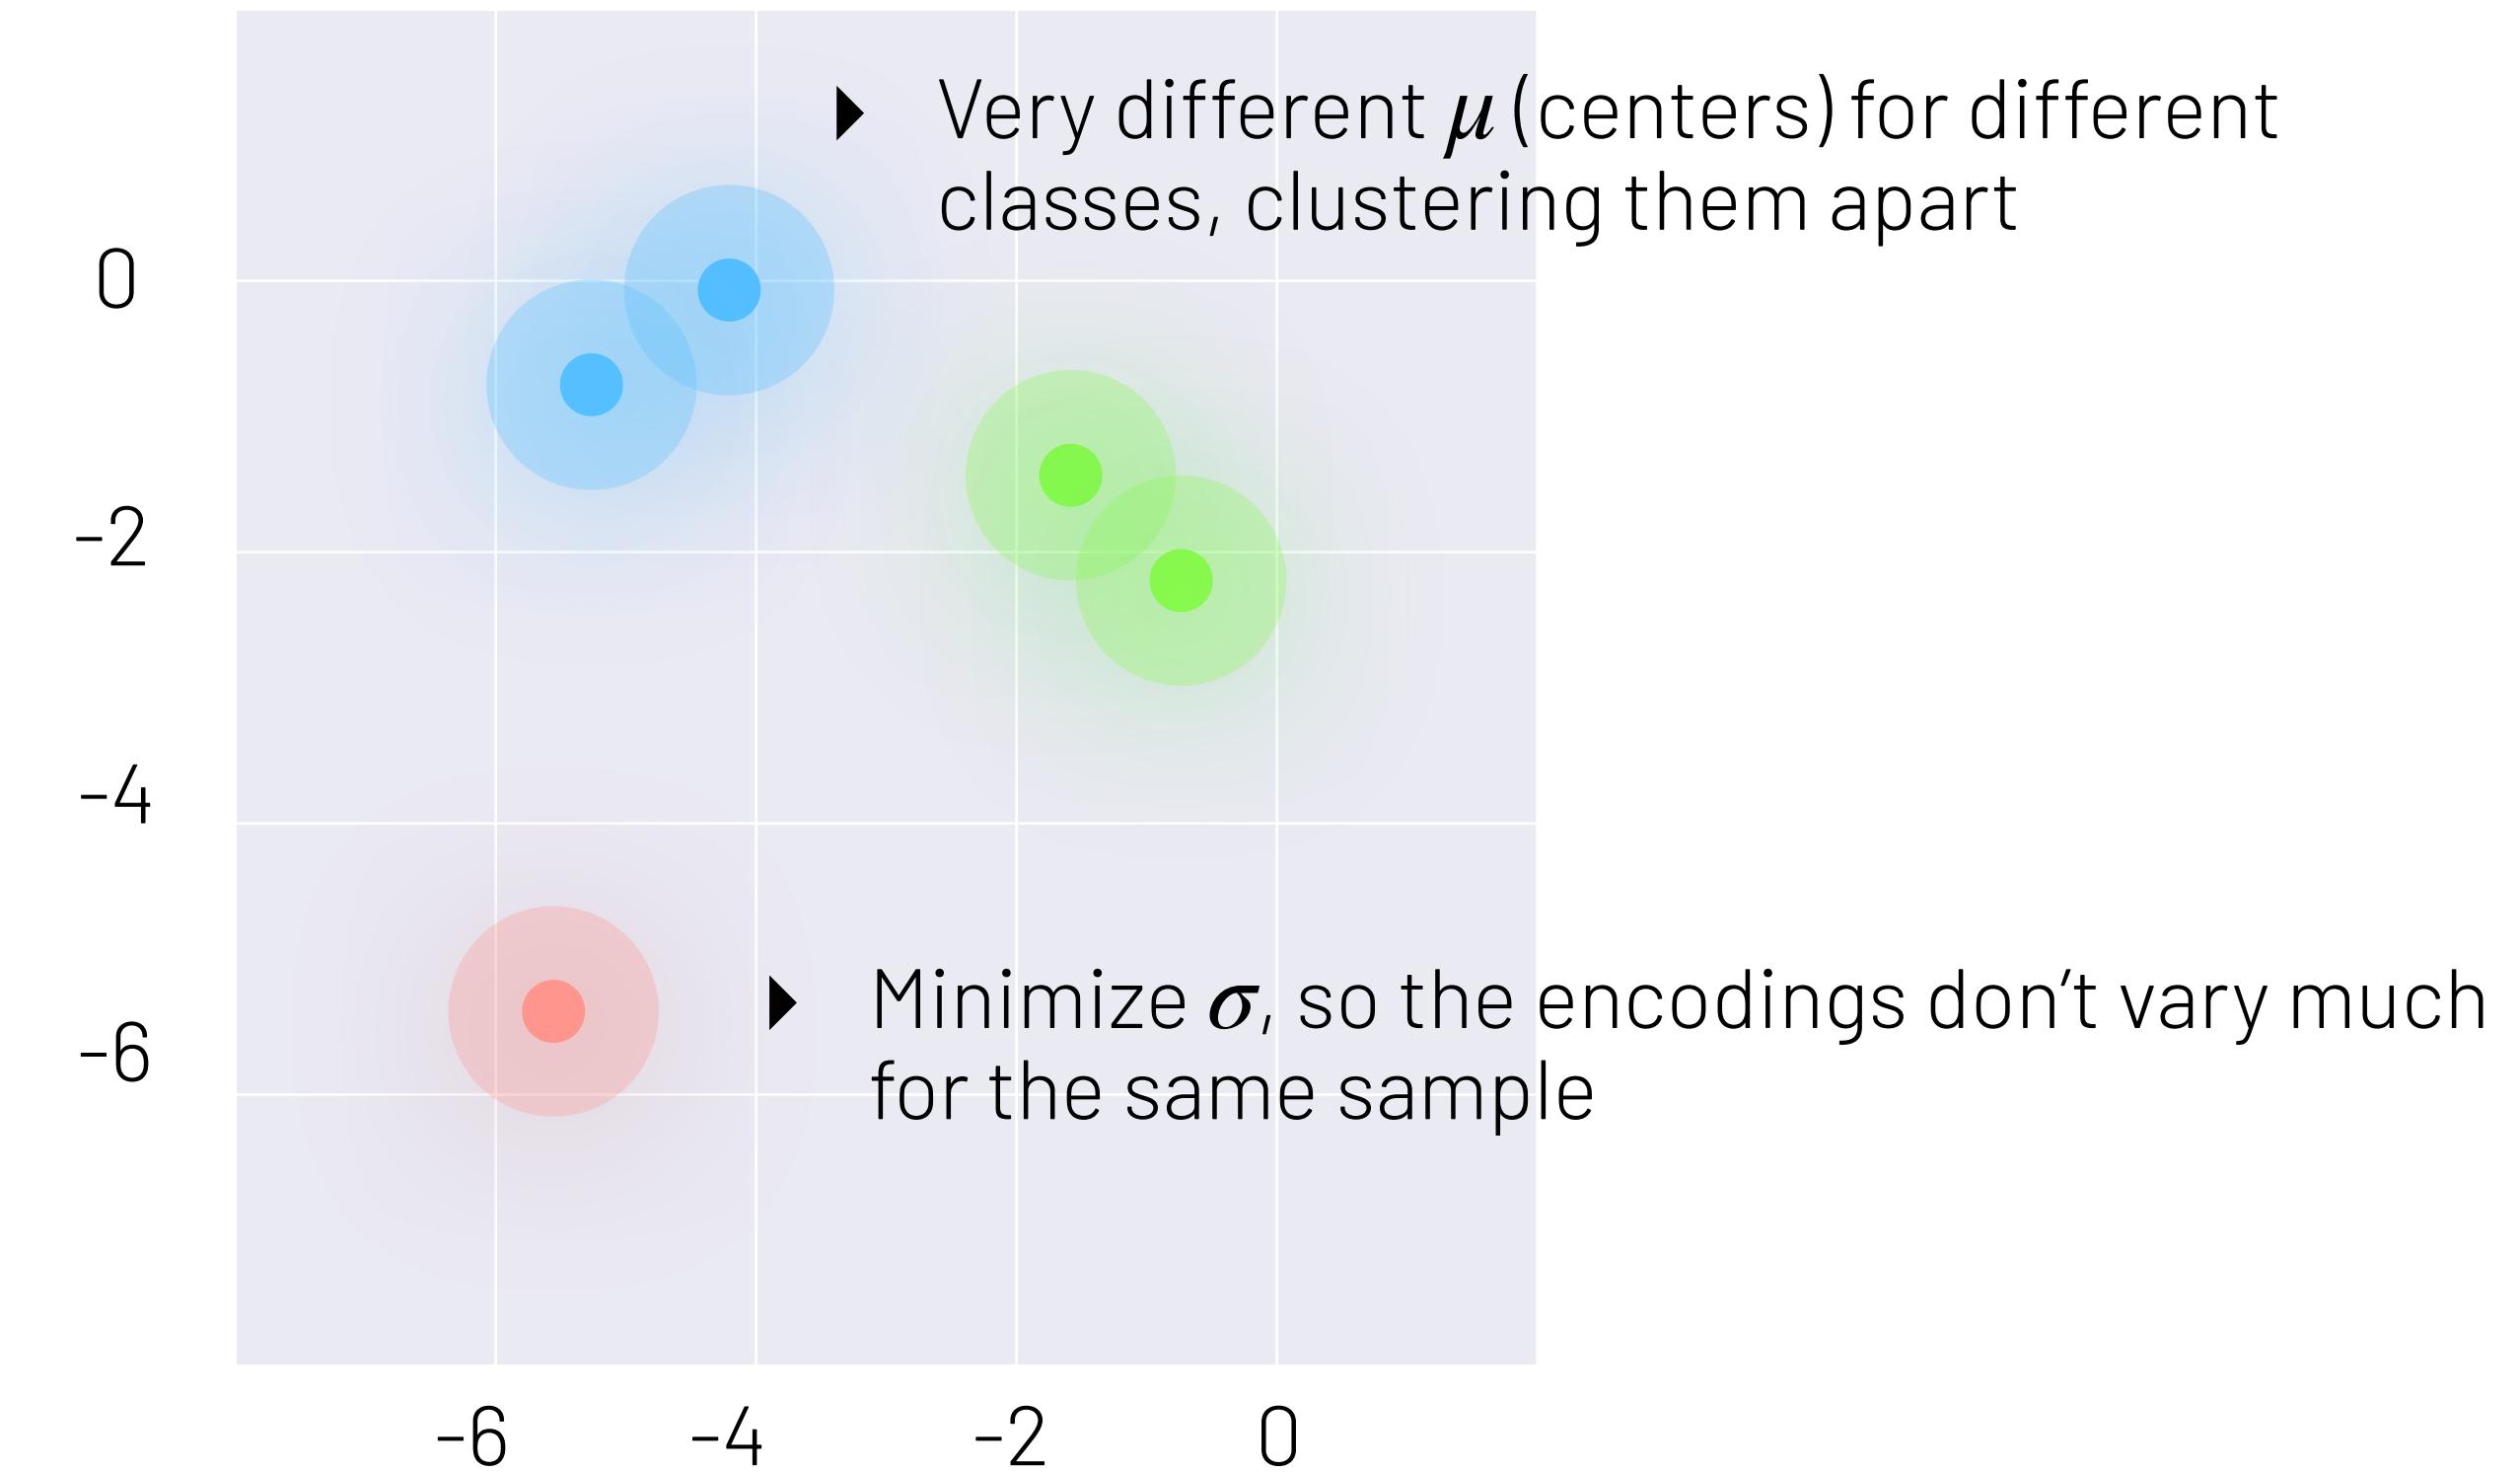
- ▶ All nearby points refer to the same sample of the class (continuity)

# Ideal latent space

Ideally, we want overlap between samples that are not very similar too, in order to interpolate between classes.

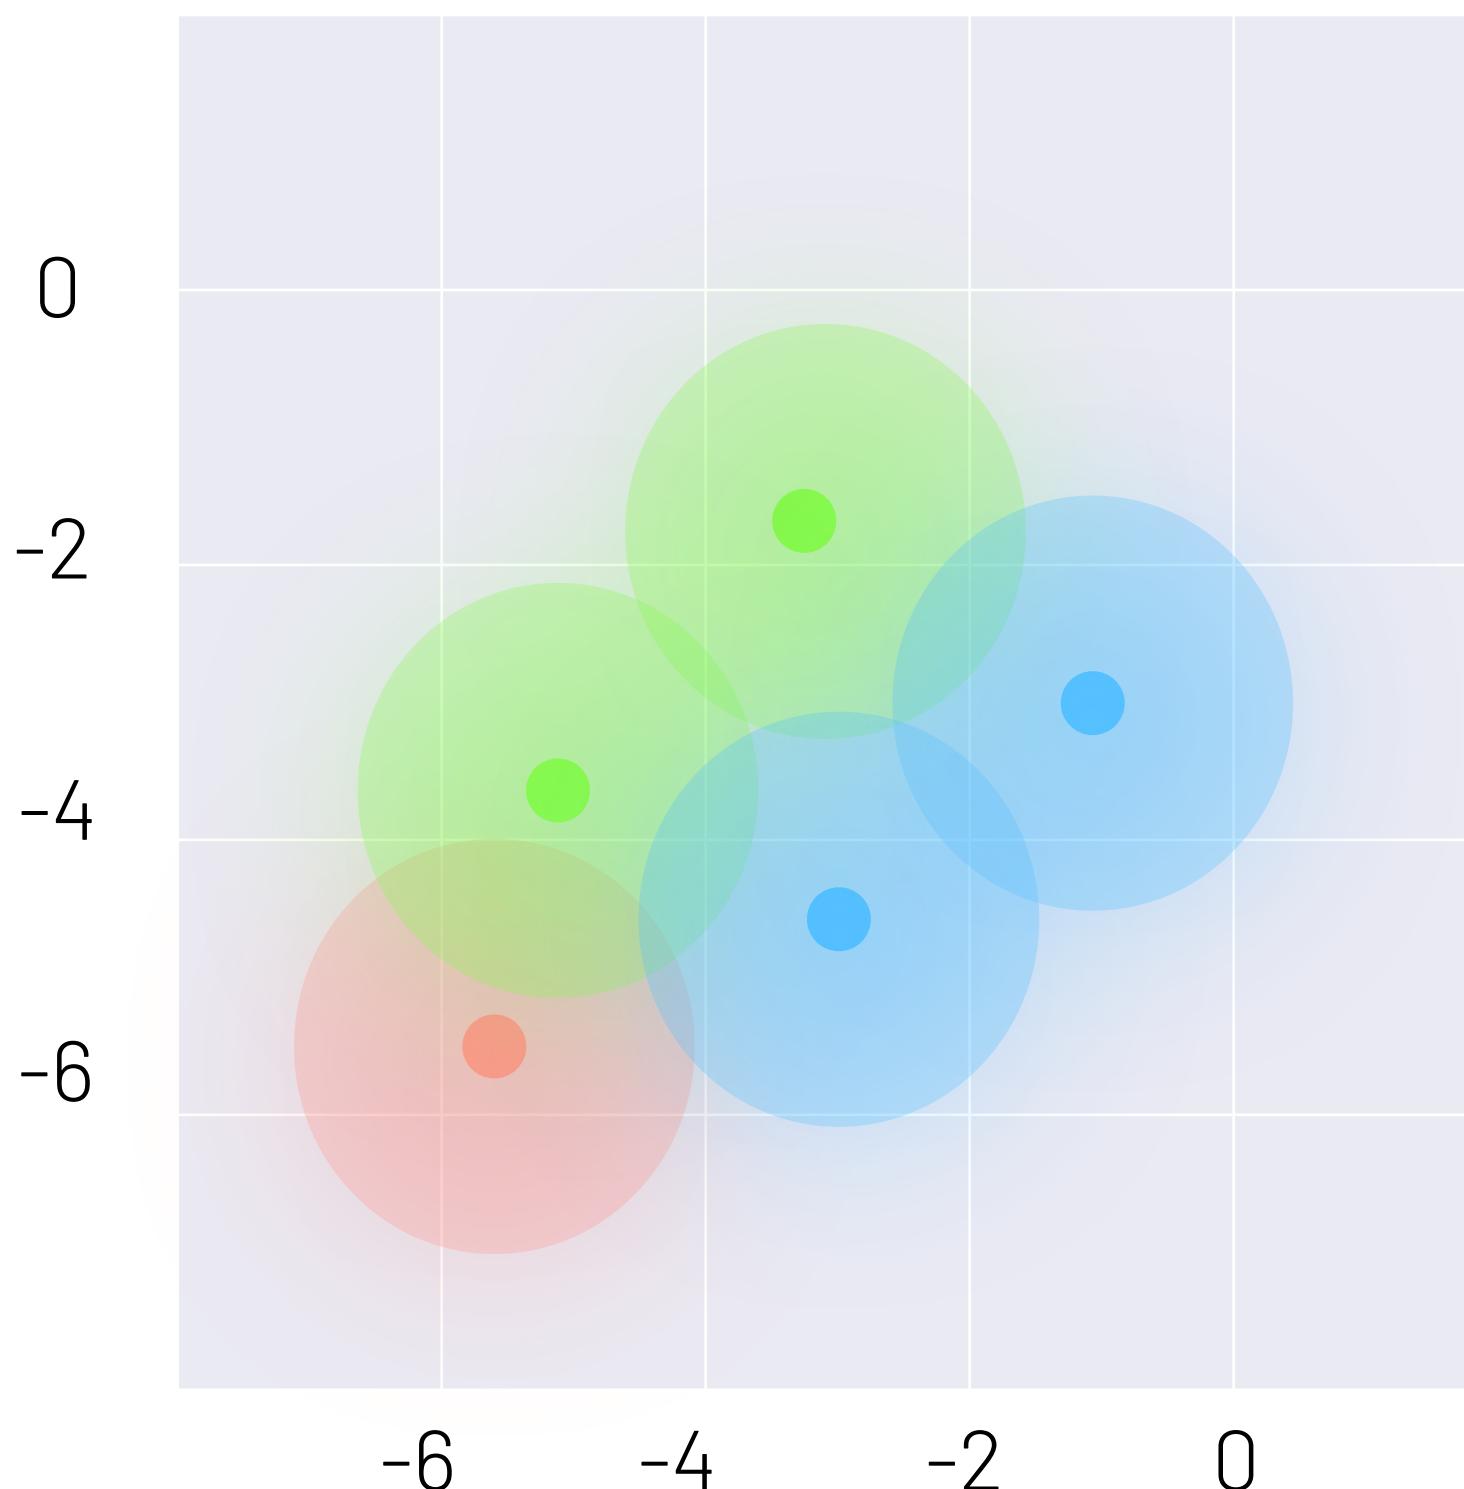


However, since there are *no limits* on what values vectors  $\mu$  and  $\sigma$  can take on, the encoder can learn:



# VAE Loss

To enforce overlap between samples, we add the **Kullback-Leibler divergence** (KL divergence) to the reconstruction loss function:



► Autoencoder Loss :

$$L(h) = MSE(x, \hat{x}) = \frac{1}{2wh} \sum_{i=1}^h \sum_{j=1}^w (x_{i,j} - \hat{x}_{i,j})^2$$

► VAE Loss:

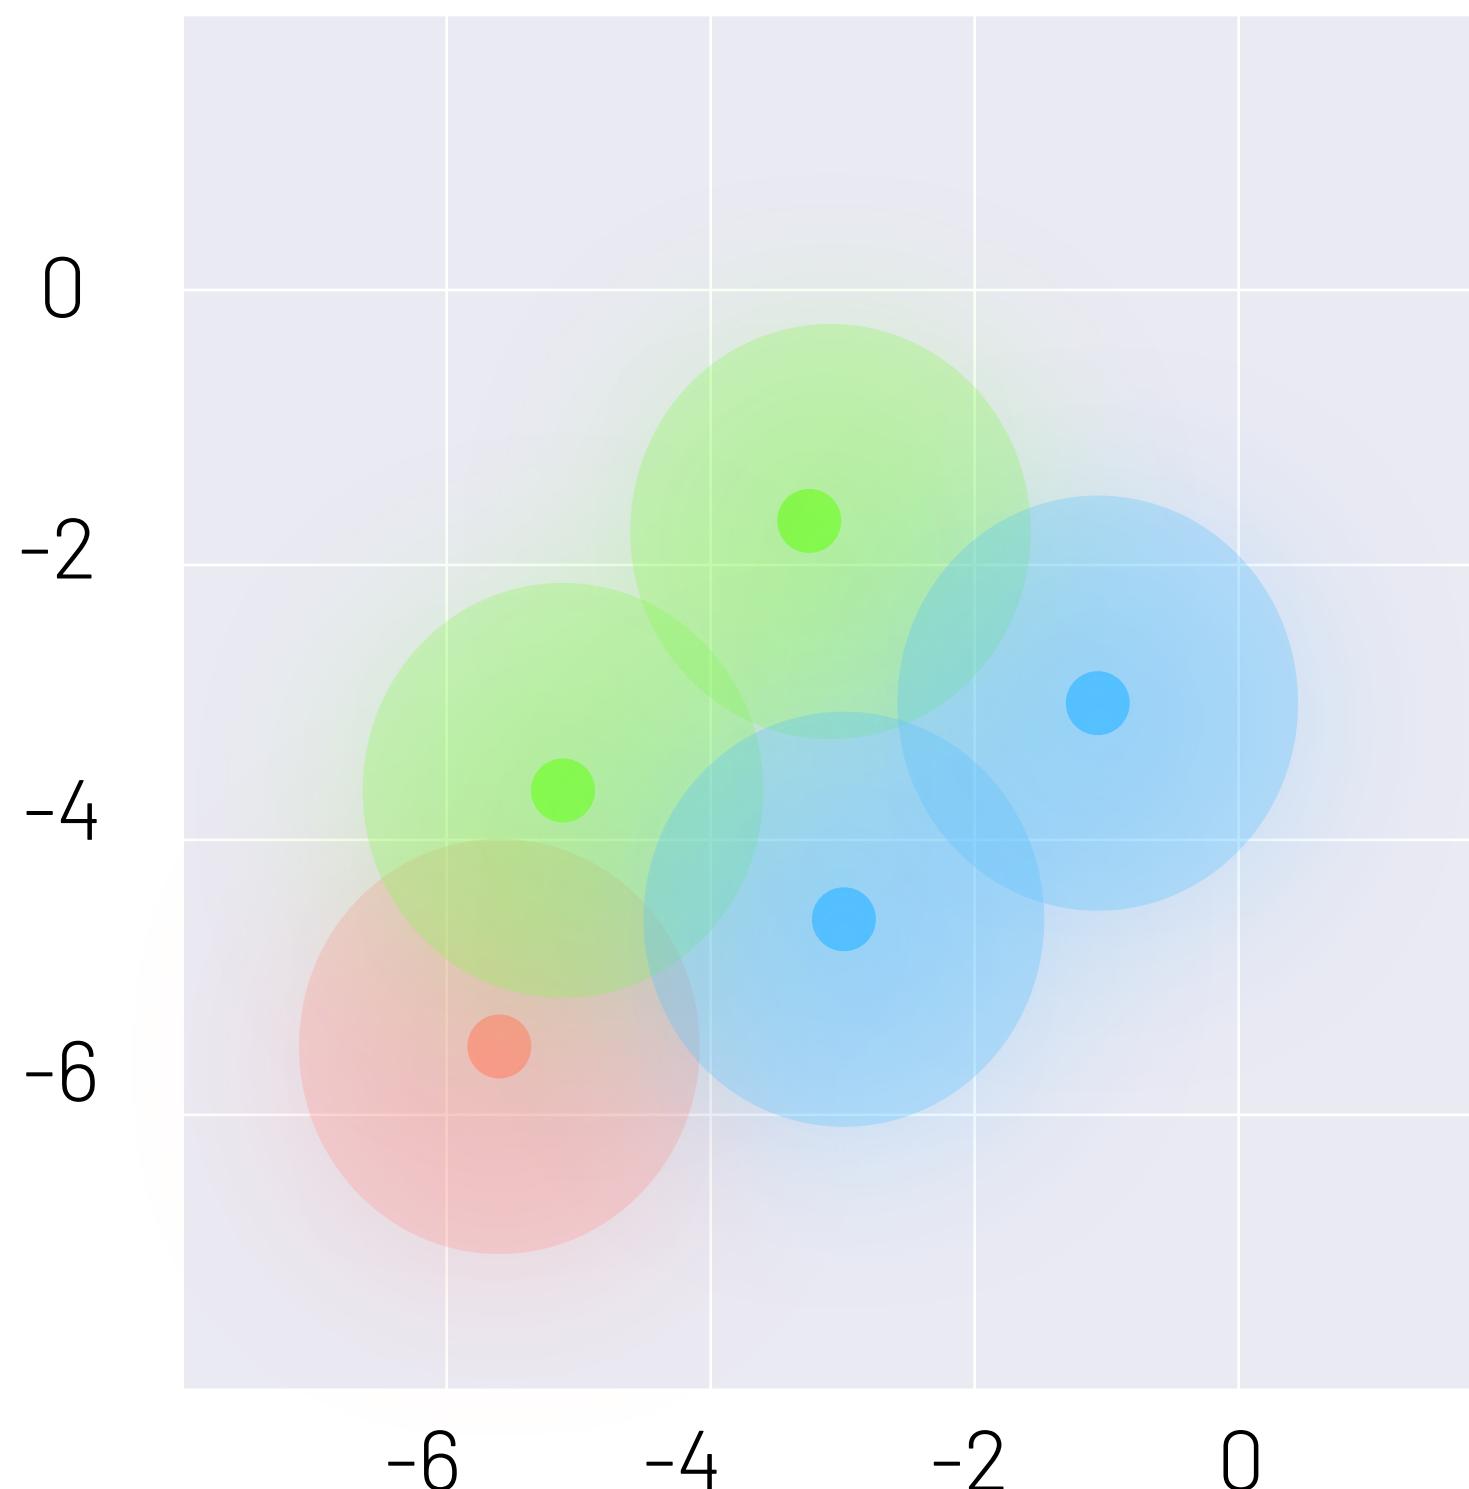
$$L(h) = MSE(x, \hat{x}) + KLD(N(\mu, \sigma), N(0,1))$$

The KL Divergence is a "distance" metric to compare two distributions.

- We want to compare the learned distribution  $N(\mu, \sigma)$  against a standard normal  $N(0,1)$
- This loss encourages the encoder to distribute all encodings evenly around the center of the latent space.

# VAE Loss

To enforce overlap between samples, we add the **Kullback-Leibler Divergence** (KL divergence) to the reconstruction loss function:



► Autoencoder Loss :

$$L(h) = MSE(x, \hat{x}) = \frac{1}{2wh} \sum_{i=1}^h \sum_{j=1}^w (x_{i,j} - \hat{x}_{i,j})^2$$

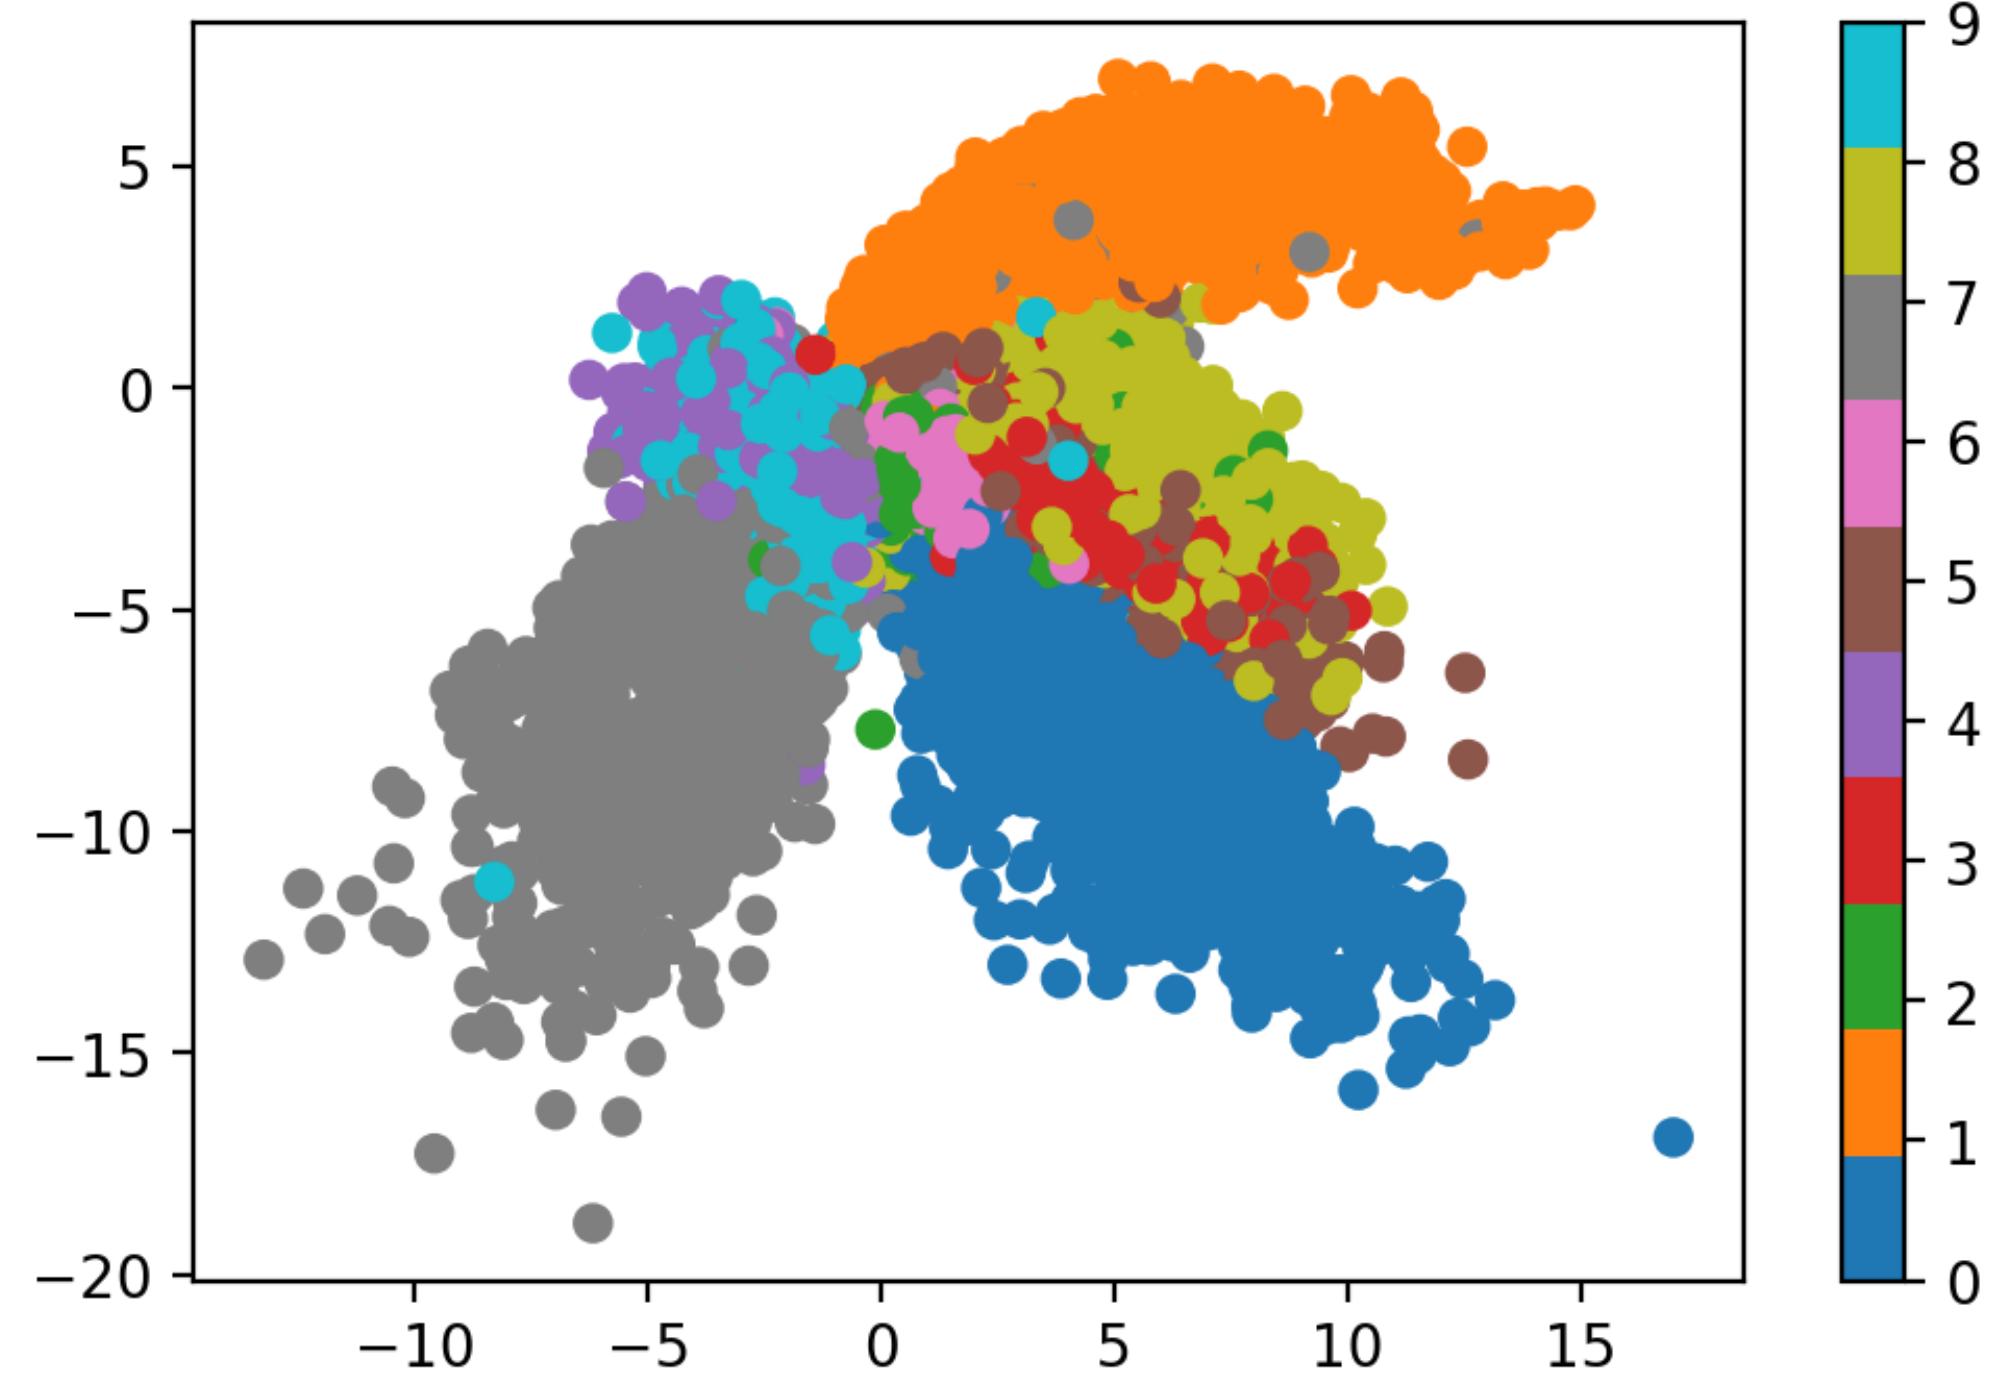
► VAE Loss:

$$L(h) = MSE(x, \hat{x}) + KLD(N(\mu, \sigma), N(0,1))$$

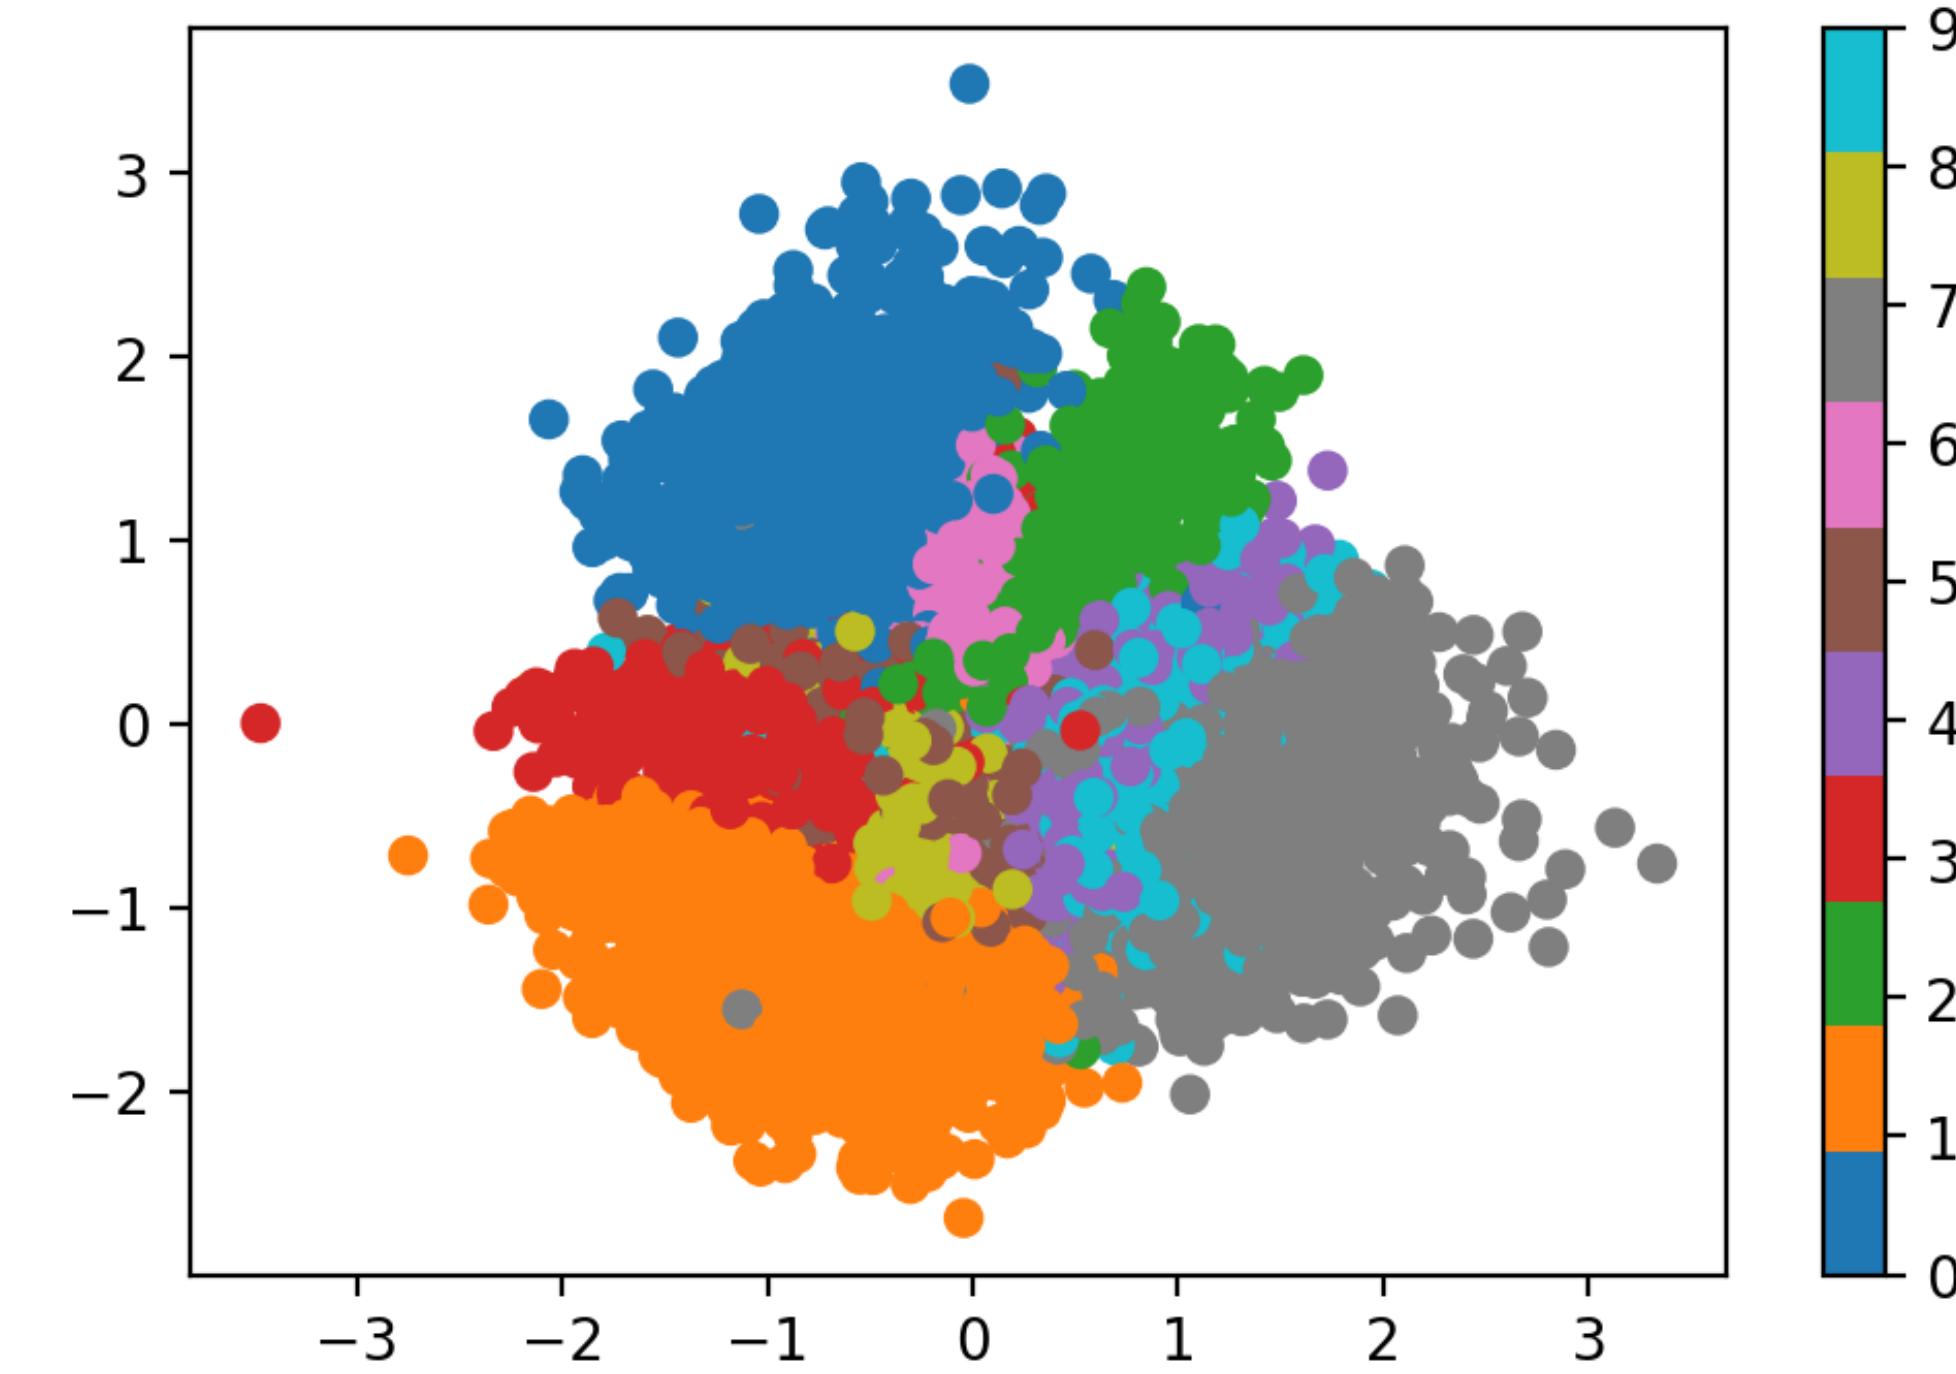
$$KLD(N(\mu, \sigma), N(0,1)) = \sum_{i=1}^n \sigma_i^2 + \mu_i^2 - \log(\sigma_i) - 1$$

# Visualizing the learned latent space

**Autoencoders**



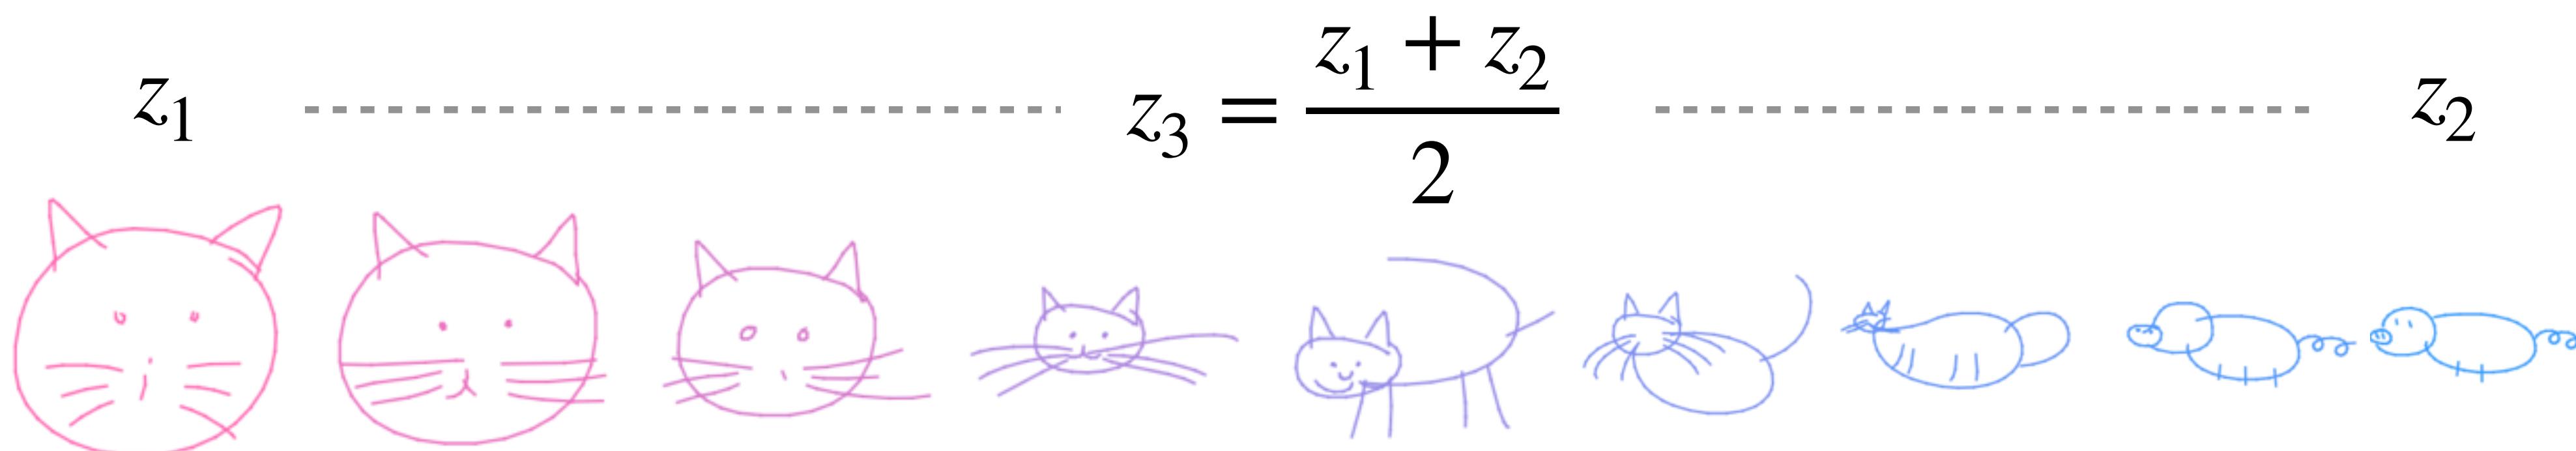
**VAEs**



- ▶ The VAE latent space maintains the similarity of nearby encodings on the *local scale* via clustering
- ▶ Yet *globally*, is very densely packed near the latent space origin

# Interpolating samples in the latent space

Since the space is continuous and smoothly transitions samples classes, we can interpolate different latent vectors  $z_1$  and  $z_2$  and get a semantically meaningful result:



$$\begin{aligned} \text{cat} + (\text{pig} - \text{cat}) &= \text{pig} \\ \text{pig} + (\text{cat} - \text{pig}) &= \text{cat} \end{aligned}$$

# Next Lecture

## L21: GANs

Generating images with Generative Adversarial Networks (GANs)